

Name:

SID:

COMP 4446 / 5046 Quiz 5 (week 12) - version A

Please see the screen for your questions. Note, we have intentionally provided boxes that are larger than necessary. Your answer does not need to fill the whole box.

1. (1 mark) What is the purpose of instruction tuning for large language models?

Solution: Modify the model to do tasks, including new tasks at inference time.

2. (1 mark) Why do we usually collect 2+ labels per example for the test set?

Solution: We want the test set to be very accurate. Collecting multiple labels helps identify mistakes or hard annotation cases.

3. (1 mark) What advantage does Cohen's Kappa have over just counting what fraction of labels match?

Solution: Cohen's Kappa accounts for chance agreement between annotators.

4. (1 mark) Where does the data for preference optimisation come from?

Solution: People looking at pairs of outputs and deciding which they prefer. We also accepted descriptions of the data sources that are specific to preference optimisation (e.g., OpenAI using user queries), and we accepted the reward model.