

Name:

SID:

## COMP 4446 / 5046 Quiz 5 (week 12) - version B

Please see the screen for your questions. Note, we have intentionally provided boxes that are larger than necessary. Your answer does not need to fill the whole box.

1. (1 mark) How does instruction tuning differ from regular fine-tuning?

**Solution:** The data is task instructions and answers rather than just text from the domain of interest.

2. (1 mark) Why do we usually collect only 1 label per example for the training set?

**Solution:** We can make a bigger training set if we only collect one label per example. We also gave partial credit for saying it was to simplify training as there is one reference.

3. (1 mark) If a dataset has a large label imbalance, what impact does that have on Cohen's Kappa?

**Solution:** It will increase the chance agreement value. We also accepted saying it won't influence it as it corrects for chance.

4. (1 mark) Where does the data for instruction tuning come from?

**Solution:** People create examples of tasks and their answers.