

TRIGGER

什么时候用 RETURN NEW 或 RETURN NULL ?		
触发器类型	推荐返回值	含义/说明
BEFORE INSERT/UPDATE	RETURN NEW 或 RETURN NULL	可以修改或取消即将执行的插入/更新操作。 NULL 表示取消操作。
AFTER INSERT/UPDATE/DELETE	RETURN NEW (或 RETURN OLD)	必须返回一个记录,但返回值不会被使用,只是语法要求。

1. We can write a trigger to update the enrolment number for a unit of study offering when a student is added (in Transcript), as shown below:

```
CREATE OR REPLACE FUNCTION updateEnrol() RETURNS trigger AS $$
BEGIN
    UPDATE UoSOffering U
    SET enrolment = enrolment+1
    WHERE U.uosCOde = NEW.uosCode AND semester = NEW.semester
        AND year = NEW.year;
    RETURN NEW;
END; $$ LANGUAGE plpgsql;
```

```
CREATE TRIGGER UpdateEnrol AFTER INSERT ON Transcript
FOR EACH ROW EXECUTE PROCEDURE updateEnrol();
```

你在判断一个unit的
max enrolment
就发生了

```
create or replace function maxcountsb() returns boolean as $maxc$
begin
    if ((SELECT COUNT(sailors.sid) FROM Sailors)
        + (SELECT COUNT(boats.bid) FROM boats) < 9)
    then return true;
    else
    return false;
    end if;
end;
$maxc$ language plpgsql;
```

5. 对于所有形式为 $X \twoheadrightarrow Y$ 的多值依赖，必须满足以下至少一个条件：

a. trivial MVD:

- $X \twoheadrightarrow Y$ 是trivial MVD，即满足以下任一条件：
 - $Y \subseteq X$ ，即Y是X的子集，或
 - $X \cup Y = R$ ，即X和Y的并集包含关系R的所有属性。

b. 超键条件：

- X是关系R的超键

employee_name	project_id	personal_phone_number
Bob	P1	047012345
Bob	P3	046098765
Bob	P1	046098765
Bob	P3	047012345
Lily	P1	045067543
Fiona	P7	043085432

Is this relation in 4NF?

No: There is **at least one non-trivial multivalued dependency**

存在多值依赖: $employee_name \twoheadrightarrow Project_id$,

而Y不是X的子集，并且X UNION Y不是R

同时 $employee_name$ is not a superkey.

MVD 和 4NF 不同

Based on the data in the relation, is this a valid MVD?

$UoS \twoheadrightarrow Tutor$

Note that according to the values in the relation, the relationship between the UoS and Tutor is **independent** from the relationship between UoS and Textbook. This means that the above MVD is valid. This also implies that the Tutor of a UoS is selected independently by the school.

Assume a new Tutor, Lijun C is added for the UoS COMP9120. What must happen to maintain this independence (i.e., MVD)?

- Add one row for each different textbook with Tutor *Lijun C* of that UoS.

<u>UoS</u>	<u>Textbook</u>	<u>Tutor</u>
COMP9120	Silberschatz	Ying Z
COMP9120	Widom	Ying Z
COMP9120	Silberschatz	Mohammad P
COMP9120	Widom	Mohammad P
COMP9120	Silberschatz	Alan F
COMP9120	Widom	Alan F
COMP5110	Silberschatz	Ying Z
COMP5110	Silberschatz	Mohammad P
COMP9120	Widom	Lijun C
COMP9120	Silberschatz	Lijun C

ONLY IN LHS -> PART OF ALL KEYS

1. 满足atomic — From 1FN
2. No partial Dependencies
A partial dependency is a non-trivial functional dependency $X \rightarrow Y$ in relation R , where X is a strict (proper) subset of some key for R , and Y is not part of any key.
也就是说, 不能只有主键中的一部分functional determine另外一个非主键element
3. A relation schema R is in Third Normal Form (3NF) if for every functional dependency $X \rightarrow Y \in F^+$, at least one of the following holds:
 - a. $X \rightarrow Y$ is a trivial functional dependency: $Y \subseteq X$
 - b. X is a superkey for R
 - c. $Y \subseteq$ some candidate key of R简而言之就是: "every non-key attribute is not transitively functionally dependent on a key"; 每个非key attribute不能间接依赖于key; 这里是允许函数依赖的"箭头"不从超键出发 (因为c) |
4. 对于所有non-trivial $X \rightarrow Y$, X 是 R 的一个superkey, 即所有函数依赖的"箭头"必须从超键出发

b) Is the relation in 3NF? ↵

gate->airline does NOT meet the 3NF restriction that **either the LHS is a superkey or at least for the RHS to be part of a key**. (hence cannot be in 3NF). In addition, 还要检查 2NF 和 1NF, contact->name 不满足 2NF(Partial dependenc), 因为 contact 是 candidate key 中的一个; 下面是老师给的例子

- c) Explain whether it is a lossless-join decomposition to decompose the relation into the following:
- R1(destination, departs, gate)
 - R2(contact, departs, pickup)
 - R3(gate, airline)
 - R4(contact, name)

这里因为 R3 和 R4 是直接和本来的 FD 相关, 因此可以从 R 中删除 airline, name 这 2 个 attributes, 变成 R7(dest, dep, gate, contact, pickup); 这里之所以不删除 gate 和 contact 是因为 R3 和 R7 的 intersection 为 gate 才能保证 gate 是 R3 的 superkey, 对于 R4 同理。↵

现在来看 R7 (dest, dep, gate, contact, pickup), 他如果分解成 R1 和 R2, 则它们的 intersection 是(departs), 根据现有的 FDs, 我们足以推断出 departs 的 closure 不能包含所有在 R1 或 R2 里面的 attributes, 因此不是 lossless 的。

Atomicity — log[↵]

- A real-world transaction is expected to happen or not happen at all.
- 通过 ROLLBACK 达成[↵]

Isolation — Concurrently, interleaved[↵]

- Transactions can run concurrently; meaning their operations can be interleaved.[↵]

Incorrect Summary

T2 reads X after N is subtracted and reads Y before N is added: an **incorrect summary** is obtained.[↵]

Lost update

```
1. 事务 T1 读取某个数据项 (如 balance = 100)。  
2. 事务 T2 也读取相同的数据项 (balance = 100)。  
3. T1 修改数据 (如 balance = balance + 50 → 150), 但尚未提交。  
4. T2 也修改数据 (如 balance = balance - 20 → 80), 并先提交。  
5. T1 提交后, balance 被覆盖为 150, 而 T2 的修改 (-20) 丢失了。  
  
最终结果应该是 130 (100 + 50 - 20), 但由于并发问题, 实际结果是 150 (T2 的更新被 T1 覆盖)。
```

Read Uncommitted — Dirty Read[↵]

- 允许事务读取其他尚未提交事务的数据, 也叫做 Dirty Read。[↵]
- PostgreSQL 不支持此隔离级别, 直接转换为 Read Committed。[↵]

Read Committed — Solve Temporary update[↵]

- 每次查询只能读取已经提交的数据。no dirty reads.[↵]
- 防止脏读, 但在同一个事务中两次读取同一行数据, 可能得到不同的结果。[↵]
- 是 PostgreSQL 的默认隔离级别。[↵]

Repeatable Read — Solve incorrect summary problems[↵]

- 只能“看到”在它开始之前已经提交的事务所做的更改。[↵]
- transaction 看不到: [↵]
 - 其他事务还没有提交的数据 — 避免 Read Uncommitted[↵]
 - 以及它执行期间, 其他事务新提交的更改 — 避免 Read Committed[↵]

2PL[↵]

- 只能先加锁(growing), 再释放锁(Shrinking)[↵]
- 在没有死锁和没有事务失败的前提下, 2PL 能够确保调度的执行是等价于某种串行执行顺序的, 即满足可串行性要求[↵]
- Basic[↵]
 - 可以在 commit 前释放锁[↵]
- Strict[↵]
 - 只能在 commit 后释放锁[↵]

Serializable

1. 事务的可见性 (MVCC)

- 事务T2开始时, 会拿到一个快照时间点 (snapshot), 只看到那个时间点之前提交的所有数据版本。
- 事务T1如果在T2开始后提交了更新, T2默认看不到T1的更新, 除非T2重启或使用特定隔离级别。

2. 那为什么串行化隔离还能避免丢失更新?

- 当T2尝试修改数据 (写操作) 时, 数据库会做序列化冲突检测, 发现T2的快照版本过旧 (和T1提交的版本冲突)。
- 这时数据库会拒绝T2的提交, 抛出 `serialization_failure` 错误。
- 事务T2必须回滚并重新执行, 这样它才能看到最新版本的x, 重新基于最新版本更新。

Blocking factor: ↵

- $b = \text{block size} / \text{record size}$ ↵
- Record spanning: happens when $b < 1$. Means “record size is bigger than the block size”; It is **not good design** when this happened ↵

Seek time: the arm assembly is moved in or out to position a read/write head on a desired track; Go to the right track ↵

Rotational delay: The platters spin; wait for platters spin to the right place ↵

Transfer time: Reading/writing the data ↵

Key approaches for physical disk-based transfer

Block transfer

- 减少 transfer time

Cylinder-based

- Principle: store relation data on the same cylinder:
- 减少 seek time

Multiple disks — **RAID**

- 减少 transfer 和 seek time

Disk scheduling — Elevator approach

- 减少 seek time

Prefetching/double buffering

- 减少 seek delay, Rotary delay

Total size of required pages: $105,264 * 4 * 1024 = 431,161,344$ bytes

Actual table size: $2,000,000 * 200 = 400,000,000$ bytes

$$\text{Overhead} = \frac{(431,161,344 - 400,000,000)}{400,000,000} = 7.79\%$$

Therefore, the **total number of index pages** is

$$390 + 2 + 1 = \mathbf{393 \text{ index pages}}$$

$$393/140,351 = \mathbf{0.2\% \text{ increase}}$$

1b) Calculate the time taken to perform a table scan (i.e., linear scan) through the table Rel1.

Hint: How many pages are needed to make up the space occupied by the table? How many pages will be read from disk during a table scan?

To do a table scan we must fetch each of the 625 pages of the table; measured in seconds, this takes $625 * 150 = 93750 \text{ ms} = 93.75 \text{ s}$ or approximately 1.5 minutes.

Scan 的是 page, 不是 record

Binary search on sorted files:

- Height $\log 140,351$ + additional pages containing retrieved records

Nested loop join — $O(b_R + |R| * b_S)$

Block-nested loop join — $O(b_R + b_R * b_S)$

assume that the buffer can only hold two pages for the two relations

Indexed-nested loop join — $O(b_R + (|R| * c))$

Must satisfy

- Join is an equi-join or natural join, and
 - CROSS JOIN 不满足
- an index is available on the inner table's join attribute

$$\lceil \log_{(B-1)}(N/B) \rceil + 1 \leftarrow \lceil N/B \rceil \text{ CEIL}(N/(B-1))$$

Materialization

Def

- Output of one operator written to disk and the next operator will read it from the disk.
- starting at the lowest-level.
- Produces sorted output.

Advantage:

Can always apply materialized evaluation

Disadvantages:

Costs can be quite high

Pipelining

Def

- Output of one operator is directly input to next operator
- evaluate several operations **concurrently**
- **unsorted output**

Advantage

Much cheaper than materialization:

- no need to store a temporary table

Issues:

If algorithm requires sorted output, pipelining may not work well if data is not already sorted.

Student(snum, sname, major, level, age)
Class(name, meets_at, room, fid)
Enrolled(snum, cname)
Faculty(fid, fname, deptid)

Express each of the following integrity constraints in SQL

- > Every class has a maximum enrolment of 30 students.
- > Every faculty member must teach at least two courses.
- > No department can have more than 10 faculty members.

```
CREATE TABLE Enrolled (snum INTEGER, cname CHAR(20),  
PRIMARY KEY (snum, cname),  
FOREIGN KEY (snum) REFERENCES Student,  
FOREIGN KEY (cname) REFERENCES Class,  
CHECK (( SELECT COUNT (snum)  
FROM Enrolled  
GROUP BY cname) <= 30))
```

```
CREATE ASSERTION TeachConstraint  
CHECK ( ( SELECT COUNT (*)  
FROM Faculty F LEFT NATURAL JOIN Class C  
GROUP BY C.fid  
HAVING COUNT (*) < 2) = 0)
```

```
CREATE TABLE Faculty (fid INTEGER, fname CHAR(20), deptid INTEGER,  
PRIMARY KEY (fnum),  
CHECK ( ( SELECT COUNT (*)  
FROM Faculty F  
GROUP BY F.deptid  
HAVING COUNT (*) > 10) = 0))
```

涉及多个表
因为要包含所有相关表
即使它在right表没有对应的

Insert anomaly: We cannot insert a new ProvNo to describe a doctor's speciality unless we also include a VisitID (i.e., we can't record details about a doctor unless they've made a visit).

Delete anomaly: If we delete the only record of a visit made by a doctor, we lose the information about their speciality.

Update anomaly: If a patient's address changes, we must update this address for every row in which that patient appears.

Assume we have three relations: Student, UoS, and Takes: Which sentence would you use to write the sql assertion?

70 ✖



A

28 ✔



B

A. For each tuple in the Student relation, the value of the attribute tot_cred must equal the sum of credits of courses that the student has completed successfully.

B. There is no student that has a total credit that is different from the sum of credits of courses that they have completed successfully.

A. 针对每个学生记录，要求 `tot_cred` = 该学生完成课程学分的和

- 说的是单个学生元组的属性值必须满足条件。
- 就是***对每条数据***有条件。
- 换句话说，检查每一条学生数据是否满足规则。

B. 不存在学生的 `tot_cred` 和其完成课程学分之和不相等

- 说的是全数据库范围内，没有一条记录违反条件。
- 这是一个整体性、全局性的声明。
- 也就是说，在所有学生里，找不到不符合规则的。

Functional dependencies

8 ✖

define an N-M relationship between
two sets of attributes

59 ✔

define an N-1 relationship between
two sets of attributes

Multivalued dependencies introduce duplication

77 ✔

True

32 ✖

False

在没有 FD 的情况下判断 MVD

在一个关系中，当某个属性 A 决定了一组值 B，而这组值与另一个属性 C 无关时，我们说存在 $A \twoheadrightarrow B$ 的多值依赖。

Transaction logs are used for

94 ✓



14 ✗




auditing purposes

ensuring durability

Multivalued dependencies are needed to

67 ✗




to represent a relationship between any two set-valued attributes

41 ✓



establish independence between two relationships

If a schedule is serializable, it is then conflict serializable



Serializable 只看“最终结果”，即 interleaved 和 serial 的结果一样；
Conflict Serializable 既看“最终结果”，也看“操作冲突的顺序结构”。

Consider the following **interleaved** execution

T1: $A = A - 100$, $B = B + 100$
T2: $A = 1.05 * A$, $B = 1.05 * B$

$$A_F = (A_i - 100) * 1.05, B_F = B_i * 1.05 + 100$$

It is not serializable: the result of the above interleaved execution is *not the same* to *either* of the following *two serial* executions.


T1: $A = A - 100$, $B = B + 100$
T2: $A = 1.05 * A$, $B = 1.05 * B$


$$A_F = 1.05 * (A_i - 100), B_F = 1.05 * (B_i + 100)!$$

T1: $A = A - 100$, $B = B + 100$
T2: $A = 1.05 * A$, $B = 1.05 * B$

$$A_F = (1.05 * A_i) - 100!, B_F = (1.05 * B_i) + 100$$

Buffer management provides

38 ✓

Virtual memory for in memory operations

39 ✗

An efficient framework to distribute blocks on disks

Elevator disk scheduling works well when data requests

38 ✗

are few

43 ✓

are unpredictable

指的是磁盘请求的到达顺序、时间和位置都是不可预知、随机变化的情况。


The disk space manager's function is to

35 ✖



manage the use of buffers

37 ✔



determine how data is stored in pages

Heuristic-based optimization mainly deals with unary operations

44 ✔



True

25 ✖



False

Unary 指 selction, projection