

COMP9120

Week 4: Relational Algebra & SQL

Semester 1, 2025

Professor Athman Bouguettaya
School of Computer Science



Warming up



THE UNIVERSITY OF
SYDNEY



Acknowledgement of Country

I would like to acknowledge the Traditional Owners of Australia and recognise their continuing connection to land, water and culture. We are currently on the land of the Gadigal People of the Eora nation and pay our respects to their Elders, past, present and emerging.

COMMONWEALTH OF AUSTRALIA

Copyright Regulations 1969

WARNING

This material has been reproduced and communicated to you by or on behalf of the University of Sydney pursuant to Part VB of the Copyright Act 1968 (**the Act**).

The material in this communication may be subject to copyright under the Act. Any further copying or communication of this material by you may be the subject of copyright protection under the Act.

Do not remove this notice.



Schedule of Zoom Drop-in Helpdesk

- › Accessible from The Zoom link on Canvas.
- › Drop-in Helpdesk:
 - Week 5, 6 and 9, 10, 11



2025 Semester 1

- [Home](#)
- [Modules](#)
- [Assignments](#)
- [Marks](#)
- [Recorded Lectures](#)
- [People](#)
- [Ed Discussion](#)
- Zoom**
- [Quizzes](#)
- [Reading List](#)
- [Pages](#)
- [Smart Search](#)
- [Unit Outline](#)
- [Discussions](#)
- [Announcements](#)
- [Rubrics](#)
- [Collaborations](#)
- [Files](#)

zoom [Home](#) [Appointments](#)

Your current Time Zone and Language are (GMT+11:00) Canberra, Melbourne, Sydney, English 

[Upcoming Meetings](#) [Previous Meetings](#) [Personal Meeting Room](#) [Cloud Recordings](#)

Show my course meetings only

Start Time	Topic	Meeting ID
Wed, Mar 26 (Recurring) 8:00 PM	COMP9120 Helpdesk Host Abbey Lin	882 0767 6381 Passcode
Wed, Apr 2 (Recurring) 8:00 PM	COMP9120 Helpdesk Host Abbey Lin	882 0767 6381
Wed, Apr 30 (Recurring) 8:00 PM	COMP9120 Helpdesk Host Abbey Lin	882 0767 6381
Wed, May 7 (Recurring) 8:00 PM	COMP9120 Helpdesk Host Abbey Lin	882 0767 6381
Wed, May 14 (Recurring) 8:00 PM	COMP9120 Helpdesk Host Abbey Lin	882 0767 6381

REMINDER

WEEKLY LIVE HELP ON ED: *Started last week!*

To access the **weekly schedule**, please go to Canvas:

<https://canvas.sydney.edu.au/courses/63042/modules>

under Modules, under NEW! LIVE ED SCHEDULE S1
2025

How: The indicated tutor (e.g., Dipankar) will be on Ed during the hours mentioned in the schedule to answer your questions live!

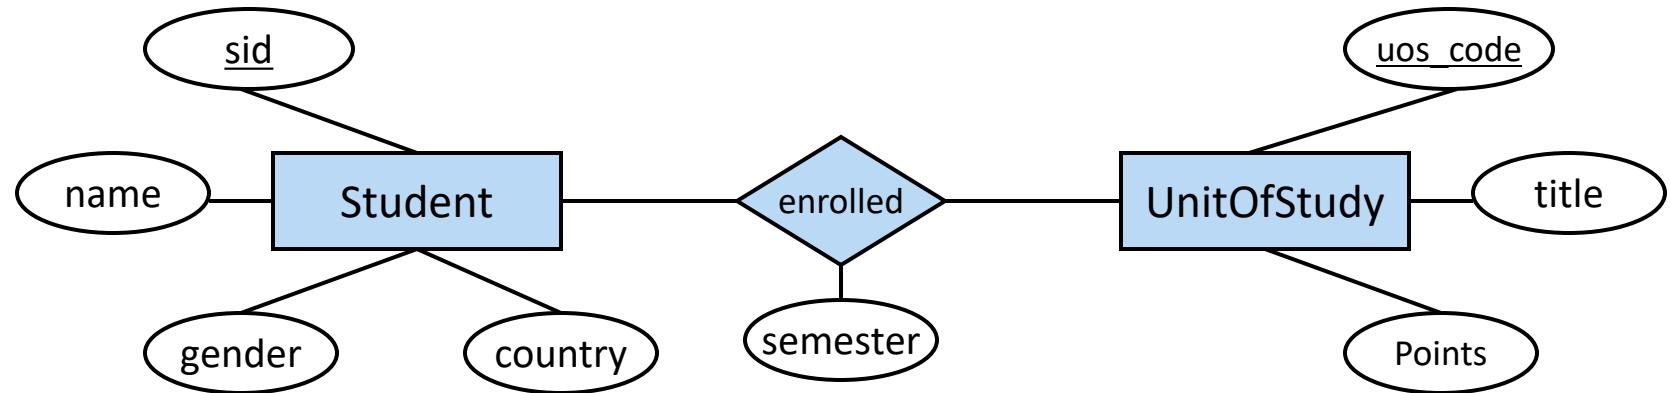
Week	Weekday	FROM	TO	Name
3	TUE	6:00 pm	8:00 pm	Dipankar
3	THUR	11:00 am	1:00 pm	Dipankar
4	TUE	6:00 pm	8:00 pm	Dipankar
4	THUR	11:00 am	1:00 pm	Dipankar
5	TUE	6:00 pm	8:00 pm	Dipankar
5	WED	6:00 pm	8:00 pm	Dipankar
5	THUR	11:00 am	1:00 pm	Dipankar
6	TUE	6:00 pm	8:00 pm	Dipankar
6	WED	6:00 pm	8:00 pm	Dipankar
6	THUR	11:00 am	1:00 pm	Dipankar
6	WED	6:00 pm	8:00 pm	Abbey
6	THUR	11:00 am	12:00 pm	Abbey
7	TUE	6:00 pm	8:00 pm	Dipankar
7	THUR	11:00 am	1:00 pm	Dipankar
8	TUE	6:00 pm	8:00 pm	Dipankar
8	THUR	11:00 am	1:00 pm	Dipankar
9	TUE	6:00 pm	8:00 pm	Dipankar
9	WED	6:00 pm	8:00 pm	Dipankar
9	THUR	11:00 am	1:00 pm	Dipankar
10	TUE	6:00 pm	8:00 pm	Dipankar
10	WED	6:00 pm	8:00 pm	Dipankar
10	THUR	11:00 am	1:00 pm	Dipankar
10	WED	6:00 pm	8:00 pm	Abbey
11	THUR	11:00 am	12:00 pm	Abbey
11	TUE	6:00 pm	8:00 pm	Dipankar
11	WED	6:00 pm	8:00 pm	Dipankar
11	THUR	11:00 am	1:00 pm	Dipankar
11	WED	6:00 pm	8:00 pm	Abbey
11	THUR	11:00 am	12:00 pm	Abbey
12	TUE	6:00 pm	8:00 pm	Dipankar
12	THUR	11:00 am	1:00 pm	Dipankar
13	TUE	6:00 pm	8:00 pm	Dipankar
13	WED	6:00 pm	8:00 pm	Dipankar
13	THUR	11:00 am	1:00 pm	Dipankar

› **Relational Algebra:** an algebra for relational model

- Six basic operators

› **Data Manipulation Language: Introduction to SQL**

- Basic SQL queries
- Join queries
- Set operations



Student			
<u>sid</u>	<u>name</u>	<u>gender</u>	<u>country</u>
1001	Ian	M	AUS
1002	Ha Tschi	F	ROK
1003	Grant	M	AUS
1004	Simon	M	GBR
1005	Jesse	F	CHN
1006	Franzisca	F	GER

Enrolled		
<u>sid</u>	<u>uos_code</u>	<u>semester</u>
1001	COMP5138	2023-S2
1002	COMP5702	2023-S2
1003	COMP5138	2023-S2
1006	COMP5318	2023-S2
1001	INFO6007	2023-S1
1003	ISYS3207	2023-S2

UnitOfStudy		
<u>uos_code</u>	<u>title</u>	<u>points</u>
COMP5138	Relational DBMS	6
COMP5318	Data Mining	6
INFO6007	IT Project Management	6
SOFT1002	Algorithms	12
ISYS3207	IS Project	4
COMP5702	Thesis	18

Relational Algebra



THE UNIVERSITY OF
SYDNEY

Exercise: Evaluating a Simple Query

Student			
sid	name	gender	country
1001	Ian	M	AUS
1002	Ha Tschi	F	ROK
1003	Grant	M	AUS
1004	Simon	M	GBR
1005	Jesse	F	CHN
1006	Franzisca	F	GER

Enrolled		
sid	uos_code	semester
1001	COMP5138	2023-S2
1002	COMP5702	2023-S2
1003	COMP5138	2023-S2
1006	COMP5318	2023-S2
1001	INFS6014	2023-S1
1003	ISYS3207	2023-S2

UnitOfStudy		
uos_code	title	points
COMP5138	Relational DBMS	6
COMP5318	Data Mining	6
INFO6007	IT Project Management	6
SOFT1002	Algorithms	12
ISYS3207	IS Project	4
COMP5702	Thesis	18

Using the above database instance, **find the titles of all units worth 6 credit points.**

Think about the steps we have to take to get the output.

title
Relational DBMS
Data Mining
IT Project Management

How does a RDBMS get the answer?

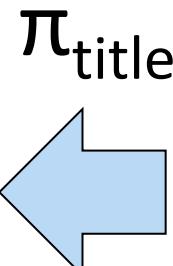
Find the titles of all units worth 6 credit points:

Two steps:

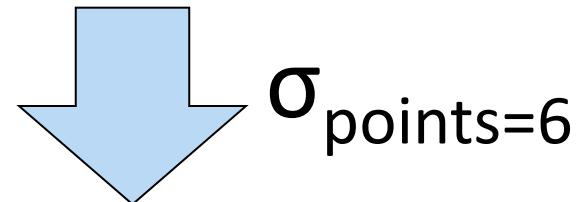
Relational Algebra expression:

$$\pi_{\text{title}}(\sigma_{\text{points}=6}(\text{UnitOfStudy}))$$

title
Relational DBMS
Data Mining
IT Project Management



uos_code	title	points
COMP5138	Relational DBMS	6
COMP5318	Data Mining	6
INFO6007	IT Project Management	6
SOFT1002	Algorithms	12
ISYS3207	IS Project	4
COMP5702	Thesis	18



uos_code	title	points
COMP5138	Relational DBMS	6
COMP5318	Data Mining	6
INFO6007	IT Project Management	6

Student			
sid	name	gender	country
1001	Ian	M	AUS
1002	Ha Tschi	F	ROK
1003	Grant	M	AUS
1004	Simon	M	GBR
1005	Jesse	F	CHN
1006	Franzisca	F	GER

Enrolled		
sid	uos_code	semester
1001	COMP5138	2023-S2
1002	COMP5702	2023-S2
1003	COMP5138	2023-S2
1006	COMP5318	2023-S2
1001	INFS6014	2023-S1
1003	ISYS3207	2023-S2

UnitOfStudy		
uos_code	title	points
COMP5138	Relational DBMS	6
COMP5318	Data Mining	6
INFO6007	IT Project Management	6
SOFT1002	Algorithms	12
ISYS3207	IS Project	4
COMP5702	Thesis	18

Find the id of all students who are enrolled in COMP5138

$\pi_{\text{sid}}(\sigma_{\text{uos_code}='COMP5138'}(\text{Enrolled}))$

Select sid
from Enrolled
where uos_code = = Comp 5138

sid
1001
1003

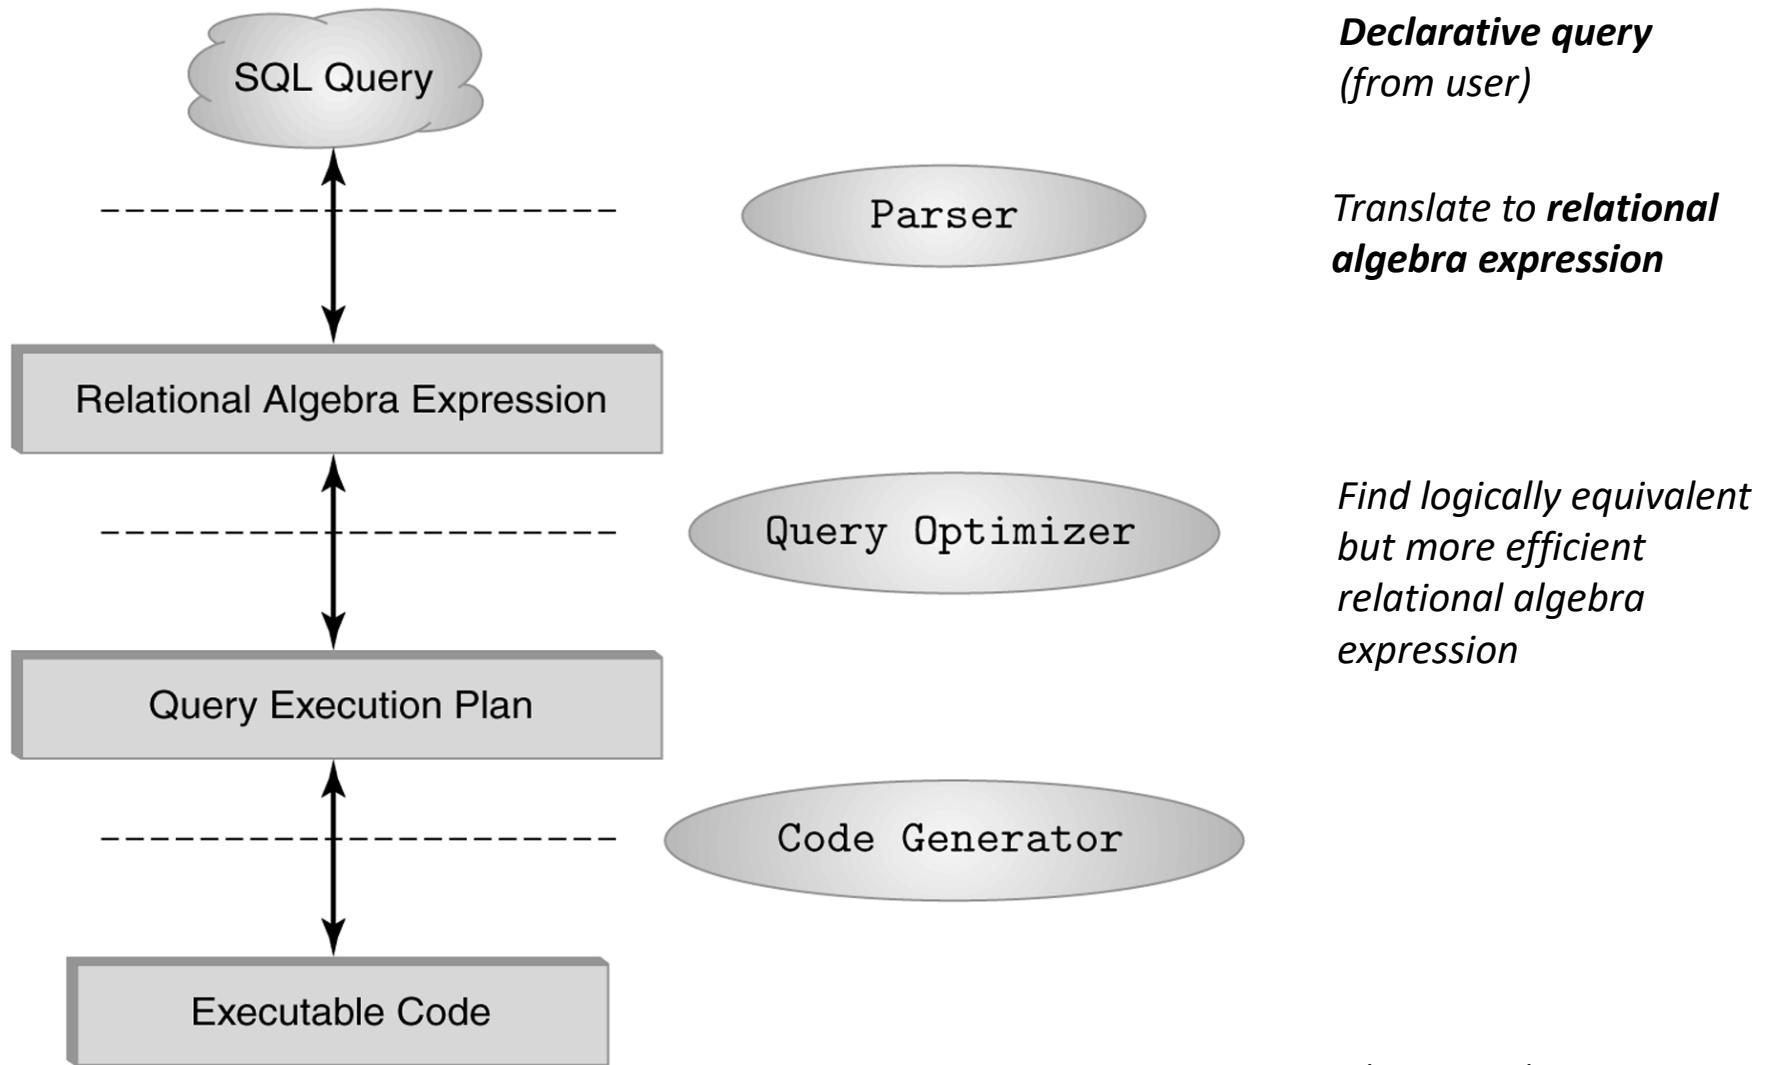
› Relational algebra (RA) is an algebra for the *relational model*

- Relational algebra operates on **sets**!
- It is composed of a set of **operators** (e.g., *selection, projection, join*)

$$\pi_{\text{title}}(\sigma_{\text{points}=6}(\text{UnitOfStudy}))$$

- It describes a step-by-step procedure (i.e., describes the **how**) for computing the desired answer - **imperative**
- **Contrast** with *tuple calculus* which specifies the *constraints* the *results* should satisfy (i.e., describes the **what**) - **declarative**

› RA allows us to translate *declarative* (SQL) queries into precise and optimizable expressions!



Unary operators

1. Operators that focus on one single relation
 - > - Selection (σ) selects a subset of rows from a relation.
 - > Operators can be chained together to form expressions that represent queries
2. A schema-level 'rename' operator $\pi_{\text{title}}(\sigma_{\text{points}=6}(\text{UnitOfStudy}))$
 - Rename (ρ) allows us to rename an attribute or a relation.

Select : σ Row
 Projection: π Column

Binary operators

1. Operators that matches tuples from two relations
 - Cross-product (\times) combines every tuple from two relations.
 - Join (\bowtie) combines matching tuples from two relations.
2. Set Operators: the relations must have the same number of attributes and be compatible
 - Union (\cup) of relations A and B returns all tuples in relation A plus those that are in relation B.
 - Intersection (\cap) of relations A and B returns all tuples in relation A that are matched in relation B (or vice versa).
 - Difference (-) of relations A and B returns tuples in relation A that have no match in relation B.

- ‘Extracts’ columns for attributes that are in the *projection list*.

- Removes columns that are *not* in the *projection list*, then ***eliminates duplicate*** rows
- Schema of the result contains *exactly* the *attributes* in the *projection list*.

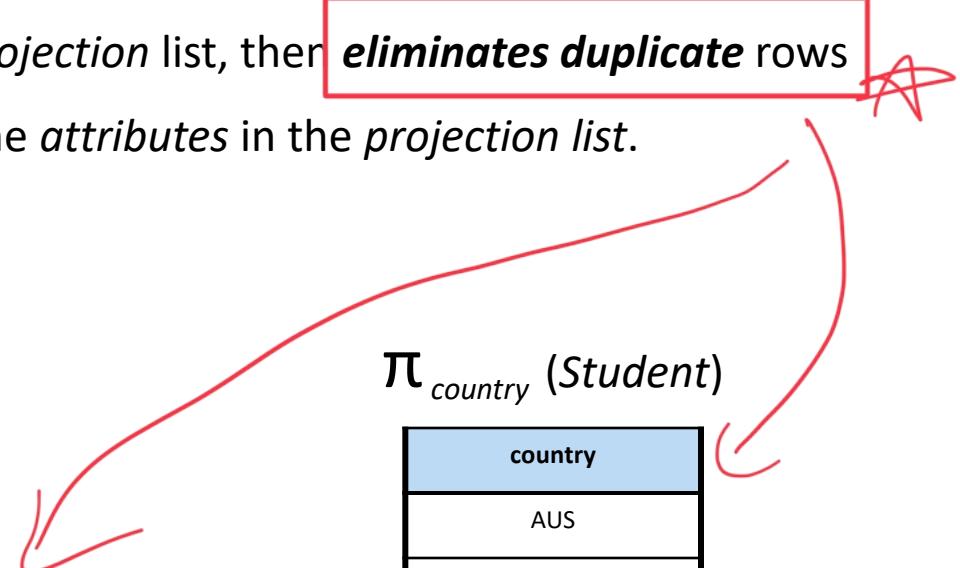
- Examples:

$\pi_{name, country} (Student)$

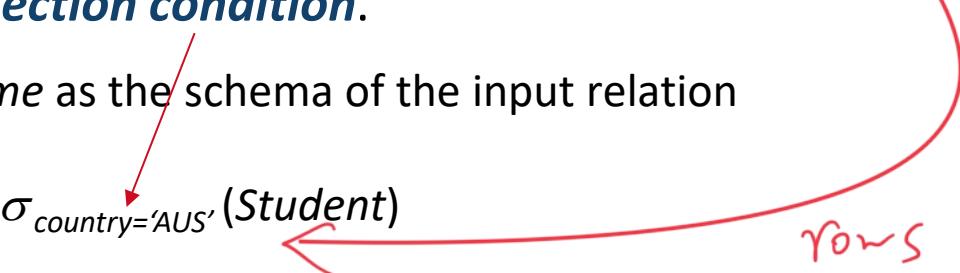
name	country
Ian	AUS
Ha Tschi	ROK
Grant	AUS
Simon	GBR
Jesse	CHN
Franzisca	GER

$\pi_{country} (Student)$

country
AUS
ROK
GBR
CHN
GER



- >Selects rows that satisfy a ***selection condition***.
 - Schema of the result is the *same* as the schema of the input relation
 - Example:

$$\sigma_{country='AUS'}(Student)$$


sid	name	gender	country
1001	Ian	M	AUS
1003	Grant	M	AUS

- Result relation can be the input for another relational algebra operation!
 - (called ***Operator composition***)
 - Example:

$$\pi_{name}(\sigma_{country='AUS'}(Student))$$

name
Ian
Grant

Student			
<u>sid</u>	<u>name</u>	<u>gender</u>	<u>country</u>
1001	Ian	M	AUS
1002	Ha Tschi	F	ROK
1003	Grant	M	AUS
1004	Simon	M	GBR
1005	Jesse	F	CHN
1006	Franziska	F	GER

- Will the following RA expression work to “Find the name of the students who live in Australia”?

$\sigma_{country='AUS'}(\pi_{name}(Student))$ → without ‘Country’ col
 (2)

$\pi_{name}(\delta_{country = 'AUS'})$

name
Ian
Ha Tschi
Grant
Simon
Jesse
Franziska

› Selection condition is a Boolean combination of terms

- Each term has the form: *attribute op constant*, or *attribute1 op attribute2*
- op can be $<$, $>$, \leq , \geq , \neq , $=$ *这些不是 SQL 的 op, 是 RA 的表达式*
- Terms are connected by *logical connectives*:
 - \wedge - means AND
 - \vee - means OR

Find out the male Australian students:

$$\sigma_{\text{gender}='M' \wedge \text{country}='AUS'} (\text{Student})$$

Student			
<u>sid</u>	name	gender	country
1001	Ian	M	AUS
1002	Ha Tschi	F	ROK
1003	Grant	M	AUS
1004	Simon	M	GBR
1005	Jesse	F	CHN
1006	Franzisca	F	GER

<u>sid</u>	name	gender	country
1001	Ian	M	AUS
1003	Grant	M	AUS

- › Sometimes also called ***Cartesian product***
- › Defined as: $R \times S = \{t s \mid t \in R \wedge s \in S\}$
 - Each tuple of R is paired with each tuple of S .
 - Resulting schema is the *concatenation* of the fields of R and S .
 - We might end up in a conflict with two fields having the same name -> use the *rename operation*
- If R or S is empty, then $R \times S$ is also empty.
- › Example:

<i>R</i>	
<i>A</i>	<i>B</i>
α	1
β	2

<i>S</i>		
<i>C</i>	<i>D</i>	<i>E</i>
α	10	a
β	10	a
β	20	b
γ	10	b

x

<i>A</i>	<i>B</i>	<i>C</i>	<i>D</i>	<i>E</i>
α	1	α	10	a
α	1	β	10	a
α	1	β	20	b
α	1	γ	10	b
β	2	α	10	a
β	2	β	10	a
β	2	β	20	b
β	2	γ	10	b

- Theta (Conditional) Join: $R \bowtie_{\theta} S = \sigma_{\theta}(R \times S)$

Example:

Student \bowtie

Lecturer

$Student.f_name = Lecturer.last_name \wedge Student.sid < Lecturer.empid$

\hookrightarrow Cartesian with selection

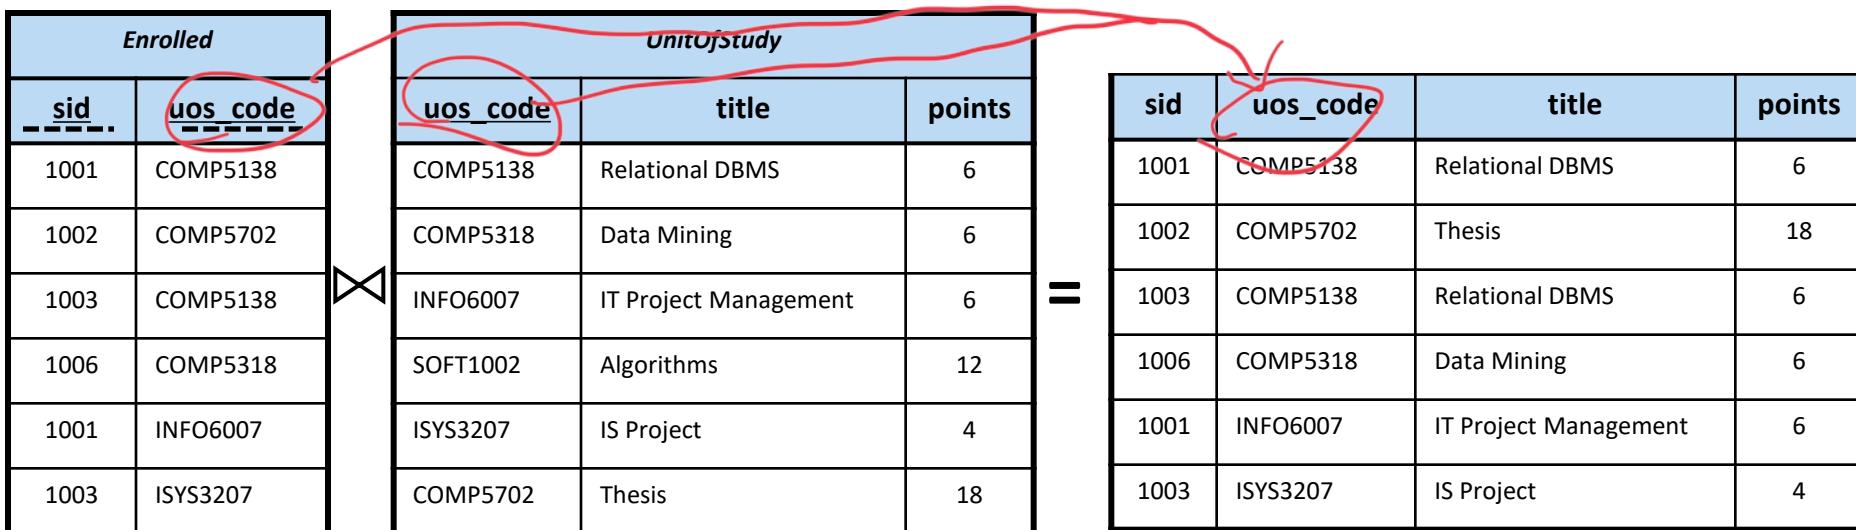
sid	given	f_name	gender	country	empid	lecturer_name	last_name	room
1001	Ian	Chung	M	AUS	47112344	Vera	Chung	321
1004	Simon	Poon	M	GBR	12345678	Simon	Poon	431
1004	Simon	Poon	M	GBR	99004400	Josiah	Poon	482
...
...
...	

- Resulting schema is the cross-product schema.
- Equi-Join: Special case of theta join where the condition θ contains only equalities.

= ↗

Natural Join: $R \bowtie S$

- *Equivalent to Equi-join on all common attributes (i.e., same named attributes), followed by a projection*
- Resulting schema is similar to equi-join, however the difference is that we retain only one common field for which equality is specified.



Enrolled

sid	uos_code
1001	COMP5138
1002	COMP5702
1003	COMP5138
1006	COMP5318
1001	INFO6007
1003	ISYS3207

UnitofStudy

uos_code	title	points
COMP5138	Relational DBMS	6
COMP5318	Data Mining	6
INFO6007	IT Project Management	6
SOFT1002	Algorithms	12
ISYS3207	IS Project	4
COMP5702	Thesis	18

=

sid	uos_code	title	points
1001	COMP5138	Relational DBMS	6
1002	COMP5702	Thesis	18
1003	COMP5138	Relational DBMS	6
1006	COMP5318	Data Mining	6
1001	INFO6007	IT Project Management	6
1003	ISYS3207	IS Project	4

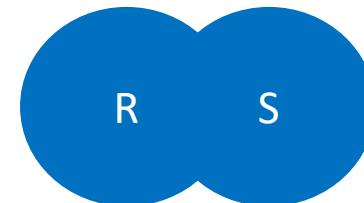
~~Join~~ will have 5 columns
Then

- The **set operations** take two input relations R and S of the *same arity* (i.e., number of attributes) *and same attribute domains* (called *union-compatible*)

- Set Union $R \cup S$

or

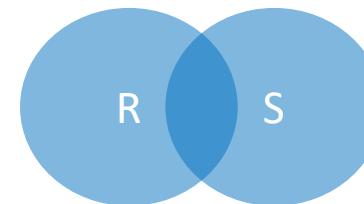
- Definition: $R \cup S = \{ t \mid t \in R \vee t \in S \}$



- Set Intersection $R \cap S$

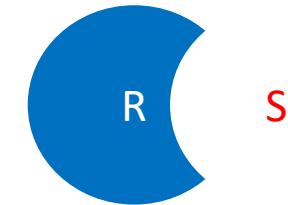
and

- Definition: $R \cap S = \{ t \mid t \in R \wedge t \in S \}$



- Set Difference $R - S$

- Definition: $R - S = \{ t \mid t \in R \wedge t \notin S \}$



- Important constraint:** R and S must have *compatible schema*

- R, S have the *same arity* (same number of fields)

- 'Corresponding' fields must have the same domains

*Even
Domain should be
same*

Exercise: Set Operations

- Consider the following relations: Use a set operation in an RA expression to return the *id* of all students who are not postgraduates.

Student			
<u>sid</u>	name	gender	country
1001	Ian	M	AUS
1002	Ha Tschi	F	ROK
1003	Grant	M	AUS
1004	Simon	M	GBR
1005	Jesse	F	CHN
1006	Franzisca	F	GER

their schema is different
~~Student – Postgraduate ?~~

Postgraduate
<u>sid</u>
1003
1004
1005

Now they have same attributes

$\pi_{\text{sid}}(\text{Student}) - \text{Postgraduate}$

<u>sid</u>
1001
1002
1006

Set difference

- Allows us to rename a relation and its attributes:

- Notation 1:

 $\rho_x(E)$ (pronounced Ro)

only RO

Change relationship E to x.

- returns the relation E under the name X

table name

- Notation 2:

 $\rho_x(A_1, A_2, \dots, A_n)(E)$

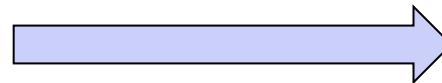
- (assumes that the relational-algebra expression E has arity n)
- returns the result of relation E under the name X , and with the attributes renamed to A_1, A_2, \dots, A_n .

- Note that rename *only modifies* the schema of a relation!

Example: Rename Operation

$$\rho_{UOS(ucode, title, credits)} (UnitOfStudy)$$

UnitOfStudy		
<u>uos_code</u>	title	points
COMP5138	Relational DBMS	6
COMP5318	Data Mining	6
INFO6007	IT Project Management	6
SOFT1002	Algorithms	12
ISYS3207	IS Project	4
COMP5702	Thesis	18

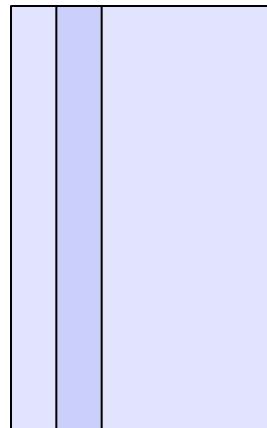


UOS		
<u>ucode</u>	title	credits
COMP5138	Relational DBMS	6
COMP5318	Data Mining	6
INFO6007	IT Project Management	6
SOFT1002	Algorithms	12
ISYS3207	IS Project	4
COMP5702	Thesis	18

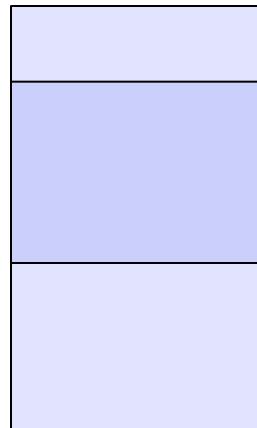


Visualisation of Relational Algebra Operators

Projection (π)

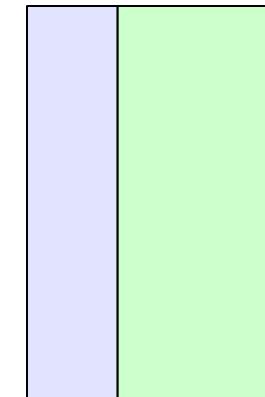


Selection (σ)



Cross-product (\times)

$$\begin{array}{c} \times \\ \text{---} \end{array}$$



Join (\bowtie)

$$\begin{array}{c} \bowtie \\ \text{---} \\ = \\ \text{---} \end{array}$$

Set Union (\cup)

$$\begin{array}{c} \cup \\ \text{---} \\ = \\ \text{---} \end{array}$$

Set Minus ($-$)

$$\begin{array}{c} - \\ \text{---} \\ = \\ \text{---} \end{array}$$

Set Intersection (\cap)

$$\begin{array}{c} \cap \\ \text{---} \\ = \\ \text{---} \end{array}$$

- › We can distinguish between *basic* and *derived* RA operators
- › **Only 6 basic operators** are required to **express everything else**:
 - **Selection** (σ) selects a subset of rows from relation.
 - **Projection** (π) removes unwanted columns from relation.
 - **Cross-product** (\times) allows us to fully combine two relations.
 - **Union** (\cup) returns tuples that are in relation 1 or relation 2.
 - **Set Difference** ($-$) returns tuples that are in relation 1, but not in relation 2.
 - **Rename** (ρ) allows us to rename one field/table name to another name.

- › Additional (derived) operations:
 - E.g.: *intersection, join*. [Not essential, but very useful]
 - Equivalence: Intersection: $R \cap S = R - (R - S)$

Join: $R \bowtie_{\Theta} S = \sigma_{\Theta}(R \times S)$

Intersection of

$$R = \{1, 2, \cancel{3}\}$$

$$S = \{\cancel{3}, 4\} = \{3\}$$

Difference

$$R - S = \{1, 2\}$$

$$R - (R - S) = \{3\}$$

$(1, 2, \cancel{3}) - (1, 2) = \{3\}$

- › Different relational algebra expressions *can be equivalent but have different execution costs* ★
- › Example: List the names of all students enrolled in ‘Relational DBMS’
 - One way is → 不 filter N 条
 - $\pi_{\text{name}} (\sigma_{\text{title}=\text{'Relational DBMS'}} ((\text{Student} \bowtie \text{Enrolled}) \bowtie \text{UnitOfStudy}))$
 - Another (*more efficient*) way is: → 不 filter 一条
 - $\pi_{\text{name}} (\text{Student} \bowtie (\text{Enrolled} \bowtie (\sigma_{\text{title}=\text{'Relational DBMS'}} (\text{UnitOfStudy}))))$

Student			
sid	name	gender	country
1001	Ian	M	AUS
1002	Ha Tschi	F	ROK
1003	Grant	M	AUS
1004	Simon	M	GBR
1005	Jesse	F	CHN
1006	Franzisca	F	GER

Enrolled		
sid	uos_code	semester
1001	COMP5138	2023-S2
1002	COMP5702	2023-S2
1003	COMP5138	2023-S2
1006	COMP5318	2023-S2
1001	INFO6007	2023-S1
1003	ISYS3207	2023-S2

UnitOfStudy		
uos_code	title	points
COMP5138	Relational DBMS	6
COMP5318	Data Mining	6
INFO6007	IT Project Management	6
SOFT1002	Algorithms	12
ISYS3207	IS Project	4
COMP5702	Thesis	18

Let's take a break!

and

have some fun!



THE UNIVERSITY OF
SYDNEY

Introduction to SQL



THE UNIVERSITY OF
SYDNEY



- › **Basic SQL Queries**

- › **Join Queries**

- › **Set Operations**

- › Relational algebra is a *theory-based* **procedural** query language
 - May be challenging to non-experts to understand and use
- › Reminder: SQL is a standard **high-level declarative** query language for RDBMS
 - Describing *what* data we are interested in, but *not how* to retrieve it.
 - Based on SEQUEL, introduced in the mid-1970's as the query language for IBM's System (Structured English Query Language)
- › *RDBMS internally maps SQL to equivalent relational algebra expressions.*
- › Many standards out there
 - ANSI SQL, SQL92 (a.k.a. SQL2), SQL99 (a.k.a. SQL3), ...
 - RDBMS vendors support various subsets

› DDL (Data Definition Language)

- Create, drop, or alter the relation schema

› DML (Data Manipulation Language)

- The insertion of new information into the database

- The deletion of information from the database

- The modification of information stored in the database

- The retrieval of information stored in the database

- A **Query** is a statement requesting the retrieval of information
- The portion of a DML that involves information retrieval is called a **query language**

› DCL (Data Control Language)

- Commands that control a database, including administering access privileges and users

- › Used for queries on single or multiple tables
- › keywords:
 - **SELECT** Lists the columns (and expressions) that should be returned from the query
 - **FROM** Indicate the table(s) from which data will be obtained
 - **WHERE** Indicate the conditions to include a tuple in the result
 - **GROUP BY** Indicate the categorization of tuples
 - **HAVING** Indicate the conditions to include a category
 - **ORDER BY** Sorts the result according to specified criteria
- › Note: the result of an SQL query is also a table / relation
 - The result table can contain duplicate rows

- › List the names of all students.

```
SELECT name
FROM Student
```

- › List the names of all Australian students.

```
SELECT name
FROM Student
WHERE country='AUS'
```

- › * in select denotes “all attributes”.

```
SELECT *
FROM Student
```

- › General form of SFW query:

```
SELECT <attributes>
FROM <one or more tables>
WHERE <conditions>
```

$$\pi_{name} (\sigma_{country='AUS'} (Student))$$

A Select-From-Where query is equivalent to the relational algebra expression:

$$\pi_{A_1, A_2, \dots, A_n} (\sigma_{condition} (R_1 \times R_2 \times \dots \times R_m))$$

- **SELECT** corresponds to **projection** (π) in RA
- **FROM** corresponds to **Cartesian product** (\times) in RA
- **WHERE** corresponds to **selection** (σ) in RA

› SQL commands are not case sensitive:

- **SELECT, Select, select**, all okay to use
- Same: Student, student

› Use single quotes for string constants:

- ‘aus’ – yes
- “aus” - no

› String constants are case sensitive:

- **Different:** ‘AUS’, ‘aus’, ‘Aus’

Remove Duplicates in *Select* Clause

- › RDBMS allows duplicates in tables and query results.
 - Query result preserves duplicates by default.
 - Example:
- › To eliminate the duplicates, use the keyword **DISTINCT** after **SELECT**.
 - Example: List the distinct countries where students come from.

```
SELECT country  
      FROM Student
```

country
AUS
ROK
AUS
GBR
CHN
GER

**SELECT DISTINCT country
 FROM Student**

country
AUS
ROK
GBR
CHN
GER

Arithmetic Expressions in *Select* Clause

- › The **SELECT** clause can contain arithmetic expressions, involving the operations $+$, $-$, $*$ and $/$, and operating on constants or attributes of tuples.
- › The query:

```
SELECT uos_code, title, points*2  
FROM UnitOfStudy
```



would return a table which is the same as the UnitofStudy table except that the credit-point-values are doubled.

- › The **WHERE** clause specifies conditions that the result must satisfy
- › Comparison operators in SQL: = , > , >= , < , <= , **<>** (or !=) 
- › Comparison results can be combined using the logical connectives **AND**, **OR**, and **NOT**.
- › Comparisons can be applied to results of arithmetic expressions
- › Example: Find all UoS codes for units taken by student 1001 in 2024-S2:

```
SELECT uos_code
FROM Enrolled
WHERE sid = 1001 AND Semester = '2024-S2'
```

- › SQL includes a string-matching operator for comparisons on character strings.

- **LIKE** is used for *string matching*
- List the titles of all “COMP” unit of studies.

```
SELECT title  
FROM UnitOfStudy  
WHERE uos_code LIKE 'COMP%'
```

format string

- › Patterns are described using two special characters (“wildcards”):

- percent (%). The % character *matches any substring*.
- underscore (_). The _ character *matches any single character*.

- › SQL supports a variety of string operations such as

- concatenation (using “||”)
- converting from upper to lower case (and vice versa)
- finding string length, extracting substrings, etc.

- › SQL allows renaming table attributes using the **AS** clause:

old_name **AS** new_name

- › This is very useful: e.g., give result *columns of expressions* a meaningful name.
- › We can also assign names (called **aliases**) to tables as shown in the following example:
- › Example:

- Find the uos_code, credit_points for COMP5318,
and rename the column name uos_code as course_code.

```
SELECT a.uos_code AS course_code, a.credit_points  
  FROM unitofstudy a  
 WHERE a.uos_code = 'COMP5318'
```

- › List all students (i.e., their names) from Australia in alphabetical order.

```
SELECT name  
FROM Student  
WHERE country='AUS'  
ORDER BY name
```

- › Two options (per attribute):
 - **ASC** ascending order (default)
 - **DESC** descending order
- › You can order by more than one attribute
 - e.g., **ORDER BY** country **DESC**, name **ASC**



› Basic SQL Queries

› Join Queries

› Set Operations

- › An *implicit join* query combines two or more tables into a single table

- The **FROM** clause lists the tables involved in the query
 - corresponds to the Cartesian product of the tables.
 - *join-predicates explicitly stated in the WHERE clause*

- › Examples:

- Find the *Cartesian product* *Student x UnitOfStudy*

```
SELECT *
  FROM Student, UnitofStudy
```

- Find the student ID, name, and gender of all students enrolled in INFO6007:

```
SELECT sid, name, gender
  FROM Student, Enrolled
 WHERE Student.sid = Enrolled.sid AND
       uos_code = 'INFO6007'
```

- Which students did enroll in what semester?

```
SELECT S.sid, S.name, E.semester
FROM Student S, Enrolled E
WHERE S.sid = E.sid
```

FROM clause: This is an implicit Join which involves multiple tables

WHERE clause performs the equality check for common columns of the two tables

$$\pi_{S.sid, S.name, E.semester} (\sigma_{S.sid = E.sid} (\rho_S (Student) \times \rho_E (Enrolled)))$$

Why do we need aliases?

- › Some queries need to refer to the same table twice
- › In this case, aliases are given to the same table name
 - Example: For each academic staff, retrieve their name, and the name of his/her immediate supervisor.

SELECT	<i>L.name, M.name</i>
FROM	Lecturer L, Lecturer M
WHERE	<i>L.manager = M.empid</i>

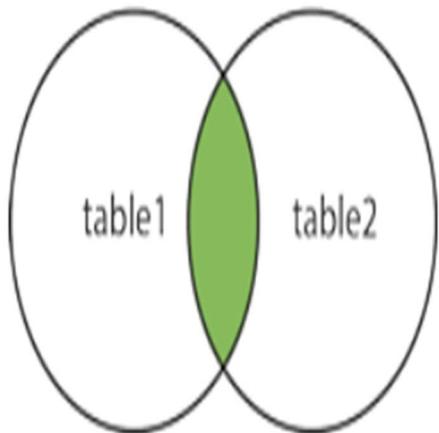
- We can think of L and M as two different copies of the table Lecturer; L represents lecturers in role of supervisees and M represents lecturers in role of supervisors (managers)



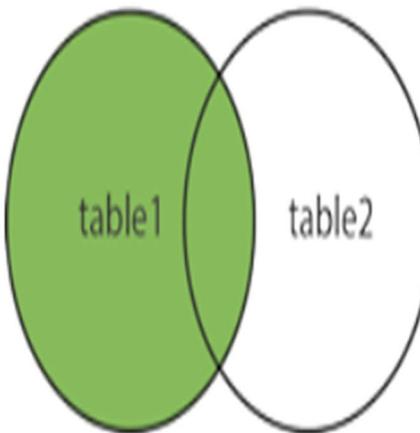
- Theta join
- Equi - join
- Natural join

Using Joins Explicitly

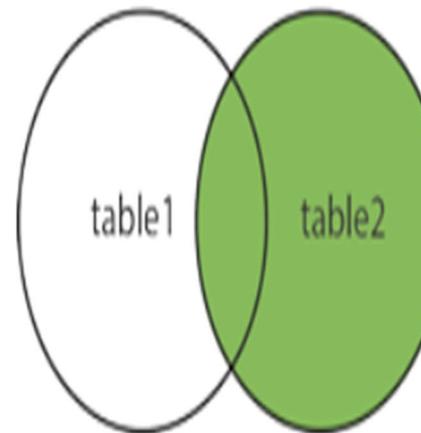
INNER JOIN



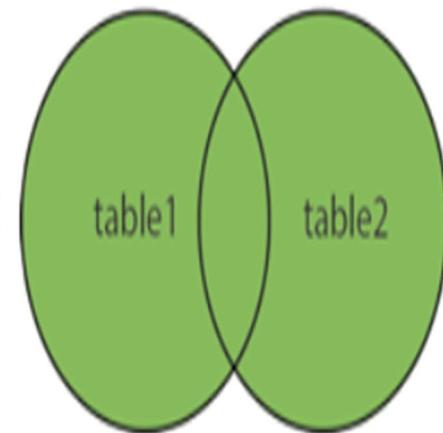
LEFT JOIN



RIGHT JOIN



FULL OUTER JOIN



- **Outer join** – a join in which rows that do not have matching values are also included (exactly once) in the result relation
- **Left outer join**: includes rows that would be found in an inner join as well as rows from the left table that don't have matches (padded with NULL values)
- **Right outer join**: includes rows that would be found in an inner join as well as rows from the right table that don't have matches (padded with NULL values)
- **Full outer join**: includes rows that would be found in an inner join as well as rows from the left table and rows from the right table that don't have matches

- › Join operators: **(a)** are specified in the **FROM** clause, and **(b)** must have both a **join type** and a **join condition**
 - Available join types: **JOIN, LEFT OUTER JOIN, RIGHT OUTER JOIN, FULL OUTER JOIN**
 - Available join conditions: **NATURAL, ON <join condition>, USING (<list of attributes>)**

e.g.: **SELECT * FROM Student JOIN Enrolled USING (sid)**

<u>sid</u>	name	gender	country	<u>uos_code</u>	semester
1001	<i>Ian</i>	<i>M</i>	<i>AUS</i>	<i>COMP5138</i>	<i>2023-S2</i>
1002	<i>Ha Tschi</i>	<i>F</i>	<i>ROK</i>	<i>COMP5702</i>	<i>2023-S2</i>
...

e.g.: **SELECT * FROM Student NATURAL LEFT OUTER JOIN Enrolled**

<u>sid</u>	name	gender	country	<u>uos_code</u>	semester
1001	<i>Ian</i>	<i>M</i>	<i>AUS</i>	<i>COMP5138</i>	<i>2023-S2</i>
1002	<i>Ha Tschi</i>	<i>F</i>	<i>ROK</i>	<i>COMP5702</i>	<i>2023-S2</i>
...
1005	<i>Jesse</i>	<i>F</i>	<i>CHN</i>	<i>Null</i>	<i>Null</i>



› Basic SQL Queries

› Join Queries

› Set Operations

- › The set operations **UNION**, **INTERSECT**, and **EXCEPT** operate on tables and correspond to the relational algebra operations \cup , \cap , $-$.
- › Each of the above operations **automatically eliminates** duplicates. 
- *First eliminate duplicates from the input tables, and then do the set operation*
 - Example: Suppose a tuple occurs 3 times in R and 1 time in S , then it occurs 0 times in $R \text{ EXCEPT } S$

理由：因为会去重，所以实际是 1-1，不是 3-1
- › To *retain all duplicates*, use the corresponding multiset versions **UNION ALL**, **INTERSECT ALL** and **EXCEPT ALL**.
- › Example: Suppose a tuple occurs m times in R and n times in S , then it occurs:

- $m + n$ times in $R \text{ UNION ALL } S$
- $\min(m, n)$ times in $R \text{ INTERSECT ALL } S$
- $\max(0, m - n)$ times in $R \text{ EXCEPT ALL } S$

$$R = \{a, b, a, c, d, a\}$$

$$S = \{a, e, f, a\}$$

Example: Set Operations

- › Find all customer names that have a loan, an account, or both:

SELECT customer_name **FROM** depositor
UNION
SELECT customer_name **FROM** borrower

Depositor(customer_name, account_balance)
Borrower(customer_name, loan_amount)

- › Find all customer names that have both a loan and an account

SELECT customer_name **FROM** depositor
INTERSECT
SELECT customer_name **FROM** borrower

- › Find all customer names that have an account but no loan

SELECT customer_name **FROM** depositor
EXCEPT
SELECT customer_name **FROM** borrower

Example: Set Operations*

- › Find students who enrolled in either ‘COMP5138’ or ‘ISYS3207’.

```
SELECT sid FROM Enrolled WHERE uos_code='COMP5138'  
UNION  
SELECT sid FROM Enrolled WHERE uos_code='ISYS3207'
```

- This is **equivalent** to the following SQL command without set operation

```
SELECT sid  
FROM Enrolled  
WHERE uos_code='COMP5138' OR uos_code='ISYS3207'
```

- › Find students who enrolled in both ‘COMP5138’ and ‘ISYS3207’.

```
SELECT sid FROM Enrolled WHERE uos_code='COMP5138'  
INTERSECT  
SELECT sid FROM Enrolled WHERE uos_code='ISYS3207'
```

- This is **not** equivalent to

```
SELECT sid  
FROM Enrolled  
WHERE uos_code='COMP5138' AND uos_code='ISYS3207'
```

不相等

Impossible, since can have
only one value



You should now be able to...

› ...**formulate basic SQL Queries**

- Select-From-Where Query
- Join queries
- Set operations

› ...know how SQL relates to the relational algebra

- Write equivalent relational algebra expressions for SQL queries

- › Ramakrishnan/Gehrke (3rd edition - the ‘Cow’ book (2003))
 - **Chapter 5**
uses the famous ‘Sailor-database’ as examples
- › Kifer/Bernstein/Lewis (2nd edition – 2006)
 - Chapter 5
includes some helpful visualisations on how complex SQL is evaluated
- › Ullman/Widom (3rd edition – 2008)
 - Chapter 6
up-to 6.5 good introduction and overview of all parts of SQL querying
- › Silberschatz/Korth/Sudarshan (5th edition - ‘sailing boat’)
 - Sections 3.1-3.6
- › Elmasri/Navathe (5th edition)
 - Sections 8.4 and 8.5.1

- › Integrity Constraints
 - Domain and CHECK constraints
 - ON DELETE and ON UPDATE actions, deferred constraints
 - Assertions
- › Readings :
 - Ramakrishnan/Gehrke (3rd edition - the 'Cow' book)
 - **Sections 3.2-3.3 and Sections 5.7-5.9**
 - *Integrity constraints are covered in different parts of the SQL discussion;*
 - Kifer/Bernstein/Lewis (2nd edition)
 - Sections 3.2.2-3.3 and Chapter 7
 - *Integrity constraints are covered as part of the relational model, but a good dedicated chapter (Chap 7) on triggers*
 - Ullman/Widom (3rd edition)
 - Chapter 7
 - *Has a complete chapter dedicated to both integrity constraints&triggers.*

See you next week!



THE UNIVERSITY OF
SYDNEY