

COMP9120

Week 11: Storage & Indexing

Semester 1, 2025

Professor Athman Bouguettaya
School of Computer Science



Warming up



THE UNIVERSITY OF
SYDNEY



Acknowledgement of Country

I would like to acknowledge the Traditional Owners of Australia and recognise their continuing connection to land, water and culture. I am currently on the land of the Gadigal people of the Eora nation and pay my respects to their Elders, past, present and emerging.

COMMONWEALTH OF AUSTRALIA

Copyright Regulations 1969

WARNING

This material has been reproduced and communicated to you by or on behalf of the University of Sydney pursuant to Part VB of the Copyright Act 1968 (the Act).

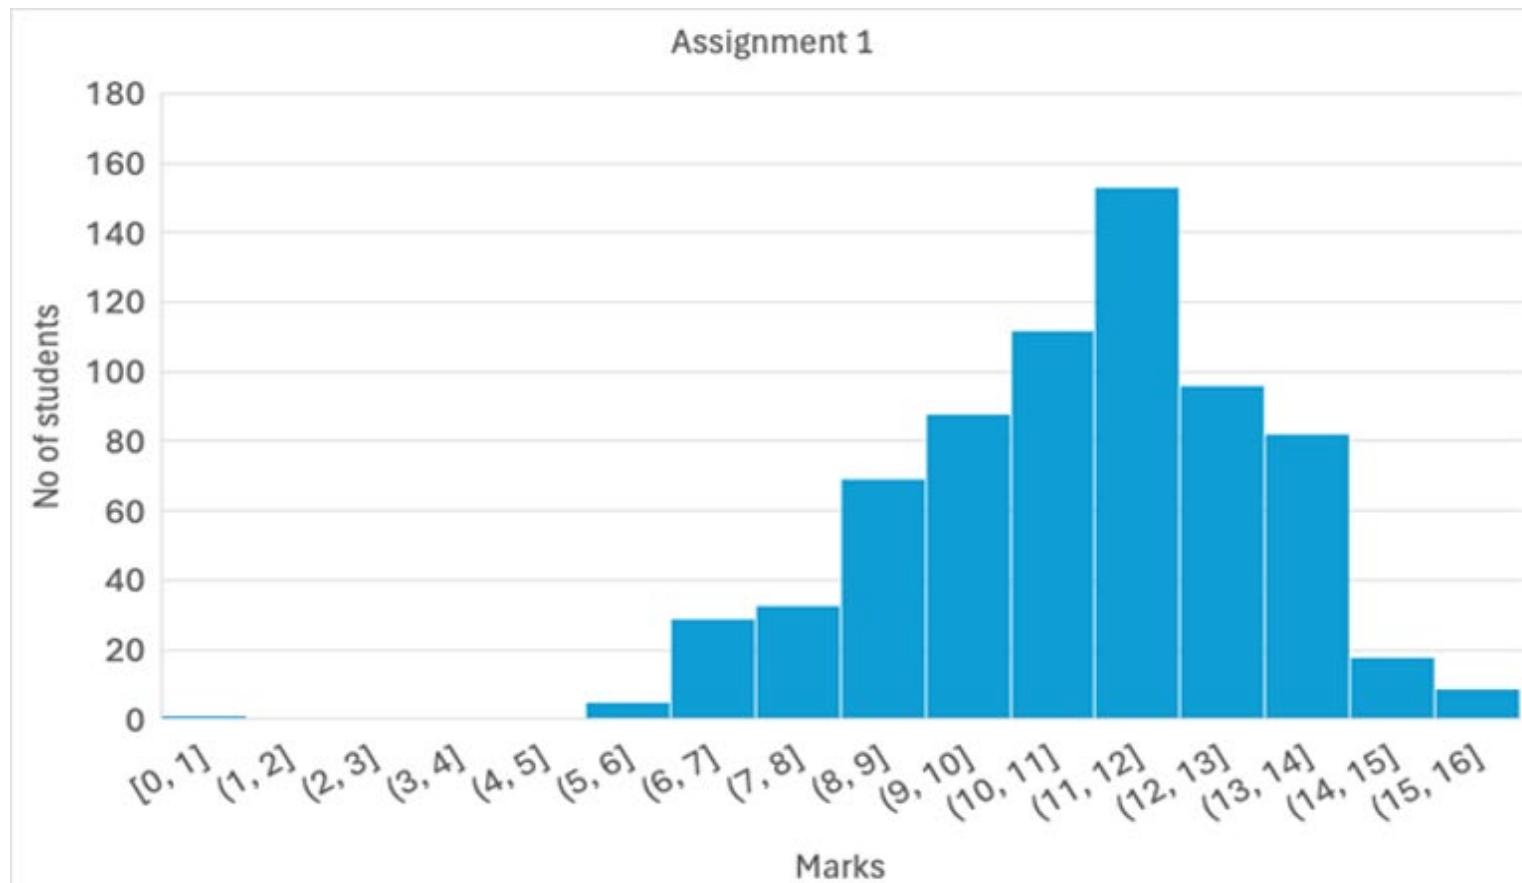
The material in this communication may be subject to copyright under the Act. Any further copying or communication of this material by you may be the subject of copyright protection under the Act.

Do not remove this notice.



Average: 69%

Median : 71.88%



› **Physical Data Organization: how is data physically stored in DBMS?**

› **Data Access Paths: how are records retrieved from DBMS?**

- Access methods for heap files (linear scan)
- Access methods for sorted files (binary search)
- Access methods for indexes (index scan)

› **B+ Tree Index**

- Primary index
- Composite search keys

› **Physical Data Organization: how is data physically stored in DBMS?**

› **Data Access Paths: how are records retrieved from DBMS?**

- Access methods for heap files (linear scan)
- Access methods for sorted files (binary search)
- Access methods for indexes (index scan)

› **B+ Tree Index**

- Primary index
- Composite search keys

How is a relation physically stored and accessed?

- Consider the table **Student**

Attributes (also called **columns, fields**)

Student				
sid	name	login	gender	address

Tuples

(also called **rows, records**)

5312666	Jones	ajon1121@cs	m	123 Main St
5366668	Smith	smith@mail	m	45 George
5309650	Jin	ojin4536@it	f	19 City Rd

- How are relational tables **physically** stored?
 - Several *organization approaches* for storing data
- How is data **accessed** and what is the **associated cost**?
 - Access strategies* and metrics to execute the query

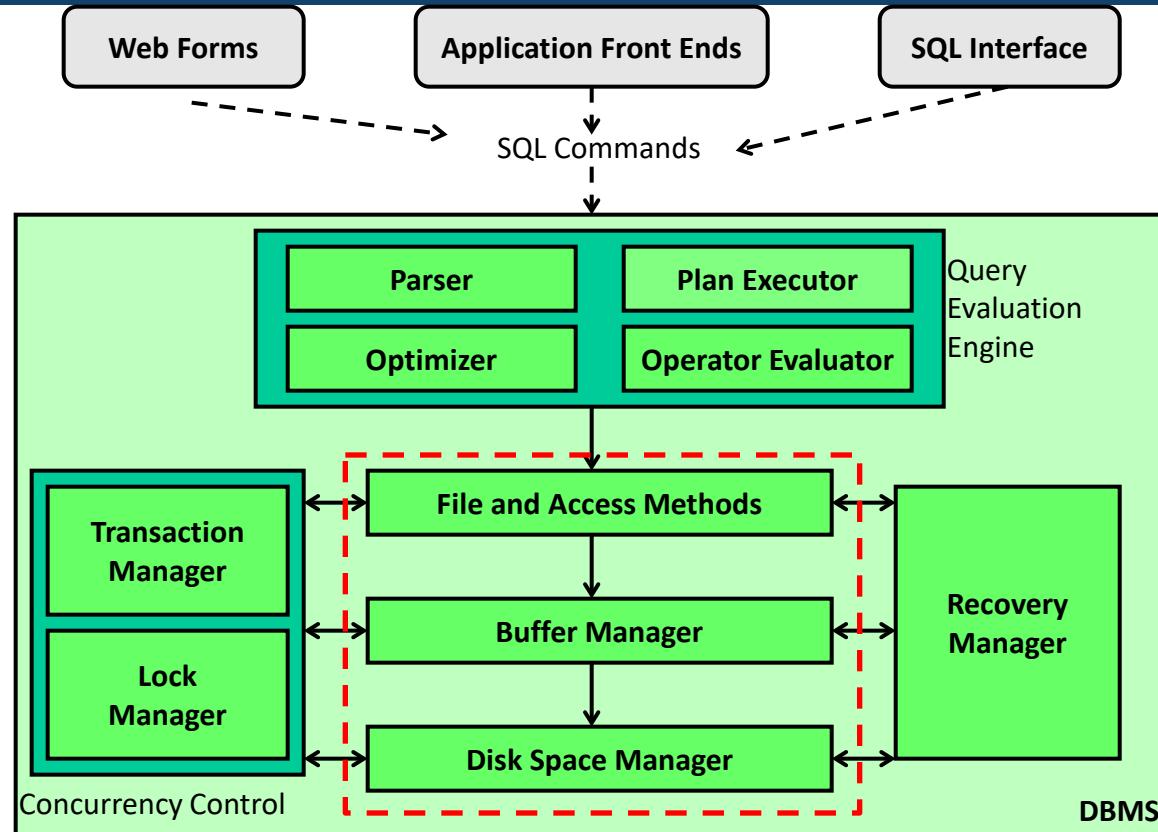
Two fundamental DBMS questions:

- How do we efficiently **organize** very **large volumes** of data?
- How do we **efficiently access** data, i.e., **minimize I/Os**?

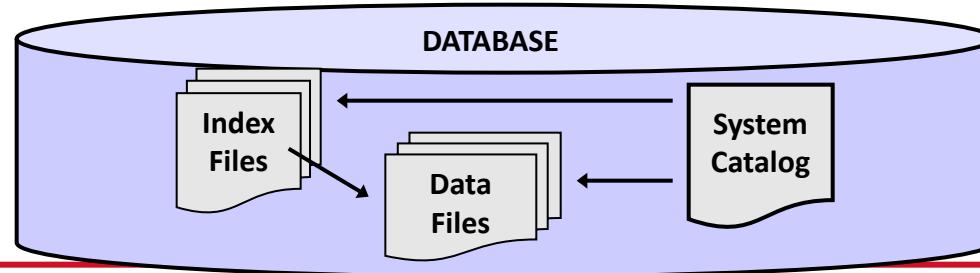
Important notes:

- **Databases need persistent storage**
 - We **cannot hold all database content in main memory**
 - We keep producing **massive** quantities of **data**, i.e., we **always produce more data** than we **can fit** in main **memory**
 - Even if **main memory** is **getting cheaper**, **emerging applications** require **massive storage** capabilities (e.g., IoT, social media, deep space exploration, science, bioinformatics, etc).

Internal Structure of a DBMS



手稿



Physical storage medium

- Permanent storage (external): nonvolatile or long-term
- Transient storage (internal). volatile or short-term

Permanent storage (secondary storage) is usually *relatively cheap*.

Transient storage (primary memory) is usually *relatively expensive*.

Transient storage:

- **Main memory**
- **Cache**

Main memory is used to store data that is being used for *on-going computations*.

Cache is used to store data in very fast computer memory chips to speedup computations.

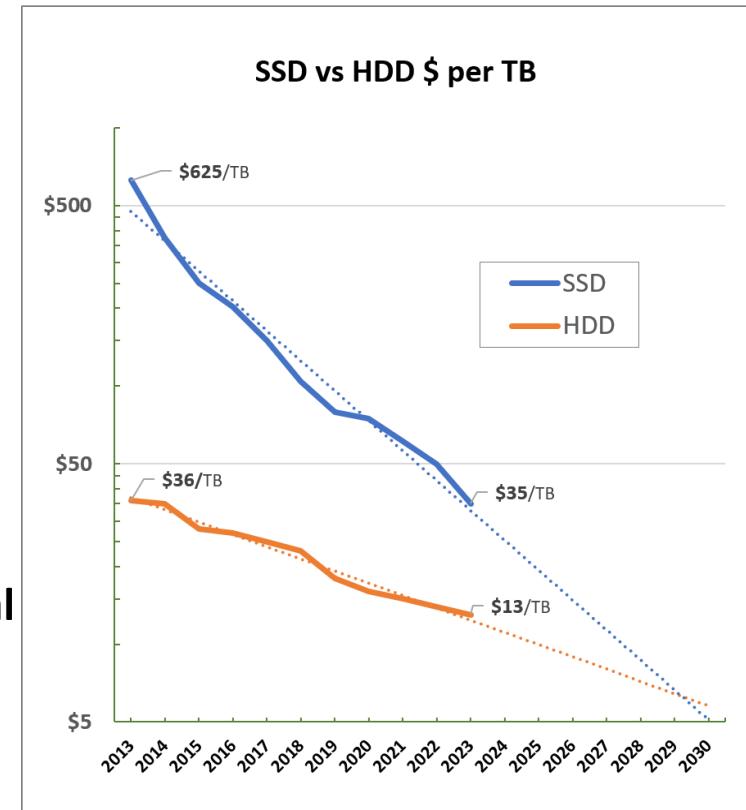
Cache sizes are **usually a lot smaller** than main memory sizes.

Secondary storage

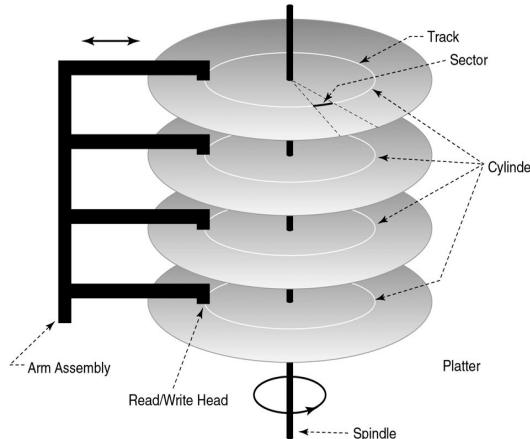
- **Magnetic disk (Hard Disk Drive - HDD):** HDD devices are **slower**, but they have a **large storage capacity**.
- **Solid-State Drive (SSD)**: SSD devices are **faster**, but they also **cost much more**, and **life span is shorter**

Tertiary storage

- **Magnetic cartridge tapes (tape silos):** only **sequential access**
- **Optical disk (juke boxes):** random access but much slower than HDD.



"live" databases are mostly **disk-based**. Therefore, we will focus on the **magnetic disk** medium.



Hard Disk (Disk):

- **Cheap:** 1TB for \$23 (2024)
- **Permanent:** Once on disk, data will be there until it is explicitly wiped out!
- **Slow:** Disk read/write are slow: 150MB/s to 250MB/s (2024)
- More coarsely addressable: **block addressable** (≥ 512 bytes)

Random Access Memory (RAM) or Main Memory:

- **Expensive:** For \$100 (2024), we get 16GB
- **Volatile:** Data is lost when a crash occurs, power goes out, etc!
- **Fast:** up to 8GB/s (2024), ~50 x faster for sequential access, ~100,000 x faster for random access, compared to disk access!
- More finely addressable: **bit addressable**

Focus in database research has mainly been on secondary storage:

Why?

Databases are ***computationally I/O bound*** – most operations are **read** and **write** operations **between disk and memory**.

CPU bound computations are ***far less significant*** compared with **I/O bound computations**

because

› Let us now turn to the representation of blocks on disks.

› **Block:**

- typically defined by the **Operating System**
- **unit of transfer between disk and main memory**
- usually spans more than one sector. E.g. of block size: 4096 bytes (4K), 8K,....
- **no concept of logical records** on disk
- once in **buffer**, a **block** is divided in **logical records** (i.e., tuples) by software

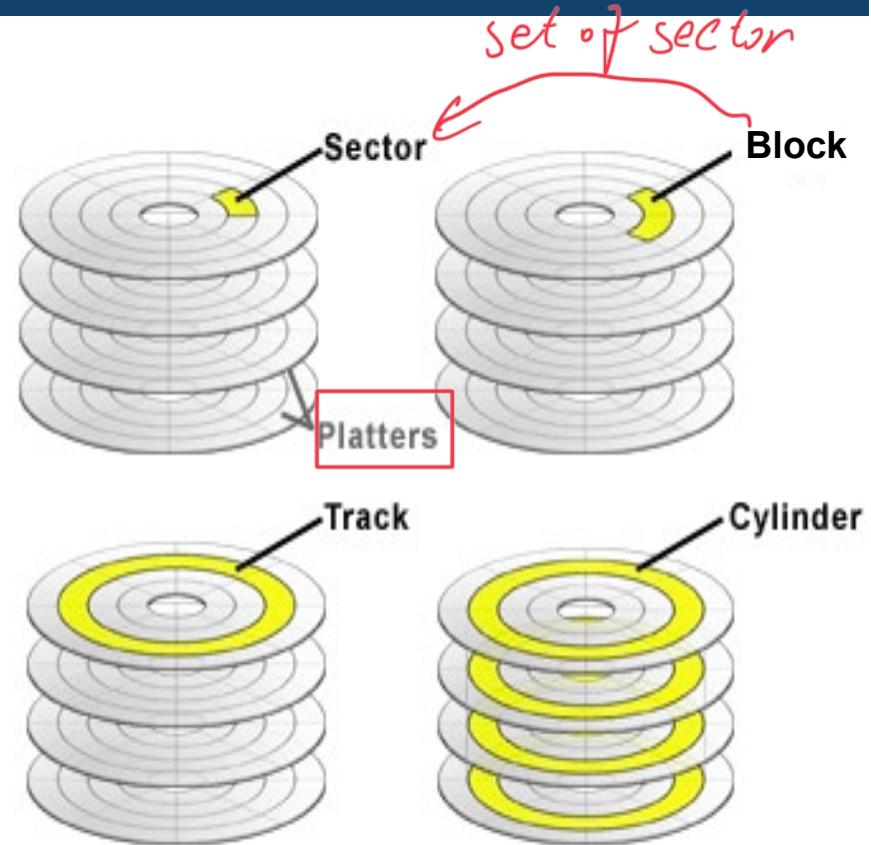
› **Blocking factor:** Block size vs. record size

- $b = \text{block size} / \text{record size}$
- **Record spanning** occurs when $b < 1$. This means that the record size is bigger than the block size, i.e., we need **more than 1 block to store a record!**
 - Important to **relational design** to prevent record spanning

bad design

Data is Organized as Blocks on Disk

- › Due to the *high access latency*, data is organized in the form of **data blocks** on disk, such that the unit of transfer of disk read/write is a **block**
 - **Block size** is set by the operating system (OS) when the disk is initialized, usually as a multiple of sector size (typically 4KB or 8KB)



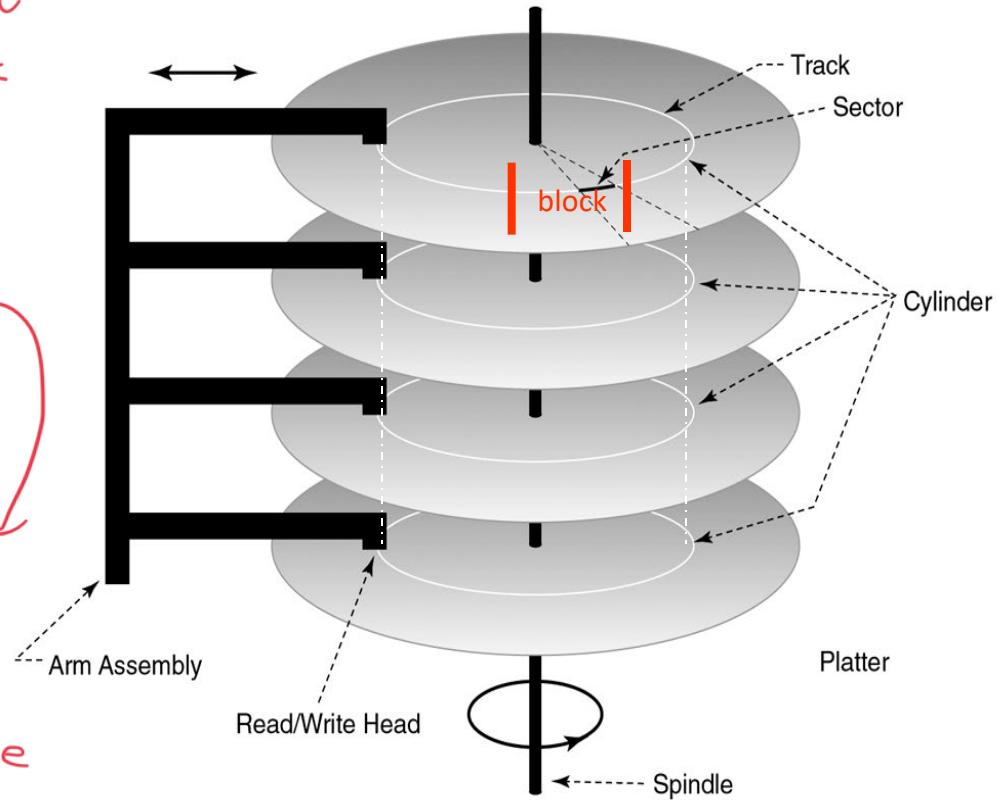
1KB = 1024 Bytes, 1MB = 1024*1024 Bytes

Data is Organized as Blocks on Disk

› Time for disk ***read/write*** *go to the right*

- 1) ***Seek time***: the arm assembly is *track* moved in or out to position a read/write head on a desired track
- 2) ***Rotational delay***: The platters spin
- 3) ***Transfer time***: Reading/writing the data

wait for platters spin to the right place



Physical organization of a disk storage unit.

Buffer management:

- › Note that **databases** are usually **too large** to be held **totally** in **main memory**:
- › The database **buffer** is a **main-memory area** used to **store database blocks**.
 - **Buffers** provide virtual memory for operations to be performed in main memory.
 - **Transfers** are done in **blocks** through a **buffer**.
 - Because **I/Os are expensive**, we try to **keep as many blocks as possible in memory** for later use.
 - Devise **strategies** that will “guess” which block the application will ask for, in such a way that **most likely it** will be in the **buffer** when **requested**.
 - **Buffer manager** is **responsible** for **managing** such **strategies**.

Buffer management:

- › Define a ***buffer replacement strategy***:
 - If the buffer is full and a program needs a block, determine what block in buffer should be **replaced**, i.e., what is the **best replacement algorithm**?
- › Two main algorithms:
 - **LRU** (Least Recently Used)
 - **MRU** (Most Recently Used).
- › There are other variants as well.



Consider this natural join: **Borrower \bowtie Customer**

Assume **each relation** is in a **separate file** and **neither** one **fits** into main memory. Assume that the *only common attribute* in the two relations is *customer-name*. Let the following **program implement** the **join** above:

```
for each tuple b of Borrower do
    for each tuple c of Customer do
        if b[customer-name] = c[customer-name] then
            let x be a tuple defined as follows:
                x[branch-name] = b[branchname]
                ....
                x[customer-city]= c[customer-city]
            add x to result
        endif
    end
end
```

Focusing on only the buffering of the *Customer* relation, which buffer replacement strategy (LRU or MRU) would be better suited for this operation?

Focusing on **only** the **buffering** of the *Customer* relation:

1st strategy: LRU (Least Recently Used)

- In this scheme, the *least recently used block* (assuming n tuples fit into **1 single block**) is the **victim** (i.e., *selected to be replaced*).

```

for each tuple b of Borrow do
    for each tuple c of Customer do
        if b[customer-name] = c[customer-name] then
            let x be a tuple defined as follows:
                x[branch-name] = b[branchname]
                ....
                x[customer-city]= c[customer-city]
            add x to result
        endif
    end
end

```

However, according to the **join** program, a **more efficient strategy** would have been to **select the victim the *most recently used* block!** Why?

- if a Customer block is used, it will **not be used again until all others have been used.***

why we choose
MRU for Natural
join



2nd strategy: MRU (Most Recently Used)

- The **most recently used record is tossed out**. This makes sense in our case:
 - This is obviously an optimal strategy for our join

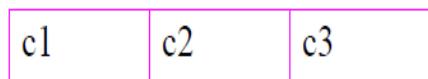
```
for each tuple b of Borrow do
    for each tuple c of Customer do
        if b[customer-name] = c[customer-name] then
            let x be a tuple defined as follows:
                x[branch-name] = b[branchname]
                ....
                x[customer-city]= c[customer-city]
            add x to result
        endif
    end
end
```

Assume the buffer consists of 3 blocks, number of **Customer** records is **4** (**c1, c2, c3, c4**) and *one Customer record fits in exactly one block.*

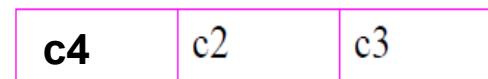
After 3 accesses, the buffer looks like (a): **c1** is the *least recently used* and **c3** is the *most recently used*. Next operation is **Read c4**. So, **we must swap one Customer record out**.

LRU: **c1** is the **least recently used** and therefore it is **selected as the victim**. After this read operation, c1 is replaced by c4 as shown in (b).

According to the join algorithm: next, **Read c1**. The **victim** is **c2** as shown in (c). This implies we have a total of **2 page faults** (i.e., 2 I/Os).



(a)



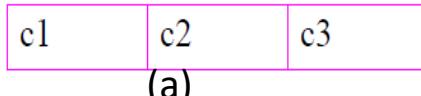
(b)



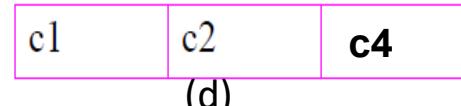
(c)

Contrast with **MRU:** **Read c4:** the **next victim** is **c3** to be swapped out as shown in (d).

Next, **Read c1 : no I/O** required in this case. Therefore, the total **page faults** is **1. Better than LRU in this case.**



(a)



(d)



Measuring the **performance** of hard disks

$$\begin{array}{lcl} \textbf{Disk Access} & = & \textbf{Seek time} + \textbf{Rotation latency} \\ (5\text{ms} - 10\text{ms}) & = & (3\text{ms} - 4\text{ms}) + (2\text{ms} - 7\text{ms}) \end{array}$$

Data Transfer Rate = 40 - 200MB/s

Typically, between **5ms** and **25ms** per **MB** transferred

multiple disks with shared interface can support **higher rates**

- › Fact of life: Disks **do fail** every now and then.
- › **MTTF** is a statistical value that describes the average operating time until the disk fails.
 - Vendors' claim: Depending on the disk model, MTTF is between one million hours (114 years!) and 2.5 million hours (285 years!) - For a single disk, MTTF is of *limited significance*.
 - A disk could fail at any time, so regular backups and RAID configurations are required to protect against data loss.
- › In practice, MTTF is helpful with a large number of disks where MTTF helps to estimate how regularly failures could occur.
 - Example: Assuming a *uniform statistical distribution* failure rate, with an MTTF of **one million hours** and **one million disks**, a **disk failure** would be expected to occur **each hour**. With the **same MTTF** but with **1,000 disks**, a **disk failure** is expected to occur **every 1,000 hours!**

$$\frac{1 \text{ million}}{1 \text{ million}} = 1/\text{hour}$$

$$\rightarrow \frac{1000}{1 \text{ million}} = \frac{1}{1000 \text{ hours}}$$

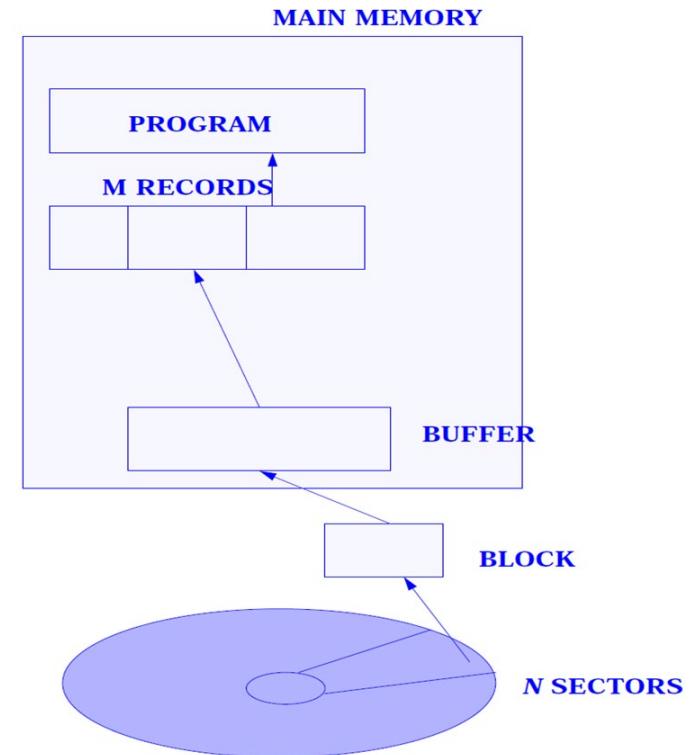
Key approaches for ***physical disk-based transfer***:

- Block transfer
- Cylinder-based
- Multiple disks
- Disk scheduling
- Prefetching/double buffering

10

Block transfer:

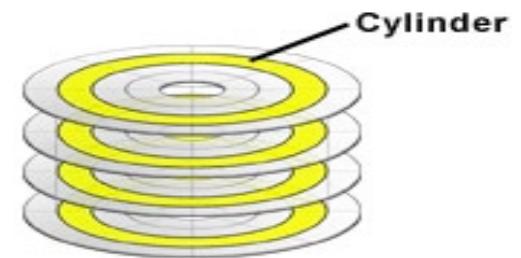
- Data is always **transferred by blocks** to *minimize transfer time*
- A disk block is a set of *contiguous sectors from a single track of one platter*
- Block sizes typically range from 4k bytes to 32k



(2)

Cylinder-based Transfer

- Important observation: *data in relations is likely to be accessed together.*
- Principle: store relation data on the *same cylinder*: *blocks in the same cylinder* effectively involve only *one seek* and *one rotation latency*.
- *Excellent* strategy if access can be .



Multiple disks

(3)

- Observation: disk drives continue to become **smaller** and **cheaper**.
- Idea: use **multiple disks** to support **parallel access** - also enhances **reliability**.

RAID: Redundant Array of Inexpensive (Independent) Disks

- *Load balance* multiple small accesses to *increase throughput*
- *Parallelize* large accesses so the *response time* is *reduced*.

RAID: Redundant Array of Inexpensive (Independent) Disks

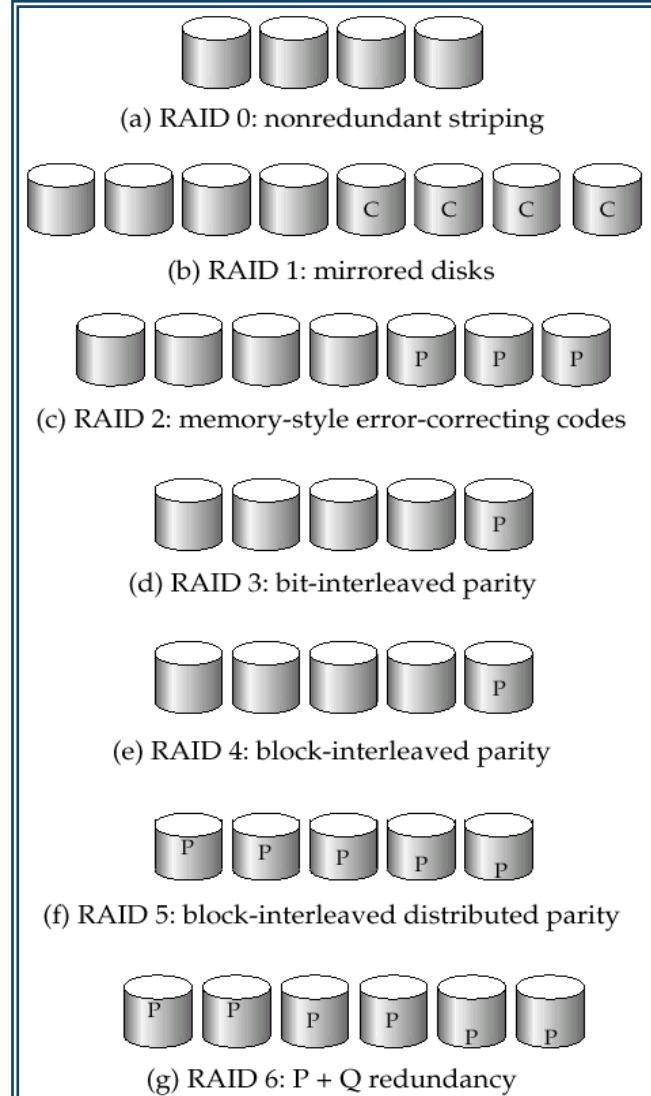
- disk organization techniques that manage a large numbers of disks, providing a view of a single disk of
 - high capacity and high speed by using multiple disks in parallel, and
 - high reliability by storing data redundantly, so that data can be recovered even if a disk fails

RAID levels:

- Schemes to provide redundancy at lower cost by using disk striping combined with parity bits
- Different RAID organizations, or RAID levels, have differing cost, performance and reliability characteristics

More info can be found at

https://en.wikipedia.org/wiki/Standard_RAID_levels



Disk scheduling: Elevator scheduling

- Principle: *schedule readings of requested blocks in the order in which they appear under the disk head.*

Elevator approach:

- The arm moves towards one direction *serving all requests on the visited tracks.*
- *Changes direction when no more pending requests.*
- **Advantage:** *reduces average access time for unpredictable requests*
- **Disadvantage:** *benefit is not uniform among requests and does not work well when there are only a few requests*

Prefetching

- **With Buffering:**

- Goal is to **keep as many blocks in memory** as possible to **reduce disk accesses**.

- **Prefetching or double buffering:**

- Situation: when *needed data is known in advance*

- Principle: speed-up access by *pre-loading* needed data.

- Problem: requires *extra memory*

Short break

please stand up and stretch

Let us also menti....



THE UNIVERSITY OF
SYDNEY

- › A table in a DBMS is **stored** as a collection of records, referred to as a file
 - Each file consists of one or more pages
 - Each page stores records of one type (i.e., from one table)
 - A page is the smallest data unit in main memory that the **buffer manager** is in charge of.

Disk space manager

- Translates **page** requests into **disk block** requests.
- Each page typically consists of one or more disk blocks

In our **subsequent discussion**, we will assume that a page consists of a single block unless otherwise indicated.

Relation(tuplekey, attribute1, ...)



Table size: assume there are **2,000,000 records** in the table **Relation** and each **record** is **200 bytes** long, including a *primary key tuplekey of 4 bytes*, an **attribute attribute1 of 4 bytes**, along with **other attributes**.

General features: assume that each **page size** is **4K bytes**, of which **250 bytes** are reserved for *header* and *array of record pointers*.

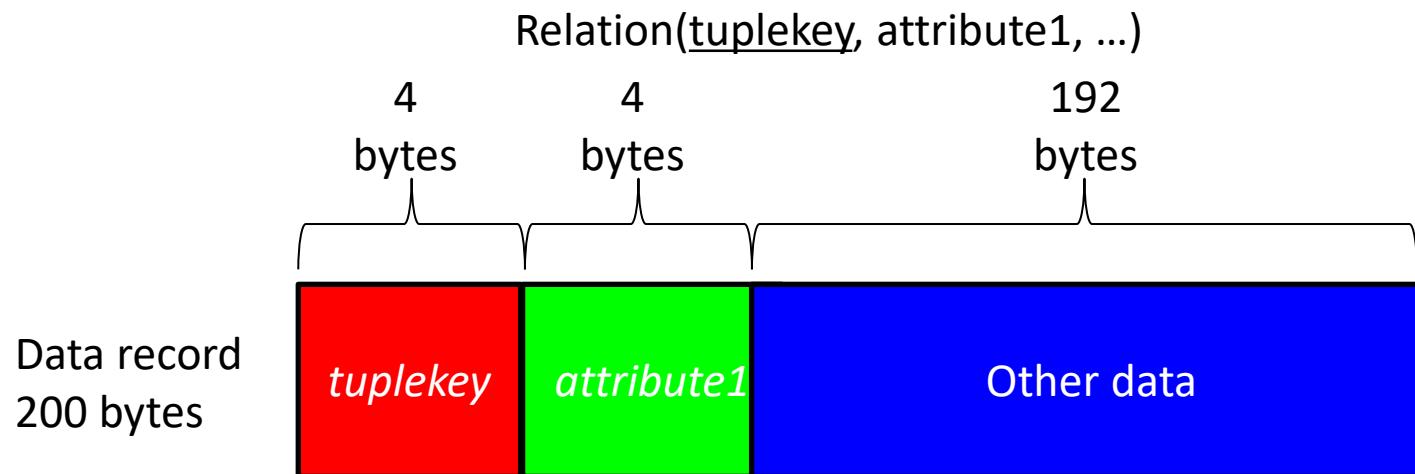
Questions:

- › *How many bytes are required for a record?*
- › *How many records can fill a page?*
 - We assume that **no record is ever split** across pages, i.e., **no record spanning.**
- › *How many pages are required to store the table?*

- › Databases typically **store** each **tuple** (row) in a **record**
 - *All fields of a record are stored together **consecutively***

From earlier

“each record is **200 bytes** long, including a primary key called **tuplekey** of 4 bytes, an attribute called **attribute1** of 4 bytes, and also other attributes”



How Many Records per Page?

Some space consumed by header/pointer data

250 bytes

*All pages in DB have the same size,
commonly 4KB (4096 bytes)*

$$4096 - 250 = 3846 \text{ bytes}$$

Remaining space available for storing data, e.g., record data, index entries

From earlier

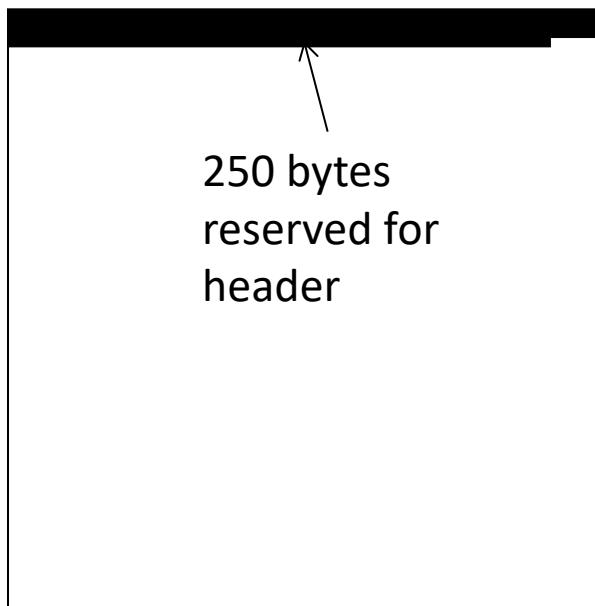
“each page is 4K bytes, of which 250 bytes are reserved for header and array of record pointers”

How Many Records per Page?

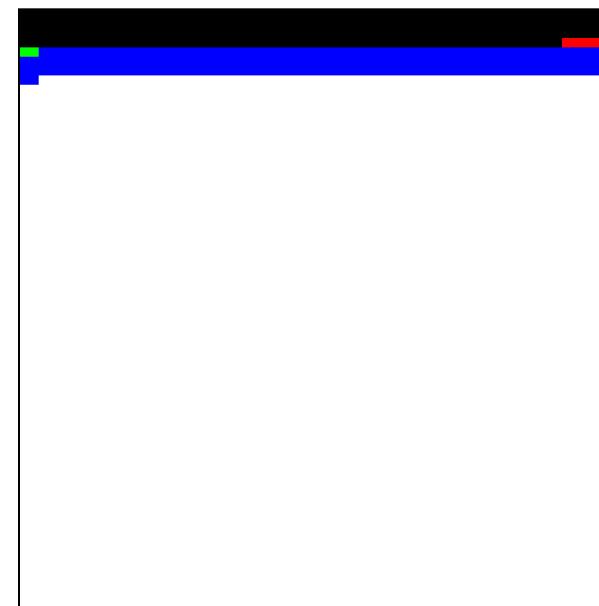
Each example data record 200 bytes



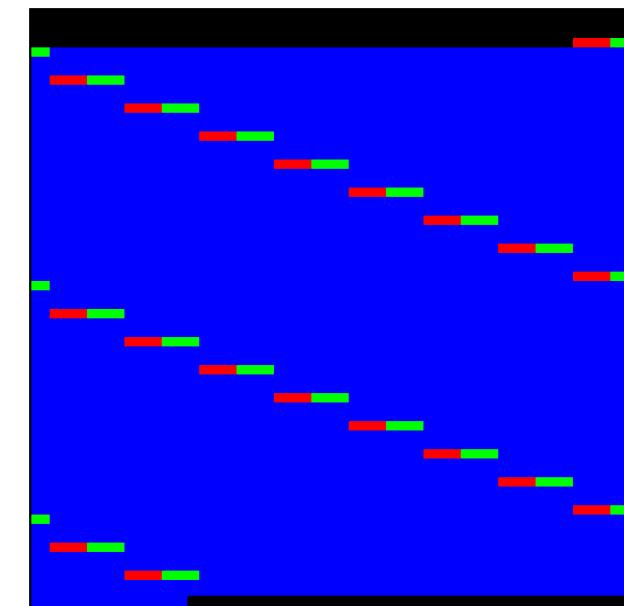
↑
250 bytes reserved for header



records: 0
empty space: 3846 bytes



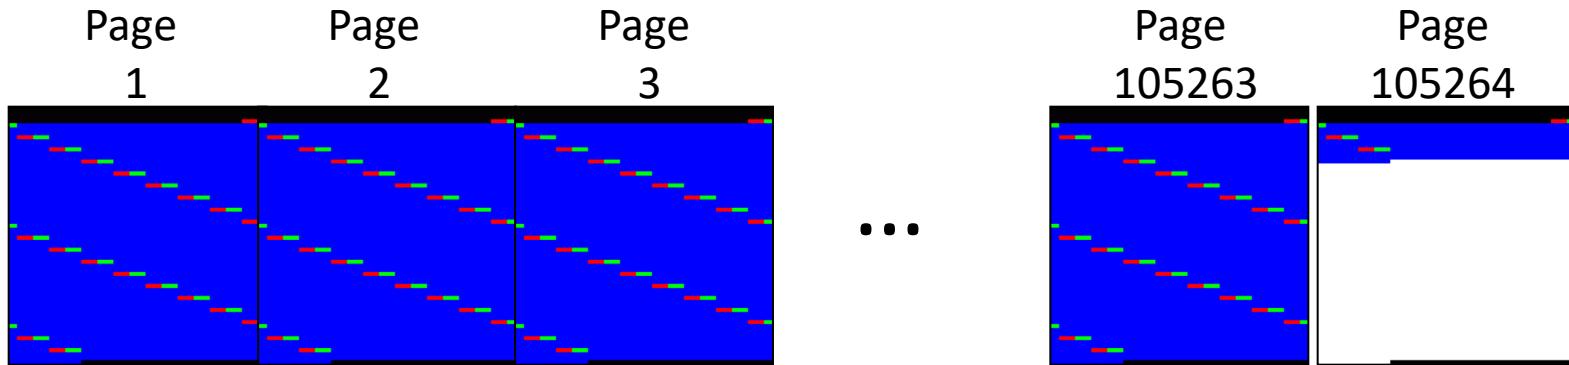
records: 1
empty space: 3646 bytes



records: 19
Remaining space: **46 bytes**

$[3846 \text{ bytes/page} \div 200 \text{ bytes/record}] = [19.23] = 19 \text{ records/page plus 46 remaining bytes}$

How Many Pages for the Table?



$$\left\lceil \frac{2,000,000 \text{ records}}{19 \text{ records/page}} \right\rceil = 105,263 \text{ full pages plus 1 page containing 3 remaining records} = 105,264$$

Total size of the required pages = 105,264 pages × 4096 bytes/page = 431,161,344 bytes

Actual Table size = 2,000,000 × 200 = 400,000,000 bytes

$$\text{Overhead} = \frac{431,161,344 - 400,000,000}{400,000,000} = 7.79\% \text{ (~8% overhead)}$$

- What is a **fill-factor** (also referred to as **occupancy** or **load factor**)?
 - The fill-factor value determines the **percentage** of space on each page to be **filled** with data, reserving the **remainder on each page** as **free space** to help in the **index expansion** when **new data is inserted**, **without the need to split the index page**.
 - The fill-factor value is a **percentage** from **1** to **100** with the value **100** meaning that the pages **are filled to capacity**.
- The **Fill-factor** is provided for **performance reasons**
- For example, specifying a **fill-factor** value of **80** means that **20 percent** of each page will **be left empty**.
- Further discussion can be found @ <https://www.geeksforgeeks.org/sql-fill-factor-and-performance/>

How Many Pages With a Fill Factor of Less Than 100?

Page

1

Page

2

Page

3

Page

105264

Page

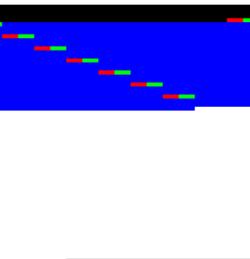
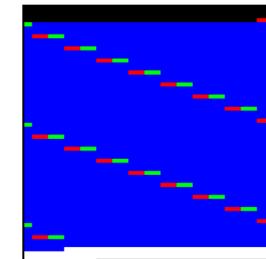
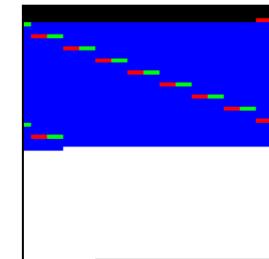
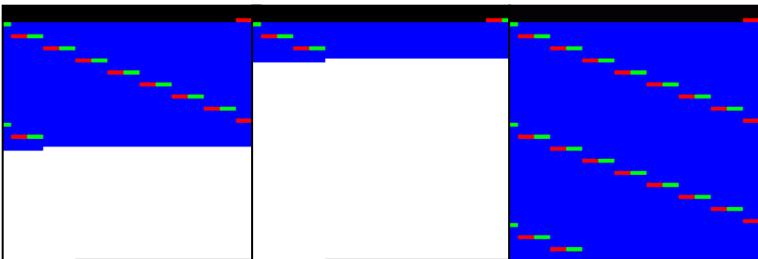
142851

Page

142858

...

...



Assume pages are 75% full on average: use records/page as the initial size of the page

$$\lfloor 19 \text{ records/page} \times 75\% \text{ average occupancy} \rfloor = \lfloor 14.25 \rfloor = 14 \text{ records/page (rounded down)}$$

$$\text{Total pages} = \left\lceil \frac{2,000,000 \text{ records}}{14 \text{ records/page}} \right\rceil = 142,858 \text{ pages}$$

$$\text{Total size} = 142,858 \text{ pages} \times 4096 \text{ bytes/page} = 585,146,368 \text{ bytes}$$

Using the previous **overhead formula**, this equates to ~47% overhead

$$\frac{585,146,368 - 400,000,000}{400,000,000} = 46.28\% \text{ (~47% overhead)}$$

*trading space
for performance*

File Organisation: a method of arranging the records stored on disk.

- › **Unordered (Heap) Files** – a record can be placed **anywhere** in the file **wherever** there is space (**random order**)
 - suitable when the typical access is a **file scan** retrieving all records.
- › **Sorted Files** – store records **in sorted order**, **based** on the **value** of the **search key** of each record
 - may **speed up** searches using **binary search**, however **updates** can be **very expensive**
- › **Indexes** – a **specialized data structure** to **organize** records via **trees** or **hashing**
 - like sorted files, they **speed up searches** based on values in certain fields (“**search key**”).
 - **updates** are much **more efficient** than in **sorted files**.

*logical sort
not physical*

› **Physical Data Organization: how is data physically stored in DBMS?**

› **Access Paths: how are records retrieved from DBMS?**

- Access methods for heap files (linear scan)
- Access methods for sorted files (binary search)
- Access methods for indexes (index scan)

› **B+ Tree Index**

- Primary index
- Composite search keys

- › An **Access Path** is a **method** for **retrieving** records, and **refers to the data structure + algorithm** used for **retrieving** and **storing** records in a table
 - **Linear scan:** *Unordered (Heap) files* can be accessed with a *linear scan*
 - **Binary search:** *Sorted files* can be accessed with a *binary search*
 - **Index scan:** *Indexed data* can be accessed with an *index scan*

› Physical Data Independence:

- › Example: **SELECT * FROM Student WHERE sid = 5309650;**
 - The choice of an **access path** to use in the execution of the above SQL statement has *no effect* on its *semantics*
 - *However*, any choice may have a *major effect* on the *execution time* of the SQL statement

- › Unordered (**heap**) file is the simplest file structure, which contains records in **no particular order**.
- › Access method is ***linear scan (file scan, table scan)***, as records are unsorted in heap files.
 - On average ***half of all pages*** in a file must be read, but in the ***worst case*** the ***whole file*** must be read.

› **Linear Scan** (sequential search): For each page,

- 1) Load the page into main memory (cost = 1 I/O);
- 2) Check each record in the page for match;

Assume we need 140,351 pages to store our Table: How many page I/Os are needed to find records for:

- **SELECT * FROM Relation WHERE tuplekey=715;**

› For **equality search**, if tuplekey is **unique**, so can terminate on first match.

- › If a matching record is present, it will on average look through half of all pages, so we require 70,176 I/Os
- › For zero matching records or non-unique attribute, it needs to check every record, so we require 140,351 I/Os

How many page I/Os (out of 140,351 pages) are needed to find records for:

- **SELECT * FROM Relation WHERE attribute1 BETWEEN 100 AND 119;**

› For **range search** need to check each record, so 140,351 I/Os

Speeding up query processing

- › Store records in a **sorted file** based on *some attribute*
 - Suppose we want to **search** for people of a specific **age**, and the records are sorted by **age**

11,12

13,14

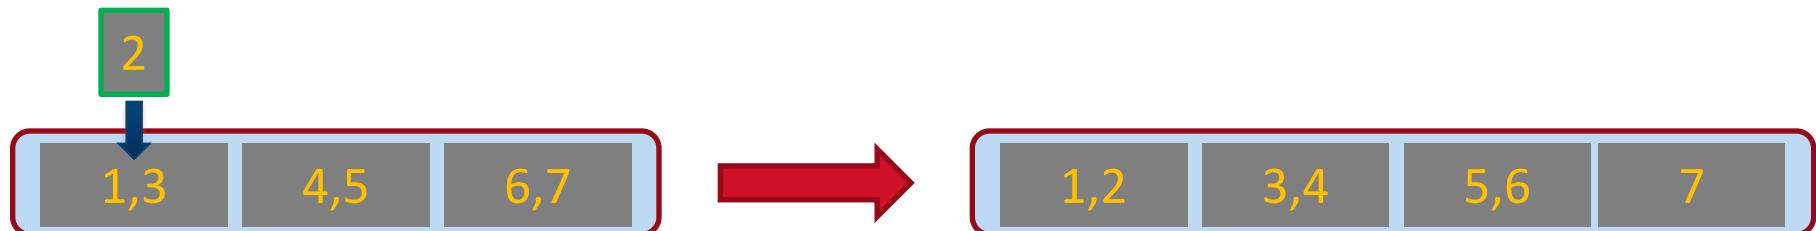
15,16

17,18

19,20

- › Access method is ***binary search***
 - The I/O cost (worst case scenario) is **$\log_2 B + \text{any additional pages containing retrieved records}$**
(**B** is the **total number of pages** in the file)
 - If $B = 140,351$ pages, the I/O cost for retrieving the page containing first record of Table is **$\log_2 B = \log_2 140,351 = 18$**

- › Issue: It is **expensive** to maintain a sorted order when **records are inserted**
 - After the **correct position** for an insert has been determined, it **needs to shift** all **subsequent records** to make space for the **new record**

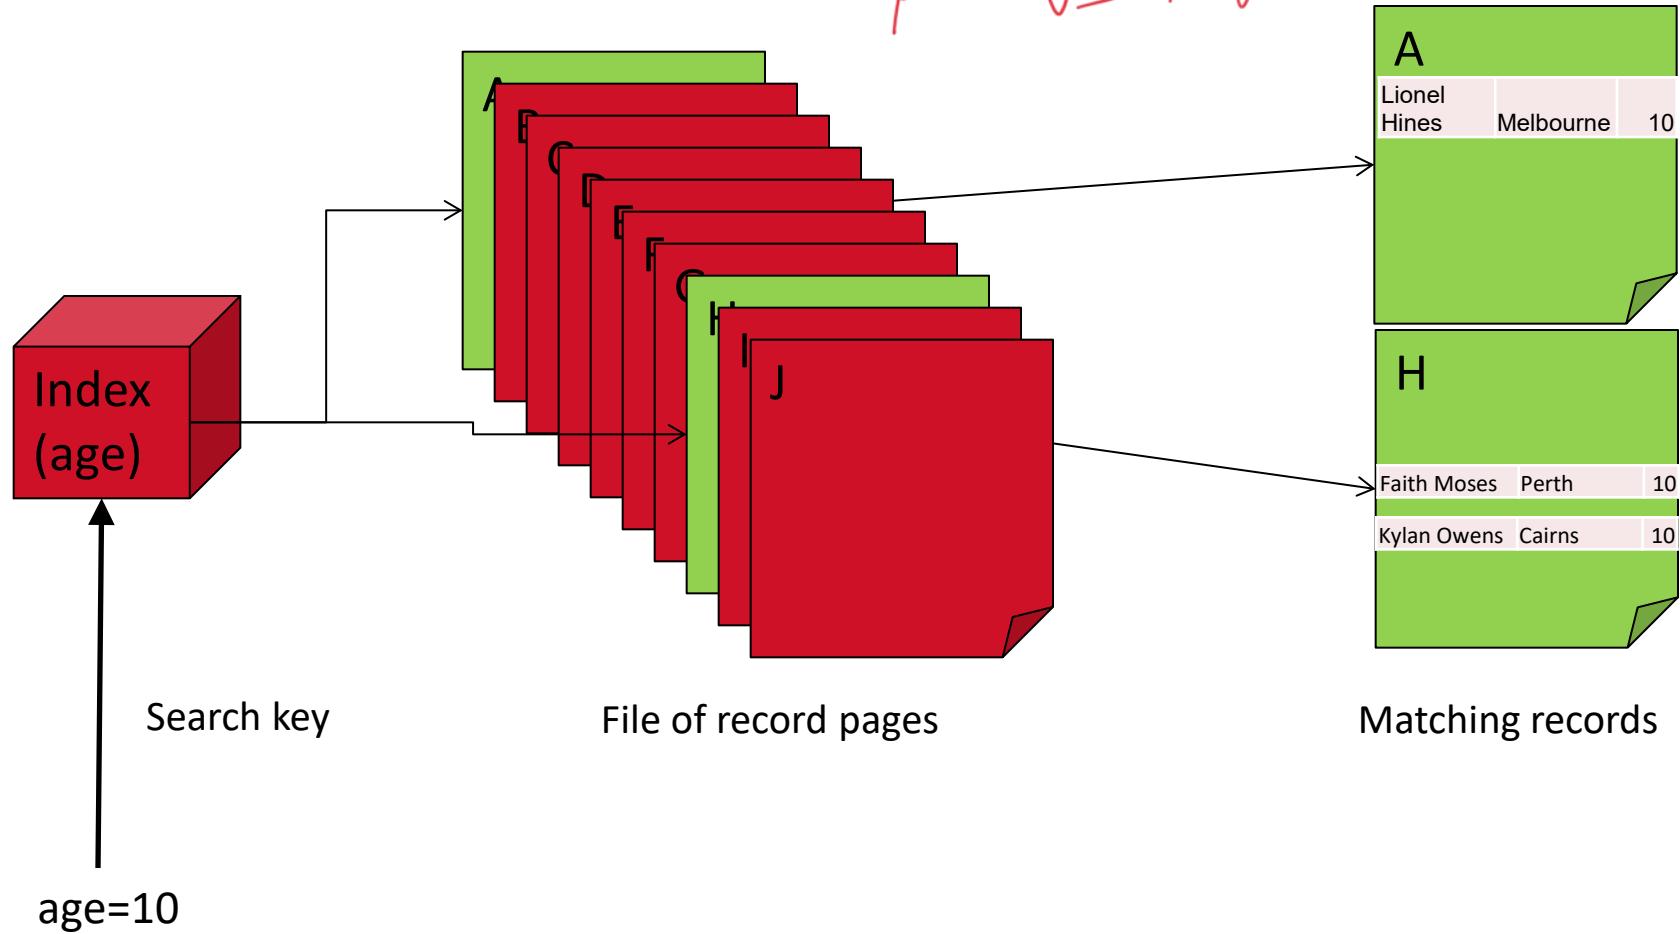


- Hence, sorted files are typically hardly used by commercial DBMS, but rather they use *index-organised (clustered) files*

Speeding Up Access via Indexing

- To speed up access (i.e., reduce time), we add additional information to facilitate query answering

space \sqrt{S} Performance

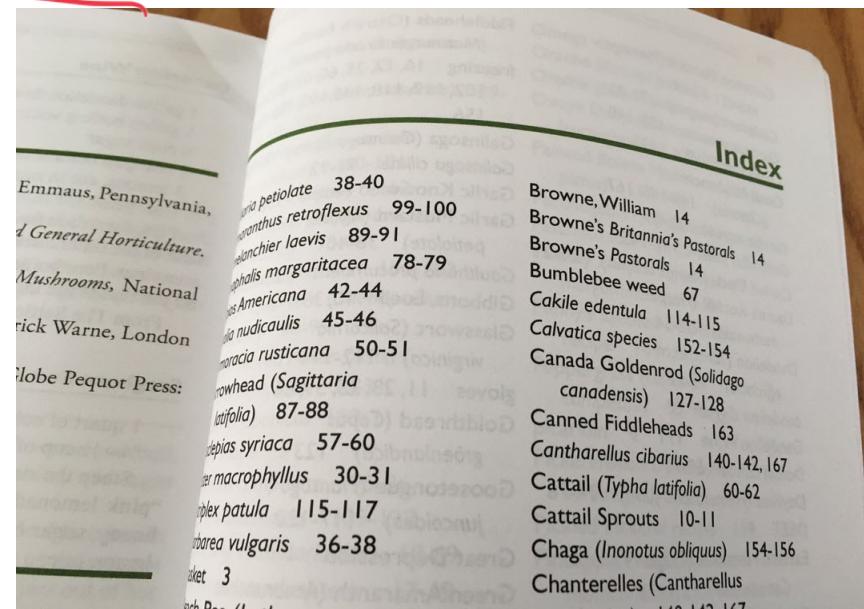


- › An **index** is a data structure mapping search key values to sets of records in a database table

- **Search key** properties
 - Any subset of fields
 - is not necessarily the same as the **primary key** of a relation
 - Primary key is *name* and search key may be *price*

↙ not necessary PK
used to access database

Product(name, maker, price)



- › An **index** is a data structure mapping search key values to sets of records in a database table

- **Search key** properties

- Any subset of fields
- is not necessarily the same as the **primary key** of a relation
 - Primary key is *name* and search key may be *price*

Product(name, maker, price)

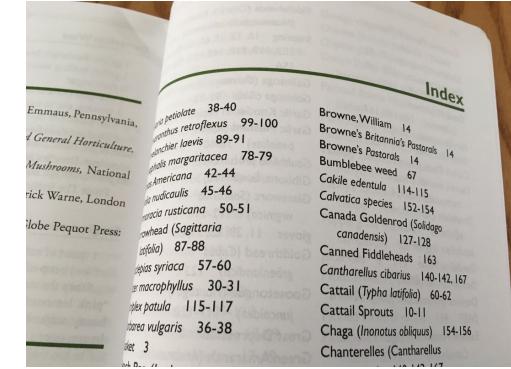
- › An index provides efficient lookup & retrieval by search key value (**index scan**)
 - usually **much faster** than linear scan

- › Typical **index types**

- {
 - **Hash index** (can only be used for **equality** search)
 - **B+ tree index** (good for **range search**, but can also be used for **equality search**)

› Downside:

- Additional storage space to store index pages
- Additional I/O to access index pages
(except if index is small enough to fit in main memory)
- Index must be *updated* when table is modified.
 - depending on index structure, this could be quite costly
- › Nevertheless, the **benefits** of using indexing far **outweigh** the **drawbacks**.
indexing is used in all DBMSs to improve **performance**.
- › **Indexing** is one of the most important **features** provided by a DBMS for **performance** purposes.



› Physical Data Organization: how is data physically stored in DBMS?

› Access Paths: how are records retrieved from DBMS?

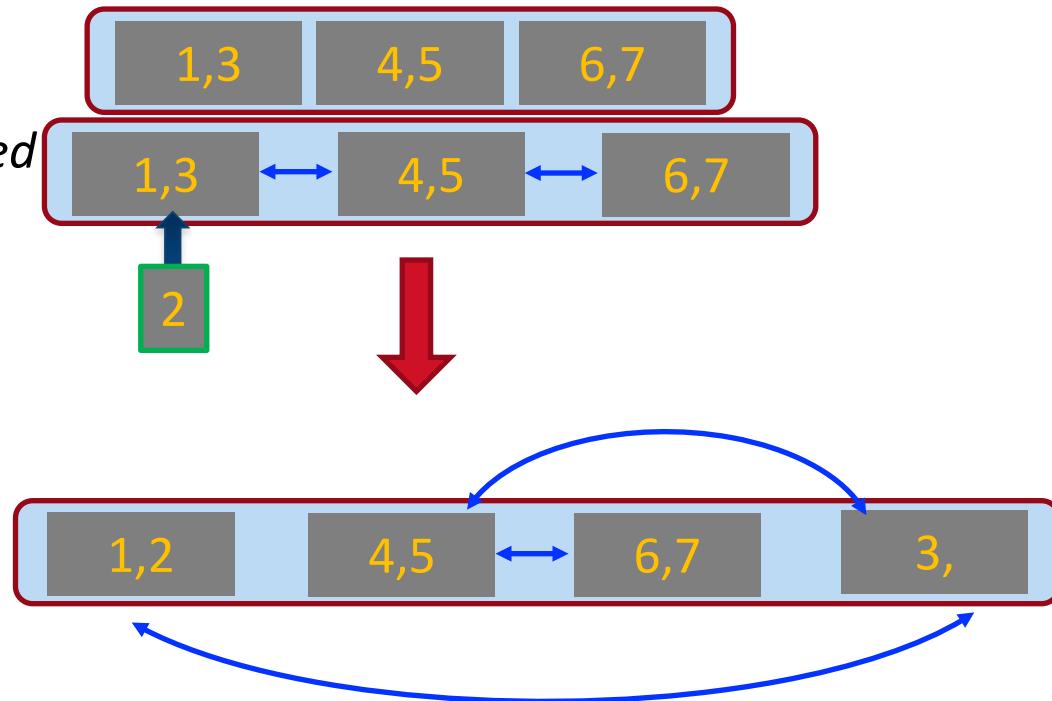
- Access methods for heap files (linear scan)
- Access methods for sorted files (binary search)
- Access methods for indexes (index scan)

› B+ Tree Index

- Primary index
- Composite search keys

- › Maintain a ***logically sorted*** file to cope with **insertions/deletions**
- › Example: Given the following set of pages

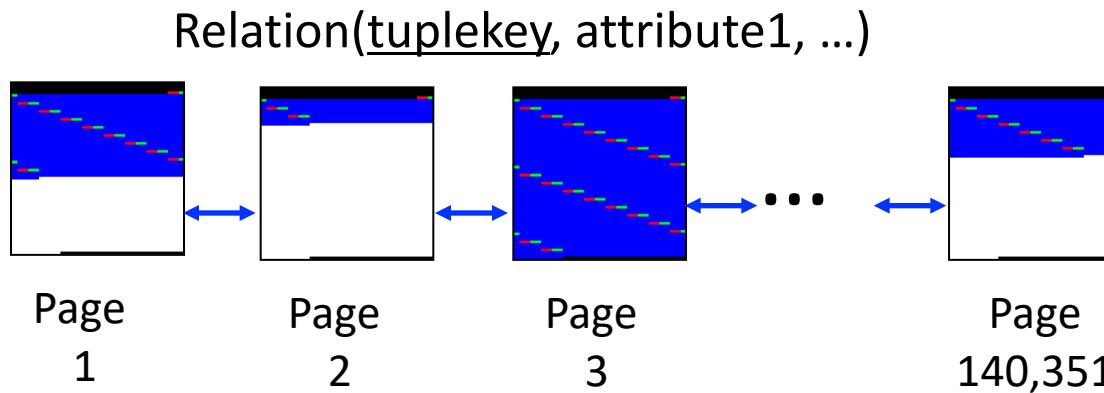
- › Not ***physically contiguous but sorted***



- › **Bad news:** since this file is ***not physically sorted***, we **cannot perform a binary search** (i.e., we can perform a binary search on an array, but not on a linked list)
- › **Good news:** we can build an **index** for the data file

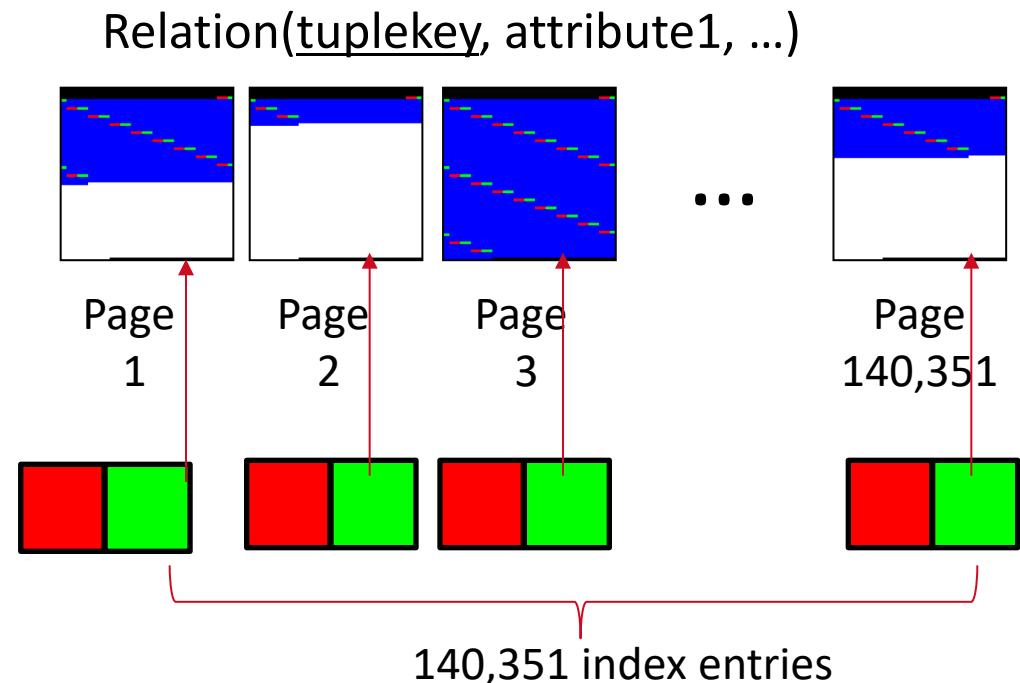
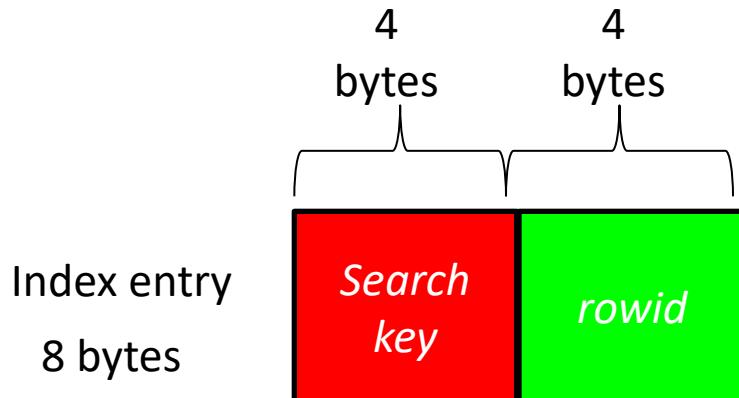
Step 1: Create a Logically Sorted Data File

- › Assume a table **Relation** is stored in a *data file* consisting of **140,351 pages** with **search key tuplekey**
 - For each page, all **records within a page** are *physically sorted* on the *search key* (let us assume in increasing order)
 - **Pages are *logically sorted*:** this means that all records in page i have smaller search key values than all records in page $i+1$
 - Note: these pages **may not be stored consecutively** on disk
 - **We call this file a logically sorted file based on search key tuplekey**



Step 2: Create One Index Entry for Each Record Page

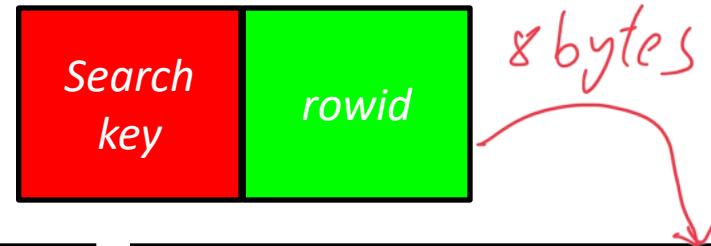
- The table Relation is stored in a ***logically sorted data file*** consisting of 140,351 pages with search key represented by ***tuplekey***
- Create one **index entry** for each page of the data file, where the index entry stores a ***pointer rowid*** which points to ***the data page***



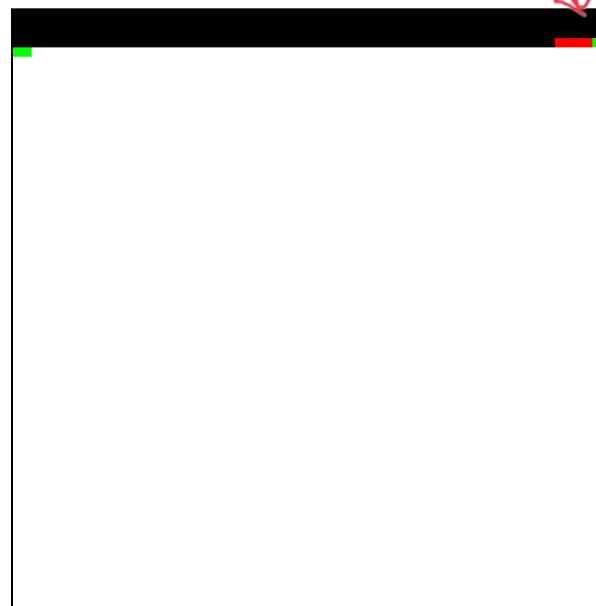
Index on ***tuplekey***: search key is the size of ***tuplekey***.
rowid is a 32-bit pointer (so 4 bytes).

Step 3: Fit Index Entries into Pages

Each example index
Entry is 8 bytes



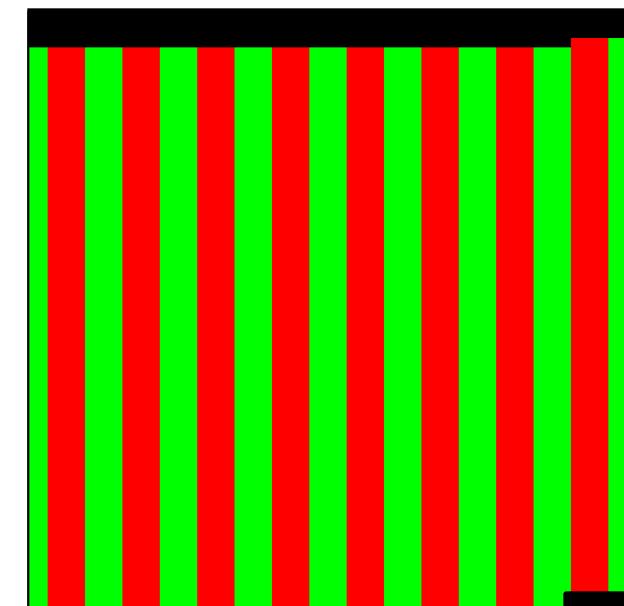
↑
250 bytes reserved



entries: 0
empty space: 3846 bytes

$$[3846 \text{ bytes/page} \div 8 \text{ bytes/entry}] = 480 \text{ entries/page}$$

entries: 1
empty space: 3838 bytes



entries: 480
empty space: 6 bytes

plus 6 remaining bytes

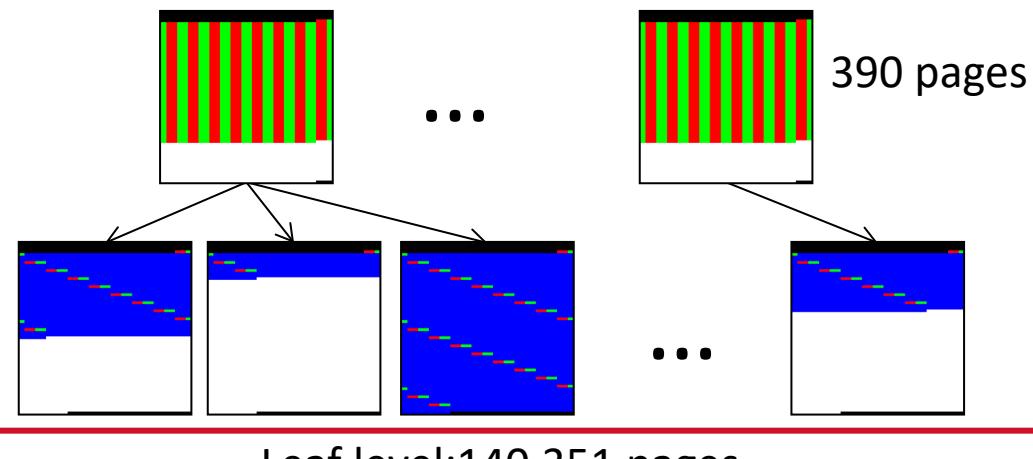
Step 3: Fit Index Entries into Pages (cont')

- Let us call a page storing *index entries* an *index page*

- Recall each index page holds up to 480 index entries
- Assume average occupancy (i.e., fill factor) of 75%
- The average number of index entries in an index page is $480 \times 75\% = 360$

- The number of **index pages** is $\lceil 140,351 / 360 \rceil = 390$ (rounded up)

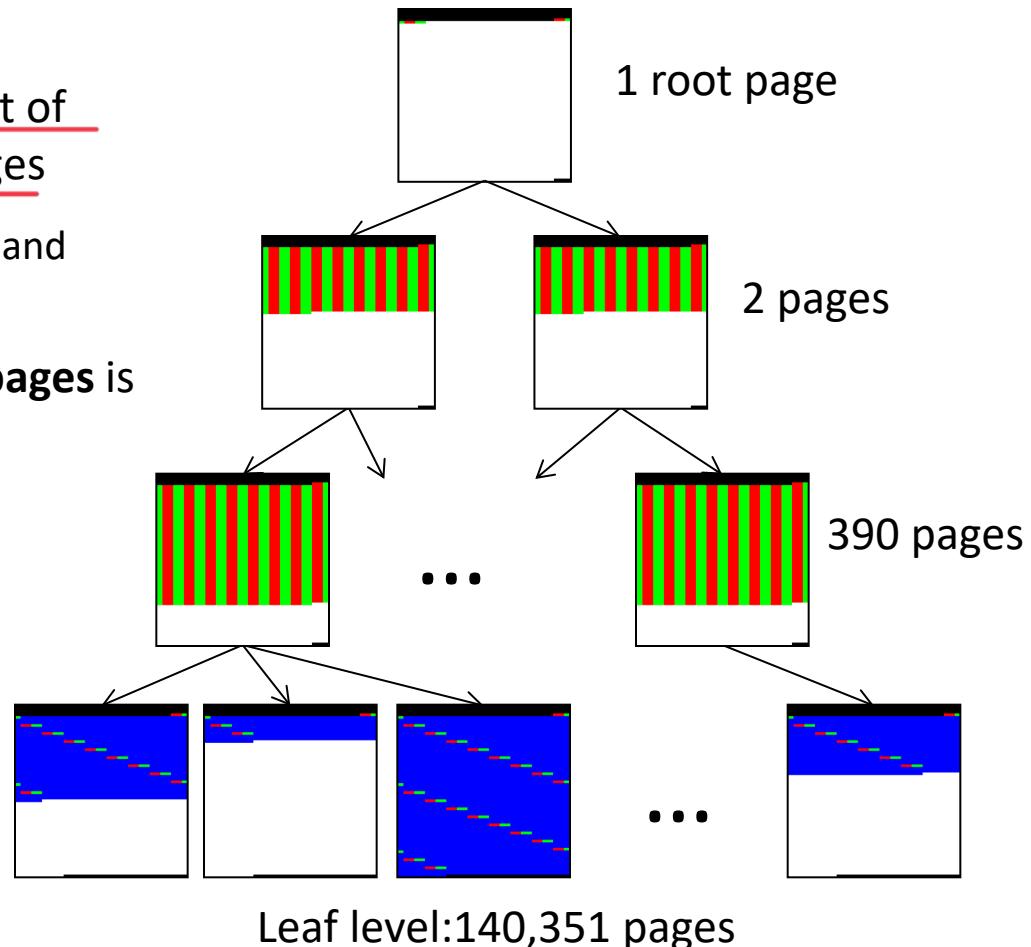
- We could store these 390 index pages as a logically sorted file and then recursively build indexes on it

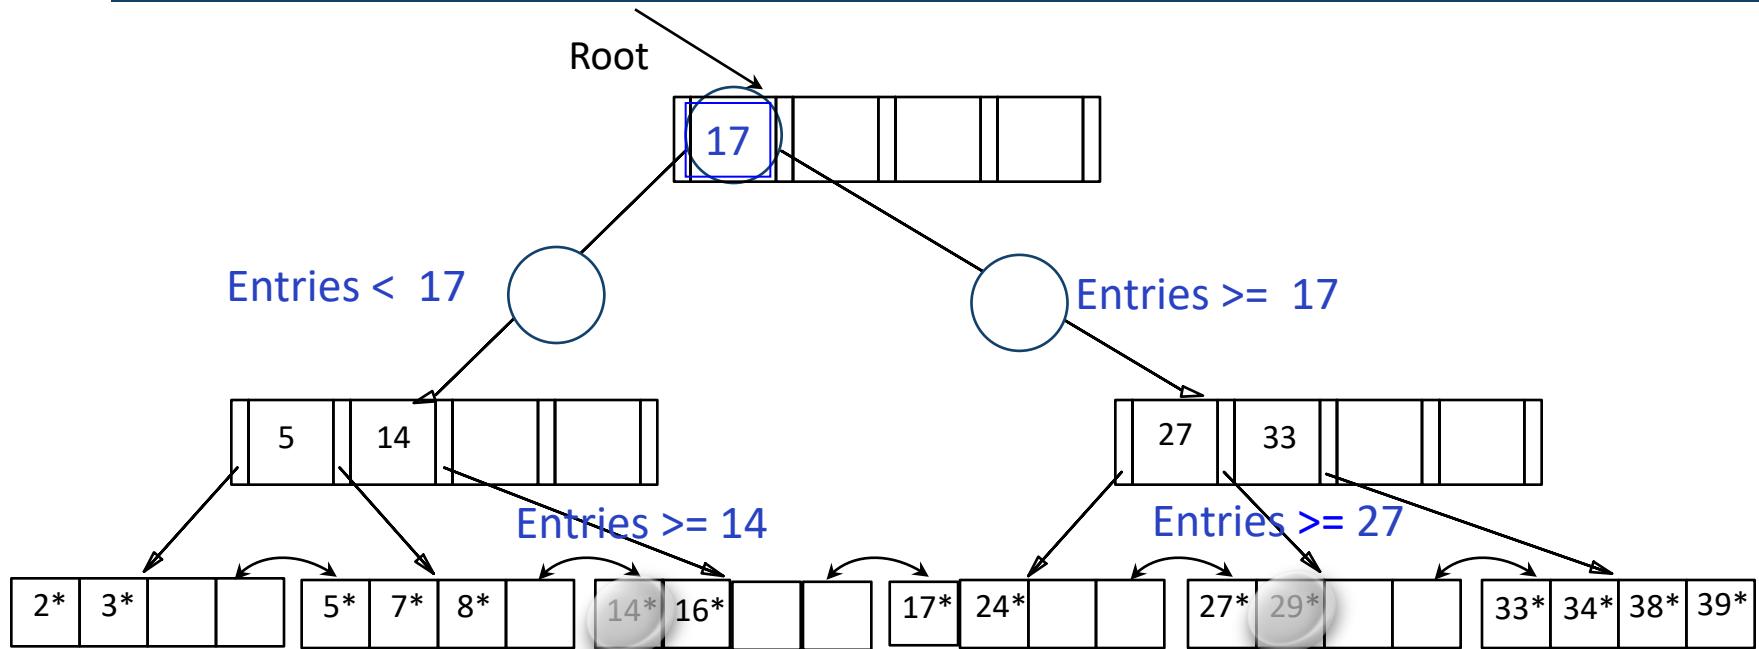


Step 4: Build B+ Tree Index

Noting that the average number of index entries in an index page is $480 \times 75\% = 360$

- › The next-level index would then consist of 390 index entries, and $390/360 = 2$ pages
 - The root level consists of 2 index entries, and $2/360 = 1$ page
- › Therefore, the **total number of index pages is** $390 + 2 + 1 = 393$ index pages
- › $393/140,351 = 0.2\%$ increase



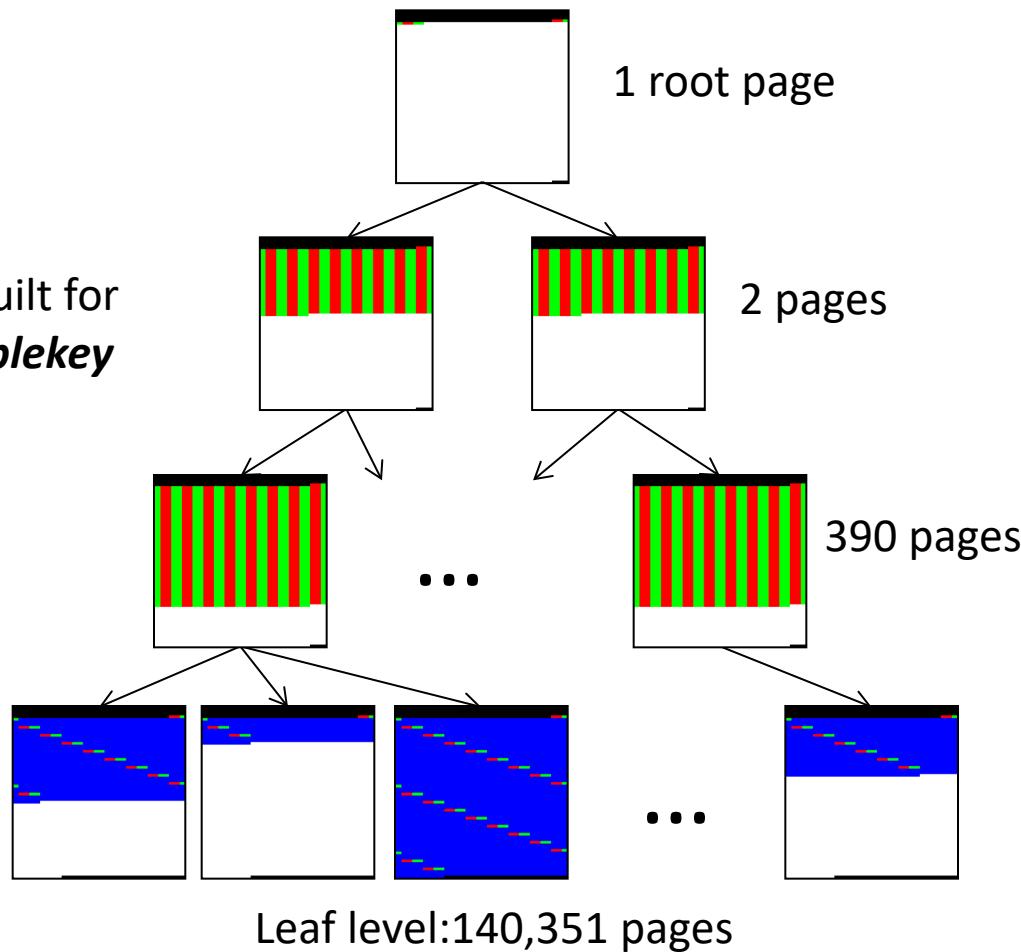


- › Note how **record/data pages in the leafs** are sorted and linked
- › **Find 129**
- › The cost of searching in a B+ tree index: *the lookup cost + number of pages containing retrieved records*
 - **Lookup cost** is equal to ***the number of levels of the B+ tree + data pages:*** In B+ tree, all paths from the root to a leaf are of the same length

Example: Cost of Equality Search in a B+ Tree Index

Example: **SELECT * FROM Relation WHERE tuplekey=715;**

The B+ tree built for
search key **tuplekey**



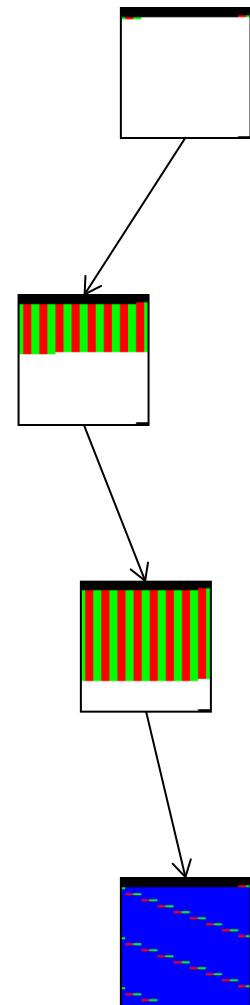
Example: Cost of Equality Search in a B+ Tree Index

SELECT * FROM Relation WHERE tuplekey=715;

› For **equality search** on tuplekey, we can use the related index:

- 1) Load index root into main memory (cost 1 I/O);
- 2) Find location of matching page in next level;
- 3) Load matching next level page (cost 1 I/O);
- 4) Find location of matching page in following level;
- 5) Load matching page on following level (cost 1 I/O);
- 6) Find location of matching record page in leaf level;
- 7) Load matching record page (cost 1 I/O);
- 8) Check each record in the record page for match.

› Total of **4 I/Os** vs **70,176 I/Os** for heap file



SELECT * FROM Relation WHERE attribute1 BETWEEN 100 AND 119;

› For **range search** on attribute1, the index above is of no use. Why? We still need to check each record page sequentially, so **140,351 I/Os**: Alternative: build another index on attribute1.

› Physical Data Organisation: how is data physically stored in DBMS?

› Access Paths: how are records retrieved from DBMS?

- Access methods for heap files (linear scan)
- Access methods for sorted files (binary search)
- Access methods for indexes (index scan)

› B+ Tree Index

- Primary index
- Composite search keys

Composite Search Key: Field Order Matters

› **Composite search key:** case where the **search key** contains **several fields**

- (age, salary): first (logically) sort on age, then among records with the same age, (logically) sort on salary
- (salary, age): first (logically) sort on salary, then among records with the same salary, (logically) sort on age
- “Index with search key (age, salary)” \neq “Index with search key (salary, age) ”

	age	salary
r1	22	2000
r2	21	2000
r3	23	2000
r4	22	3000
r5	22	1000

Heap file order

	age	salary
r2	21	2000
r5	22	1000
r1	22	2000
r4	22	3000
r3	23	2000

Sorted order on search key
(age, salary)

	age	salary
r5	22	1000
r2	21	2000
r1	22	2000
r3	23	2000
r4	22	3000

Sorted order on search key
(salary, age)

› Given an **index** with search key (**age, salary**)

- Can we use the index to lookup search condition “age = 22”? **YES!**
- Can we use the index to lookup search condition “age >= 22”? **YES!**
- Can we use the index to lookup search condition “salary = 2000”? **NO!**
- Can we use the index to lookup search condition “salary >= 2000”? **NO!**
- Can we use the index to lookup search condition “age = 22 AND salary = 2000”? **YES!**
- Can we use the index to lookup search condition “age = 22 AND salary >= 2000”? **YES!**

	age	salary
r2	21	2000
r5	22	1000
r1	22	2000
r4	22	3000
r3	23	2000

Sorted order on search key
(age, salary)



Important indexing rule: An index with a search key can be used to lookup a search condition if all records satisfying the search condition are **consecutive** in the sorted order of the table’s records based on that search key

- › Kifer/Bernstein/Lewis (2nd edition)
 - Chapter 9 (9.1-9.4)
 - *Kifer/Bernstein/Lewis gives a good overview of indexing*
- › Ramakrishnan/Gehrke (3rd edition)
 - Chapter 8
 - *The Ramakrishnan/Gehrke is very technical on this topic, providing a lot of insight into how disk-based indexes are implemented.*
- › Ullman/Widom (3rd edition - '1st Course in Databases')
 - Chapter 8 (8.3 onwards)
 - *Mostly overview, with simple cost model of indexing*
- › Silberschatz/Korth/Sudarshan (5th ed)
 - Chapter 11 and 12

Next Week: Query Processing and Evaluation

- › Ramakrishnan/Gehrke – Chapters 13 and 14
- › Kifer/Bernstein/Lewis – Chapter 10
- › Garcia-Molina/Ullman/Widom – Chapter 15

See you next week!



THE UNIVERSITY OF
SYDNEY