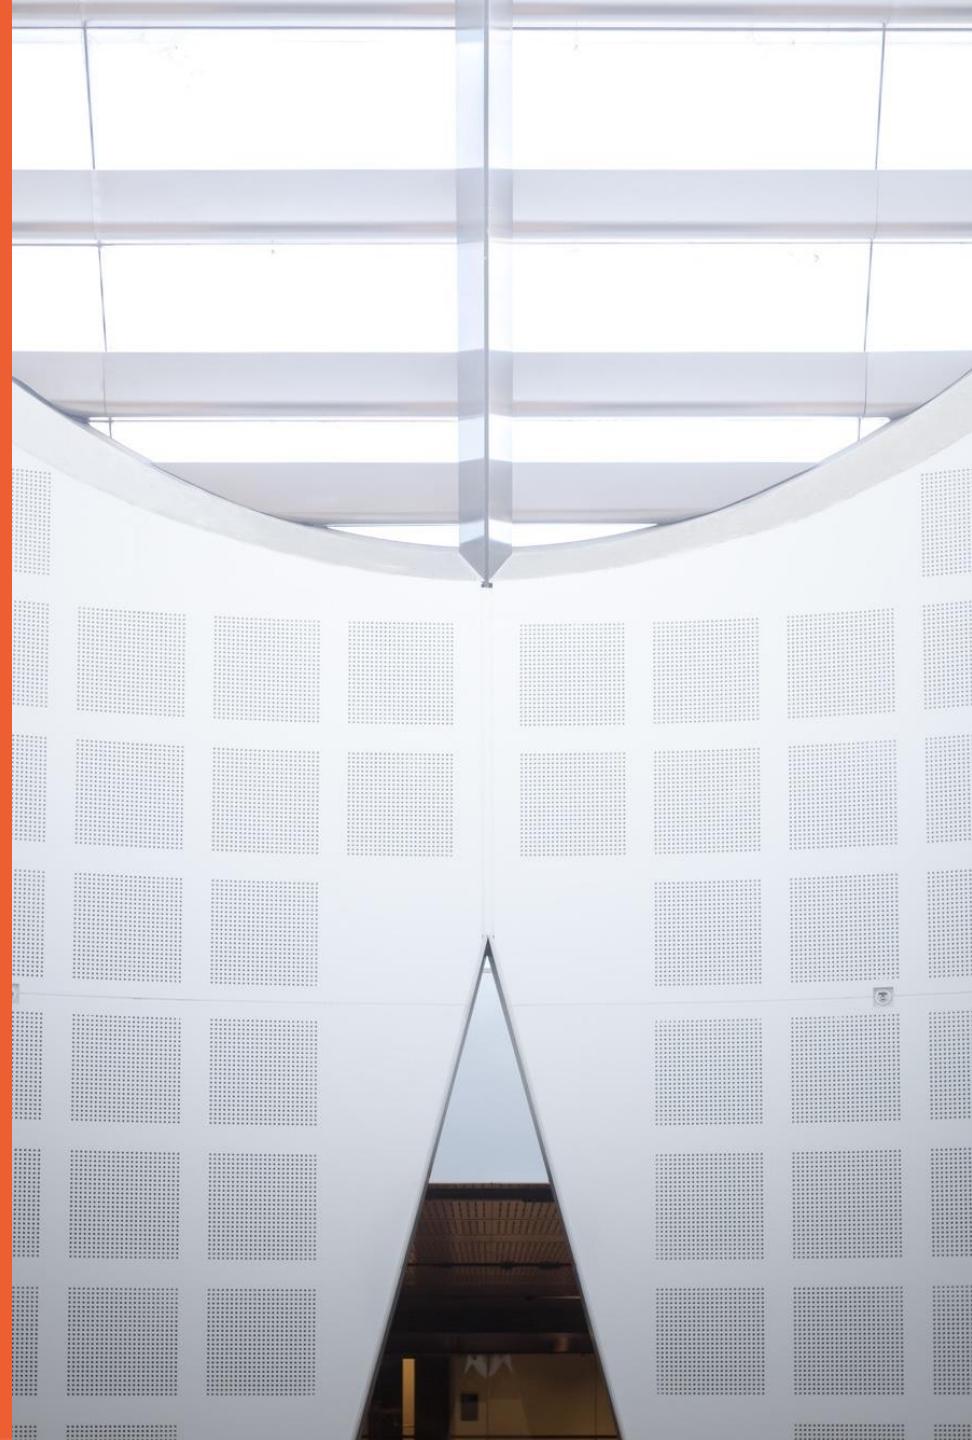


COMP9121: Design of Networks and Distributed Systems

Week 10: Application layer 2
Network Security
Wei Bao
School of Computer Science



THE UNIVERSITY OF
SYDNEY

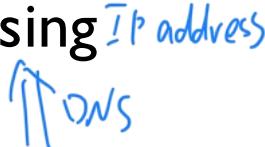


DNS

DNS: domain name system

Internet hosts, routers:

- IP address (32 bit) - used for addressing datagrams
- “name”, e.g., www.yahoo.com - used by humans



people: many identifiers:

- name, passport #

Q: how to map between IP address and name, and vice versa ?

Domain Name System:

- *distributed database* implemented in hierarchy of many *name servers*
- *application-layer protocol*: hosts, name servers communicate to *resolve* names (address/name translation)

DNS: services, structure

DNS services

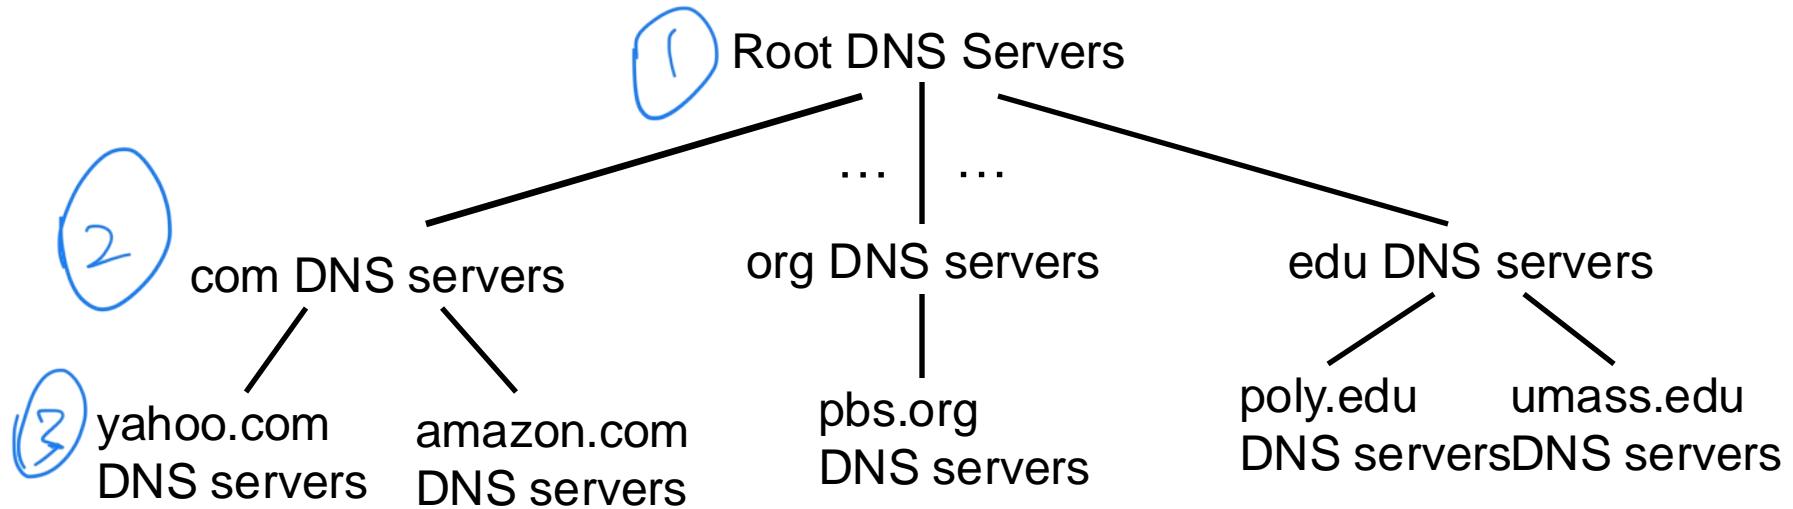
- hostname to IP address translation, Function
- host aliasing
 - canonical, alias names
- mail server aliasing
- load distribution
 - replicated Web servers:
 - { many IP addresses
 - correspond to one name

→ 一个名字可以对应很多个
server, 可以动态分配不同的 server

why not centralize DNS?

- ① single point of failure
- ② distant centralized database
- ③ scalability

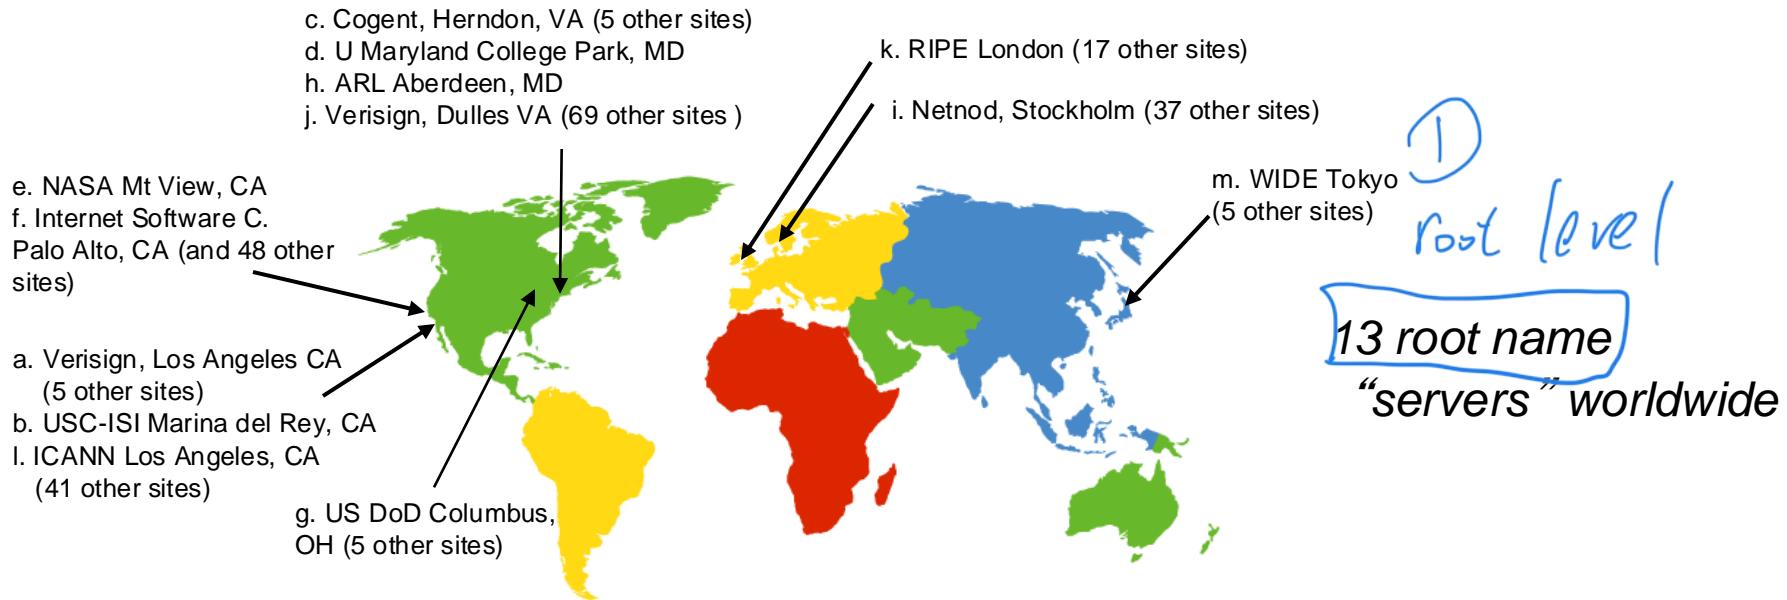
DNS: a distributed, hierarchical database



client wants IP for www.amazon.com; 1st approx:

- client queries root server to find com DNS server
- client queries .com DNS server to get amazon.com DNS server
- client queries amazon.com DNS server to get IP address for www.amazon.com

DNS: root name servers



TLD, authoritative servers

2

top-level domain (TLD) servers: .com, .ca, .AU

- responsible for com, org, net, edu, aero, jobs, museums, and all top-level country domains, e.g.: uk, fr, ca, jp
- Network Solutions maintains servers for .com TLD
- Educause for .edu TLD

3

authoritative DNS servers: Amazon.com

- organization's own DNS server(s), providing authoritative hostname to IP mappings for organization's named hosts
- can be maintained by organization or service provider

④

Local DNS name server

A. 如果面对都是 client 的话

Author DNS

B. 如果面对有 client + provider 的话

缓存

- does not strictly belong to hierarchy
- each ISP (residential ISP, company, university) has one
 - also called “default name server”
- when host makes DNS query, query is sent to its local DNS server
 - has local cache of recent name-to-address translation pairs (but may be out of date!)
 - acts as proxy, forwards query into hierarchy

Better, 因为☆承担了主要压力

DNS name resolution example

Method 1

- host at cis.poly.edu wants IP address for gaia.cs.umass.edu

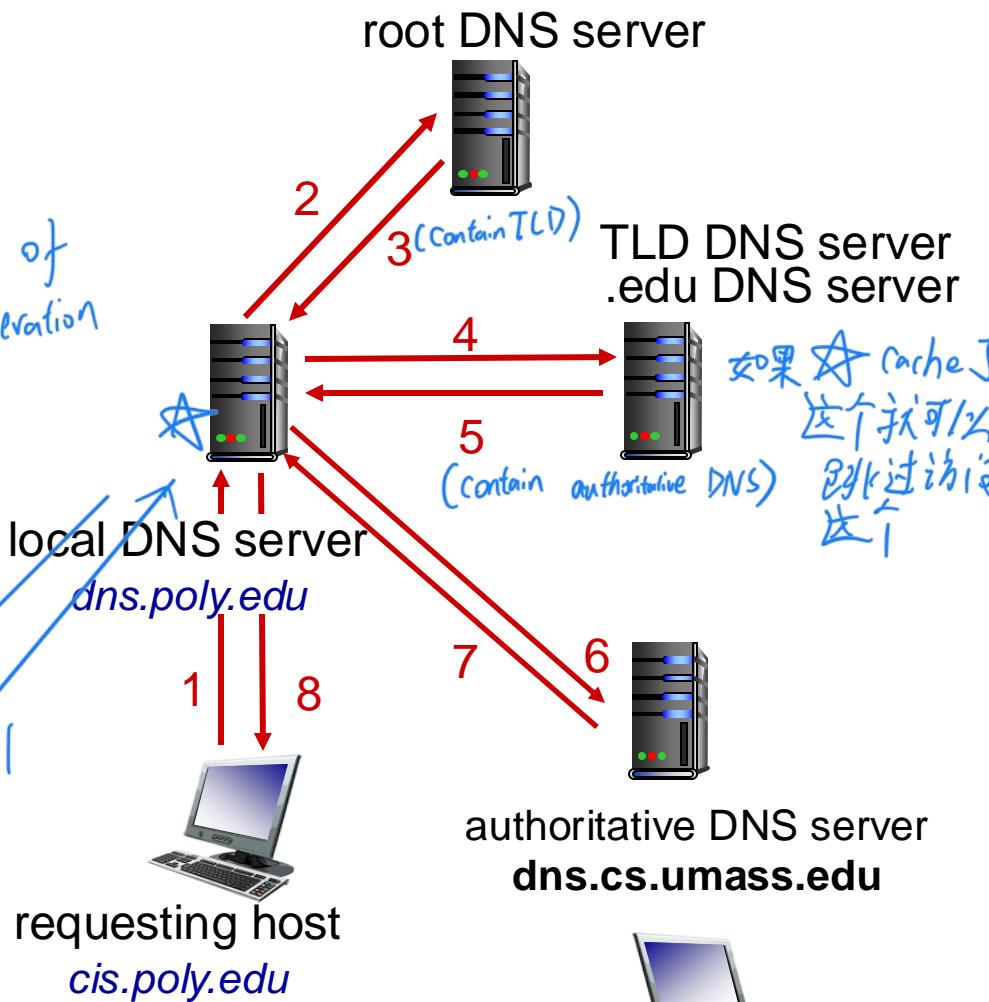
iterated query:

- ❖ contacted server replies with name of server to contact
- ❖ “I don't know this name, but ask this server”

如果这样
another user

访问同一IP, 只用traverse

performance of
this operation



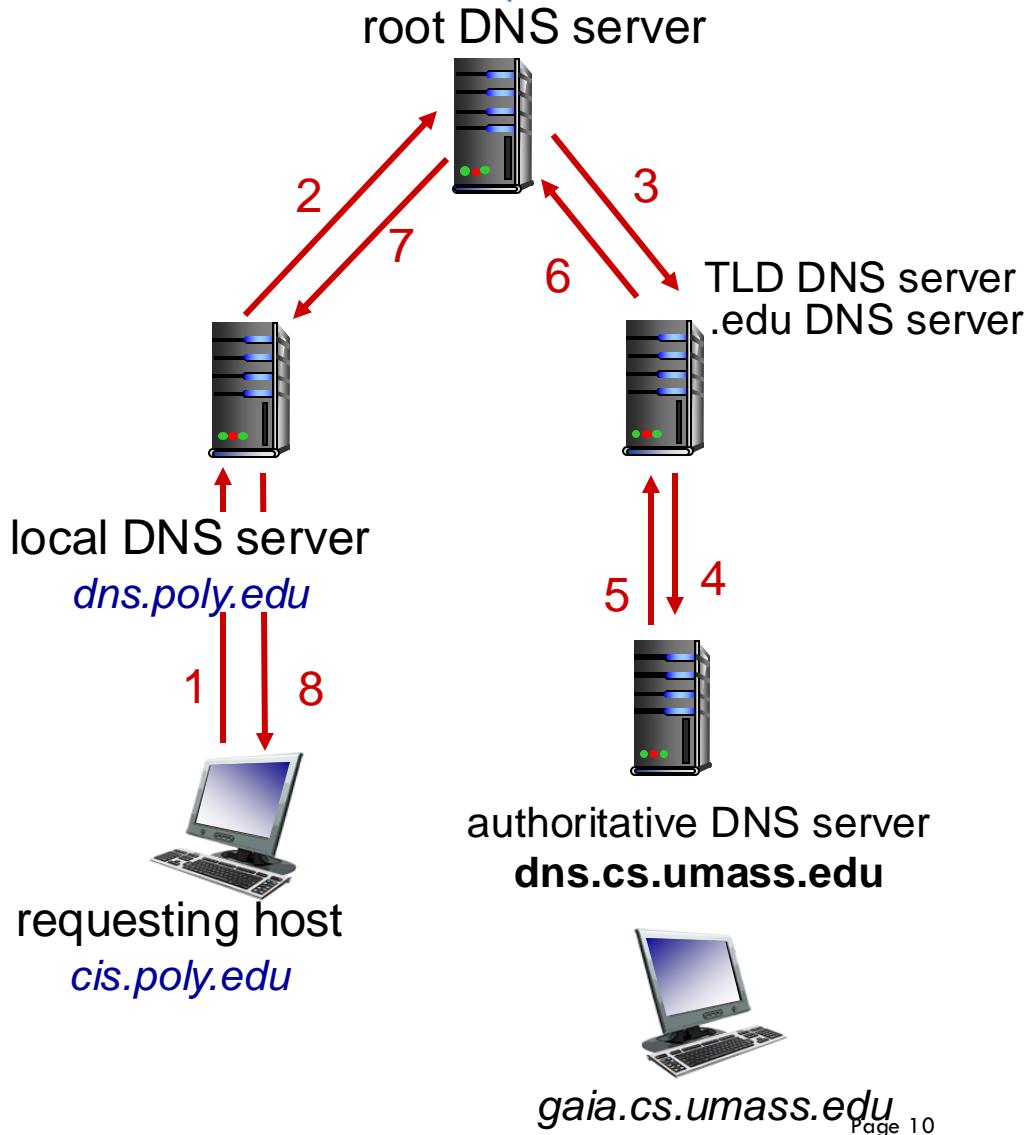
承担主要压力，不希望
入这样

DNS name resolution example (cont'd)

Method 2

recursive query:

- ❖ puts burden of name resolution on contacted name server
- ❖ heavy load at upper levels of hierarchy?



DNS caching, updating records

Prefer UPP, because we want to be fast

- once (any) name server learns mapping, it caches mapping
 - cache entries timeout (disappear) after some time (TTL)
- cached entries may be *out-of-date* (best effort name-to-address translation!)
 - if name host changes IP address, may not be known Internet-wide until all TTLs expire
- update/notify mechanisms proposed IETF standard
 - RFC 2136

DHCP 可以帮助 client 知道
local DNS

DNS records

DNS: distributed db storing resource records (RR)

RR format: `(name, value, type, ttl)`

type=A

- **name** is hostname
- **value** is IP address

type=NS

- **name** is domain (e.g.,
`foo.com`)
- **value** is hostname of
authoritative name server
for this domain

type=CNAME

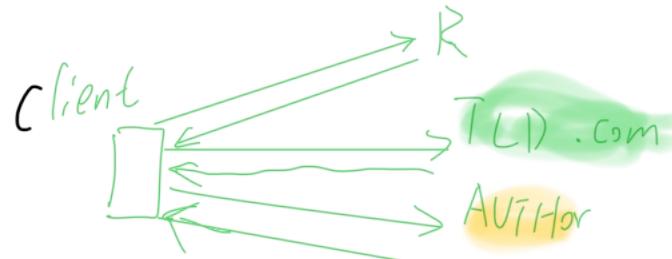
- **name** is alias name for some
“canonical” (the real) name
- `www.ibm.com` is really
`servereast.backup2.ibm.com`
- **value** is canonical name

type=MX

- **value** is name of mailserver
associated with **name**

Inserting records into DNS

- example: new startup “Network Utopia”
- register name networkuptopia.com at *DNS registrar* (e.g., Network Solutions)
 - provide names, IP addresses of authoritative name server
 - registrar inserts two RRs into .com TLD server:
(networkutopia.com, dns1.networkutopia.com, NS)
(dns1.networkutopia.com, 212.212.212.1, A)
 - create at authoritative server
type A record for www.networkutopia.com;
(www.networkutopia.com, 212.212.212.22, A)
(www.home.networkutopia.com, www.networkutopia.com, CNAME)



Peer-to-Peer

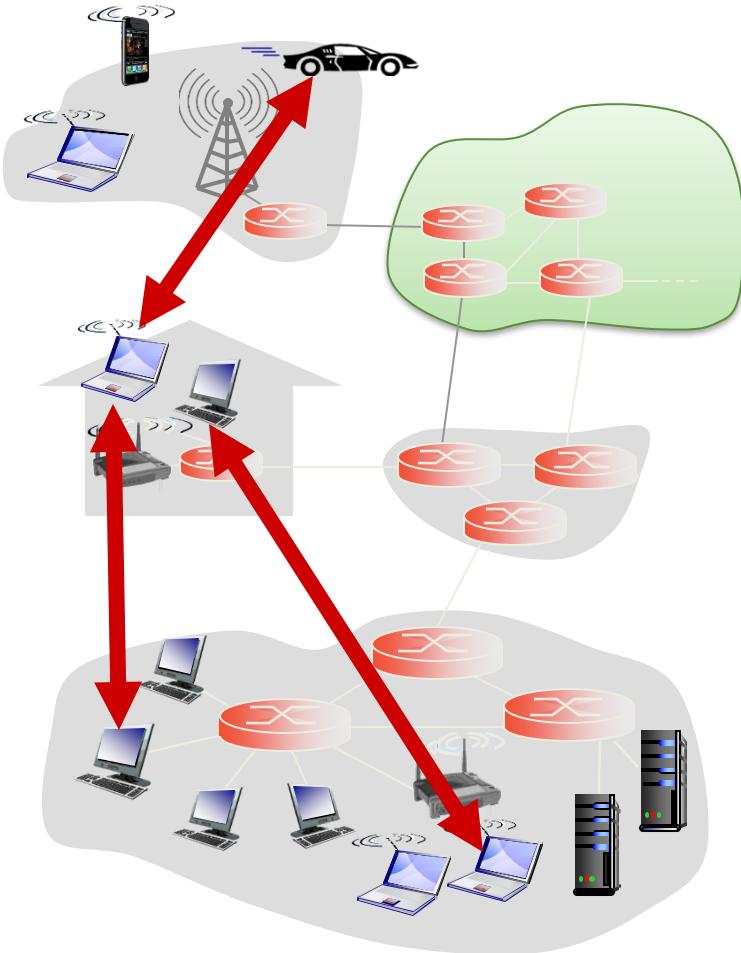
P2P

Pure peer-to-peer model architecture

- no always-on server
- arbitrary end systems directly communicate
- peers are intermittently connected and change IP addresses

examples:

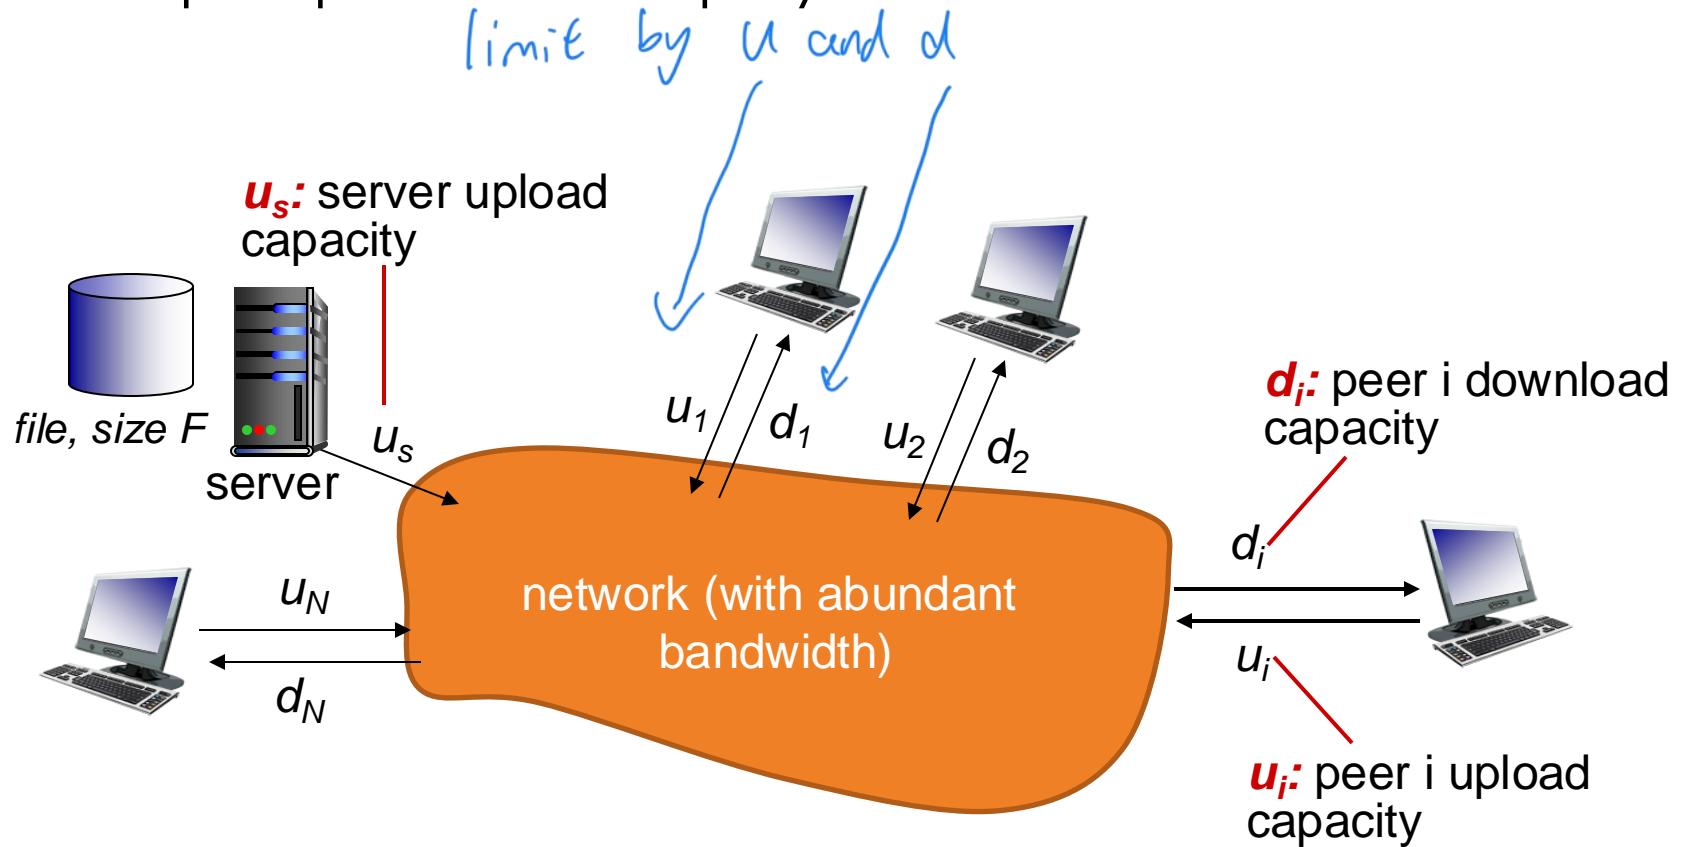
- file distribution (BitTorrent)
- Streaming (Zattoo, KanKan)
- VoIP (Skype)



File distribution: client-server vs. p2p

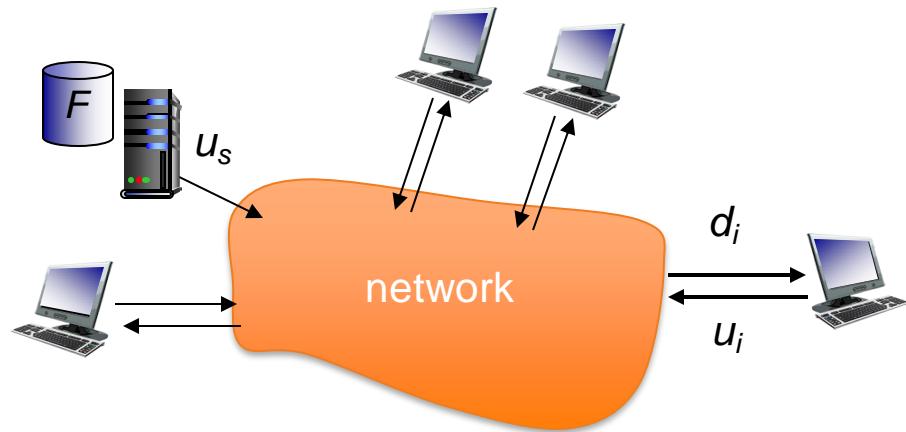
Question: how much time to distribute file (size F) from one server to N peers?

- peer upload/download capacity is limited resource



File distribution time: client-server

- **server transmission:** must sequentially send (upload) N file copies:
 - time to send one copy: F/u_s
 - time to send N copies: NF/u_s
- ❖ **client:** each client must download file copy
 - d_{\min} = min client download rate
 - (worst case) client download time: F/d_{\min}



time to distribute F to N clients using client-server approach

$$D_{c-s} \geq \max\{NF/u_s, F/d_{\min}\}$$

Distribution Time

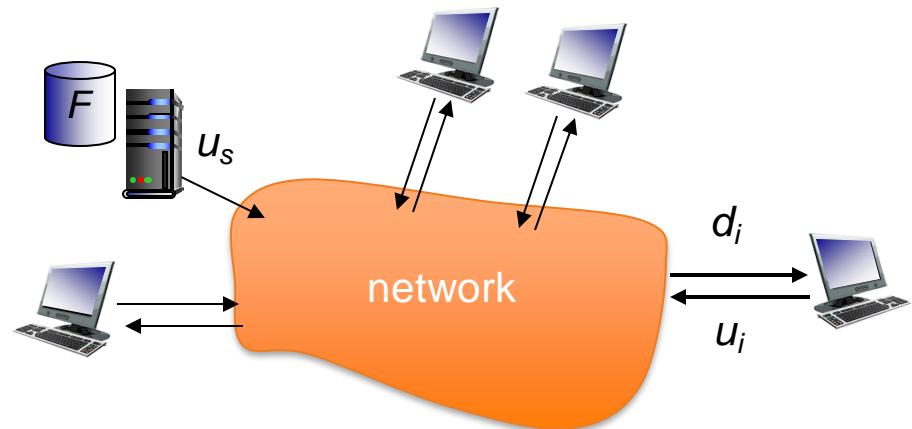
increases linearly in N

向上坡度

更多用户导致 Distributed Time 增加

File distribution time: p2p

- **server transmission:** must upload at least one copy
 - time to send one copy: F/u_s
- ❖ **client:** each client must download file copy
 - client download time: F/d_{\min}
- ❖ **clients:** as aggregate must download NF bits = upload NF bits
 - Max upload rate $u_s + \sum u_i$
 - $NF/(u_s + \sum u_i)$



time to distribute F to N clients using P2P approach

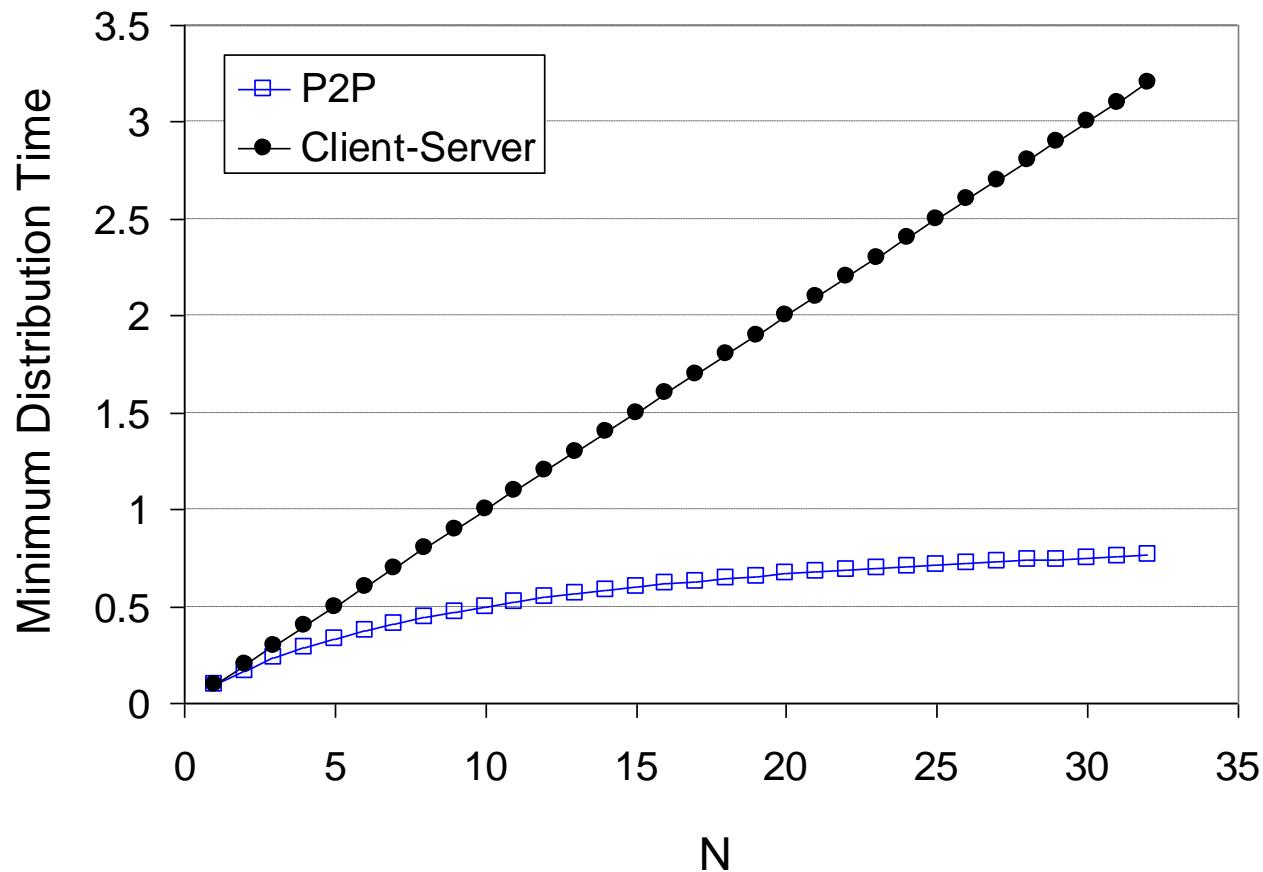
$$D_{P2P} \geq \max\{F/u_s, F/d_{\min}, NF/(u_s + \sum u_i)\}$$

increases linearly in N ...

... but so does this, as each peer brings service capacity $Download$

Client-server vs. p2p

client upload rate = u , $F/u = 1$ hour, $u_s = 10u$, $d_{min} \geq u_s$



P2P file distribution: BitTorrent

BitTorrent, a file sharing application

- 20% of European internet traffic in 2012.
- Used for Linux distribution, software patches, distributing i
- Goal: quickly replicate large files to large number of clients
- Web server hosts a .torrent file (w/ file length, hash, tracker's URL...)
- A tracker tracks downloaders/owners of a file
- Files are divided into chunks (256kB-1MB)
- Downloaders download chunks from themselves (and owners)
- Tit-for-tat: the more one shares (server), the faster it can download (client)



上傳更多, 下載更快

P2P file distribution: BitTorrent

- › file divided into 256KB chunks
- › peers in torrent send/receive file chunks



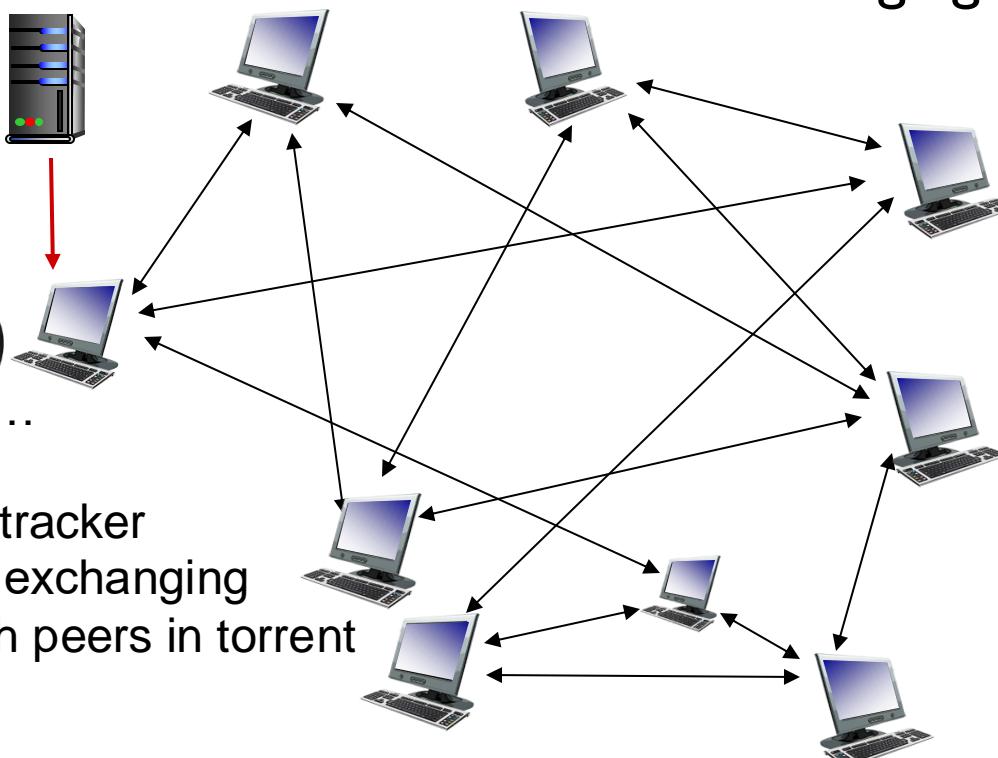
tracker: tracks peers
participating in torrent

帮助别人
加入 Torrent
网络



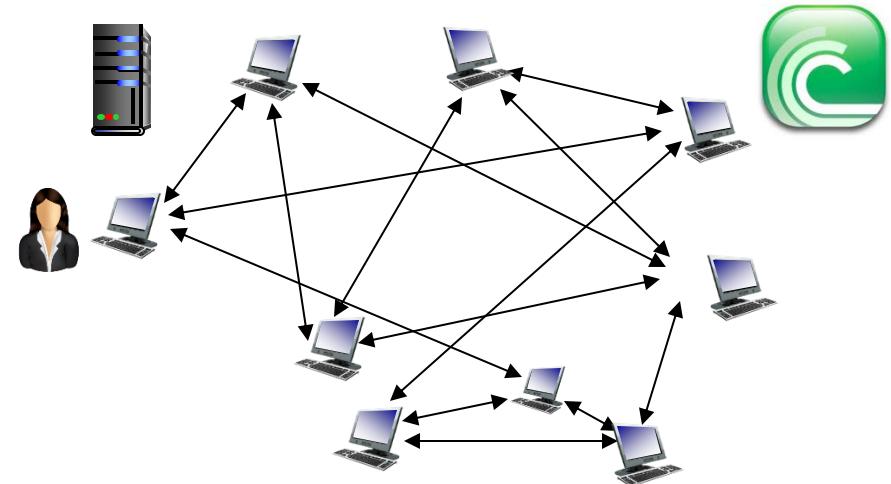
Alice arrives ...
... obtains list
of peers from tracker
... and begins exchanging
file chunks with peers in torrent

torrent: group of peers
exchanging chunks of a file



P2P file distribution: BitTorrent

- peer joining torrent:
 - has no chunks, but will accumulate them over time from other peers
 - registers with tracker to get list of peers, connects to subset of peers (“neighbors”)



- › while downloading, peer uploads chunks to other peers
- › peer may change peers with whom it exchanges chunks
- › **churn:** peers may come and go
- › once peer has entire file, it may (selfishly) leave or (altruistically) remain in torrent

BitTorrent: requesting, sending file chunks

requesting chunks:

- at any given time, different peers have different subsets of file chunks
- periodically, Alice asks each peer for list of chunks that they have
- Alice requests missing chunks from peers, rarest first

sending chunks: tit-for-tat

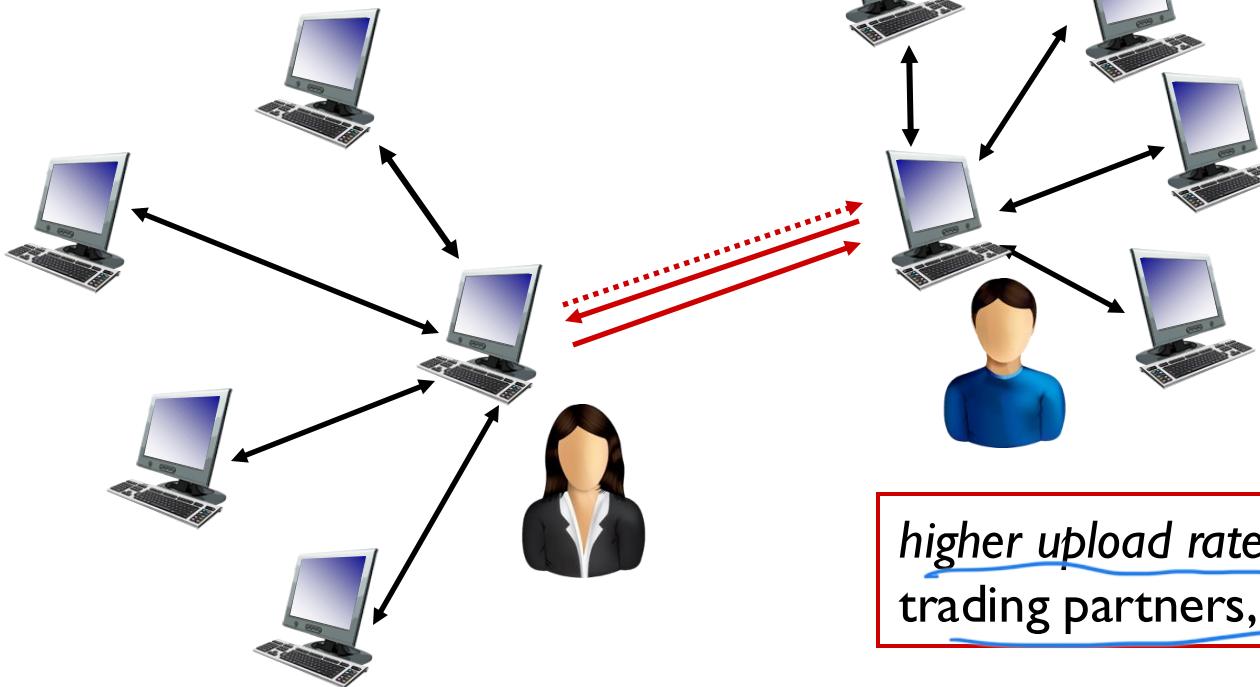


- › Alice sends chunks to those four peers currently sending her *Examp)b* chunks *at highest rate*
 - ① other peers are choked by Alice (do not receive chunks from her)
 - › re-evaluate top 4 every 10 secs
- › every 30 secs: randomly select another peer, starts sending chunks
 - ② “optimistically unchoke” this peer newly chosen peer may join top 4

BitTorrent: tit-for-tat

randomly chose Bob

- (1) Alice “optimistically unchoke” Bob
- (2) Alice becomes one of Bob’s top-four providers; Bob reciprocates
- (3) Bob becomes one of Alice’s top-four providers

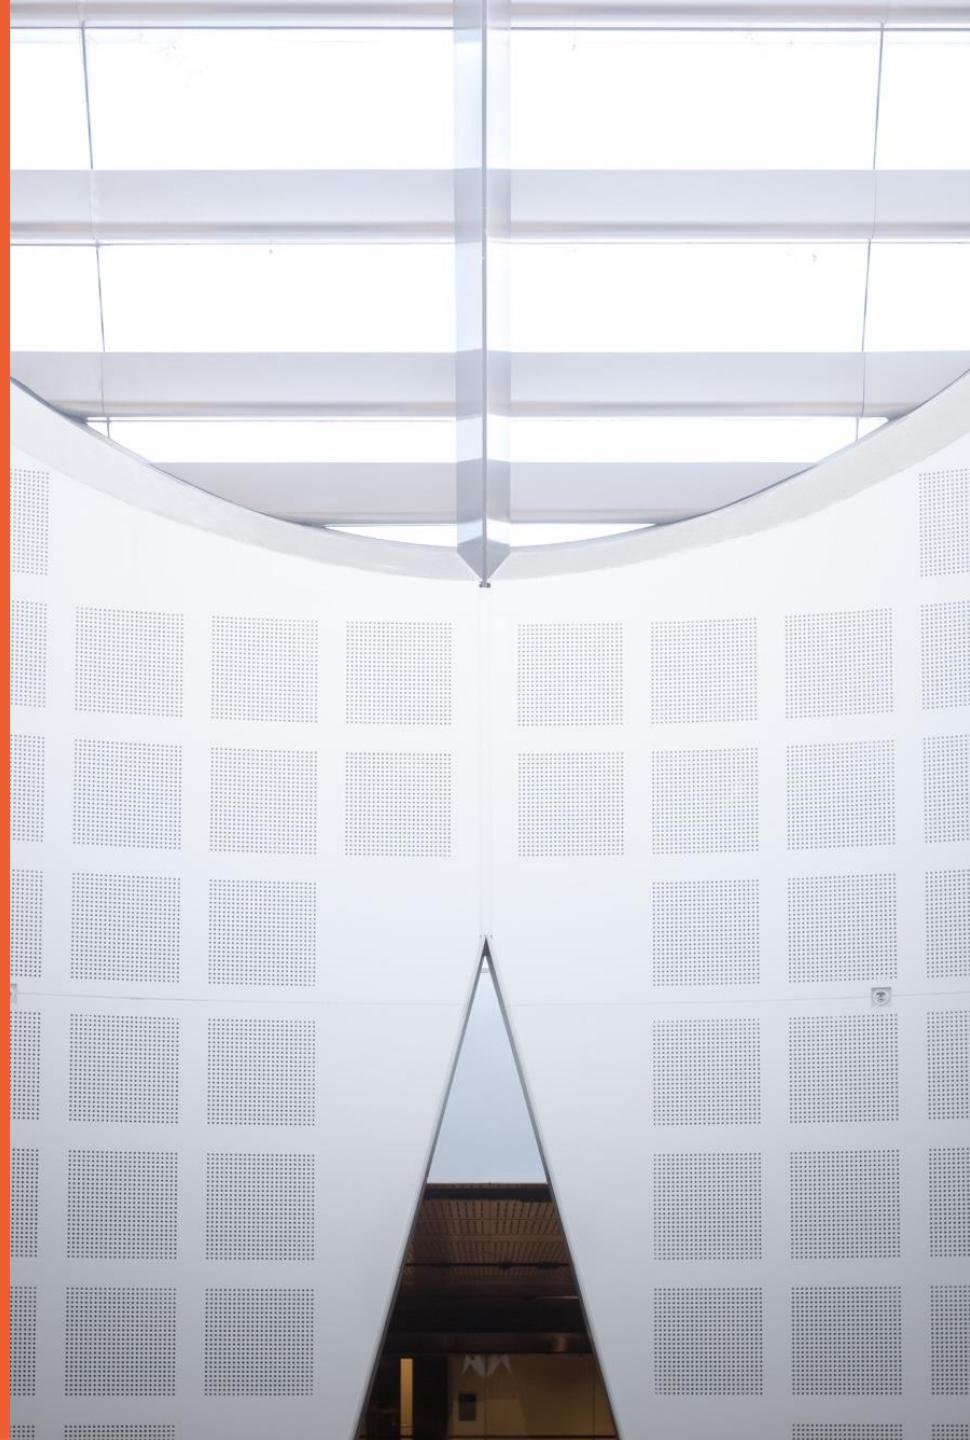


higher upload rate: find better
trading partners, get file faster !

Network Security



THE UNIVERSITY OF
SYDNEY



Network Security

Chapter goals:

- understand principles of network security:
 - cryptography and its *many* uses beyond “confidentiality”
 - authentication
 - message integrity
- security in practice:
 - firewalls and intrusion detection systems
 - security in application, transport, network, link layers

What is network security?

confidentiality: only sender, intended receiver should
“understand” message contents

- sender encrypts message
- receiver decrypts message

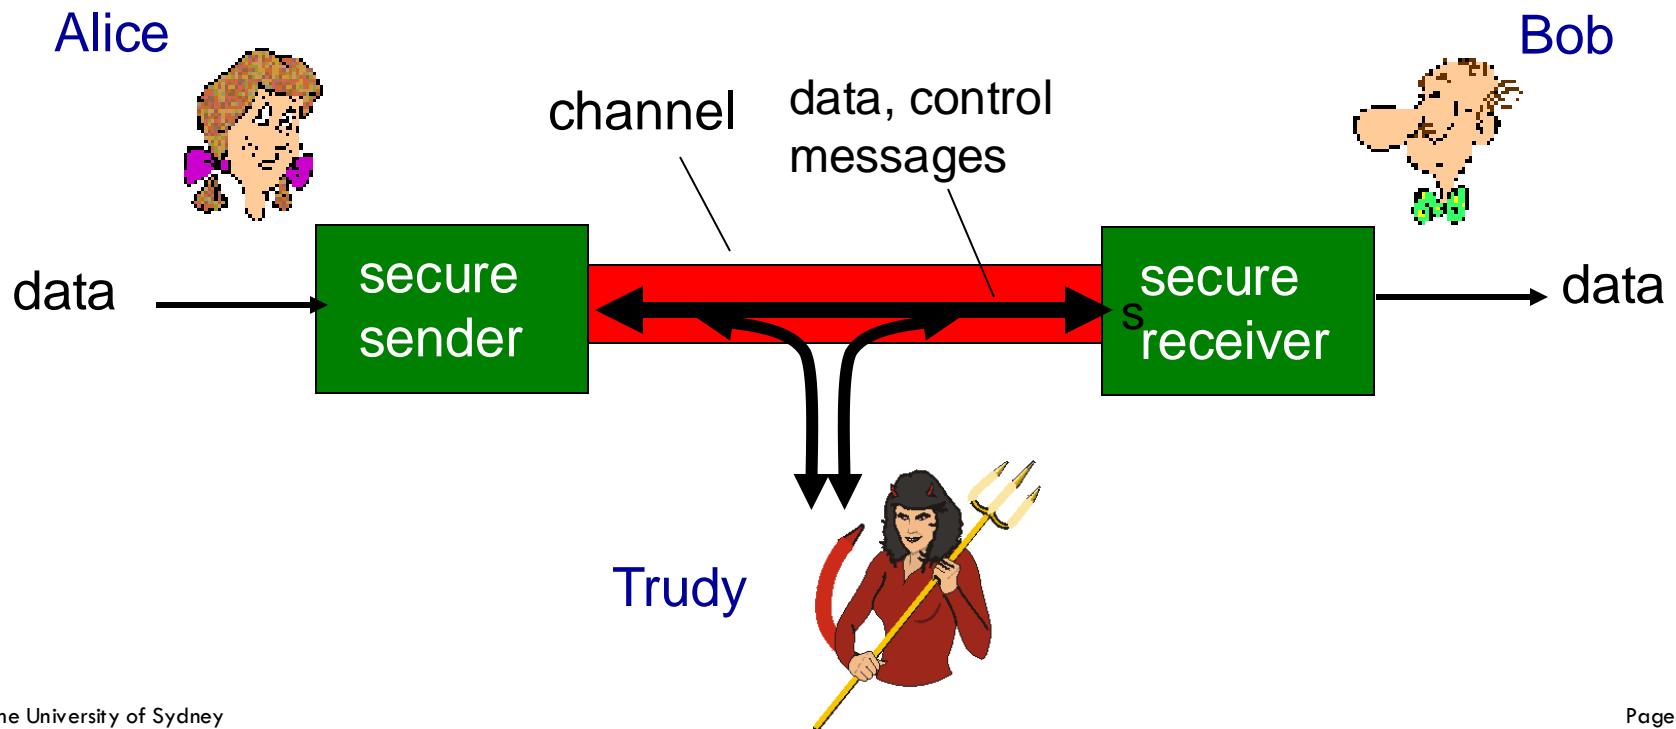
authentication: sender, receiver want to confirm identity of
each other

message integrity: sender, receiver want to ensure message
not altered (in transit, or afterwards) without detection

access and availability: services must be accessible and
available to users

Friends and enemies: Alice, Bob, Trudy

- well-known in network security world
- Bob, Alice (lovers!) want to communicate “securely”
- Trudy (intruder) may intercept, delete, add messages



Who might Bob, Alice be?

- ... well, *real-life* Bobs and Alices!
- Web browser/server for electronic transactions (e.g., on-line purchases)
- on-line banking client/server
- DNS servers
- routers exchanging routing table updates

There are bad guys (and girls) out there!

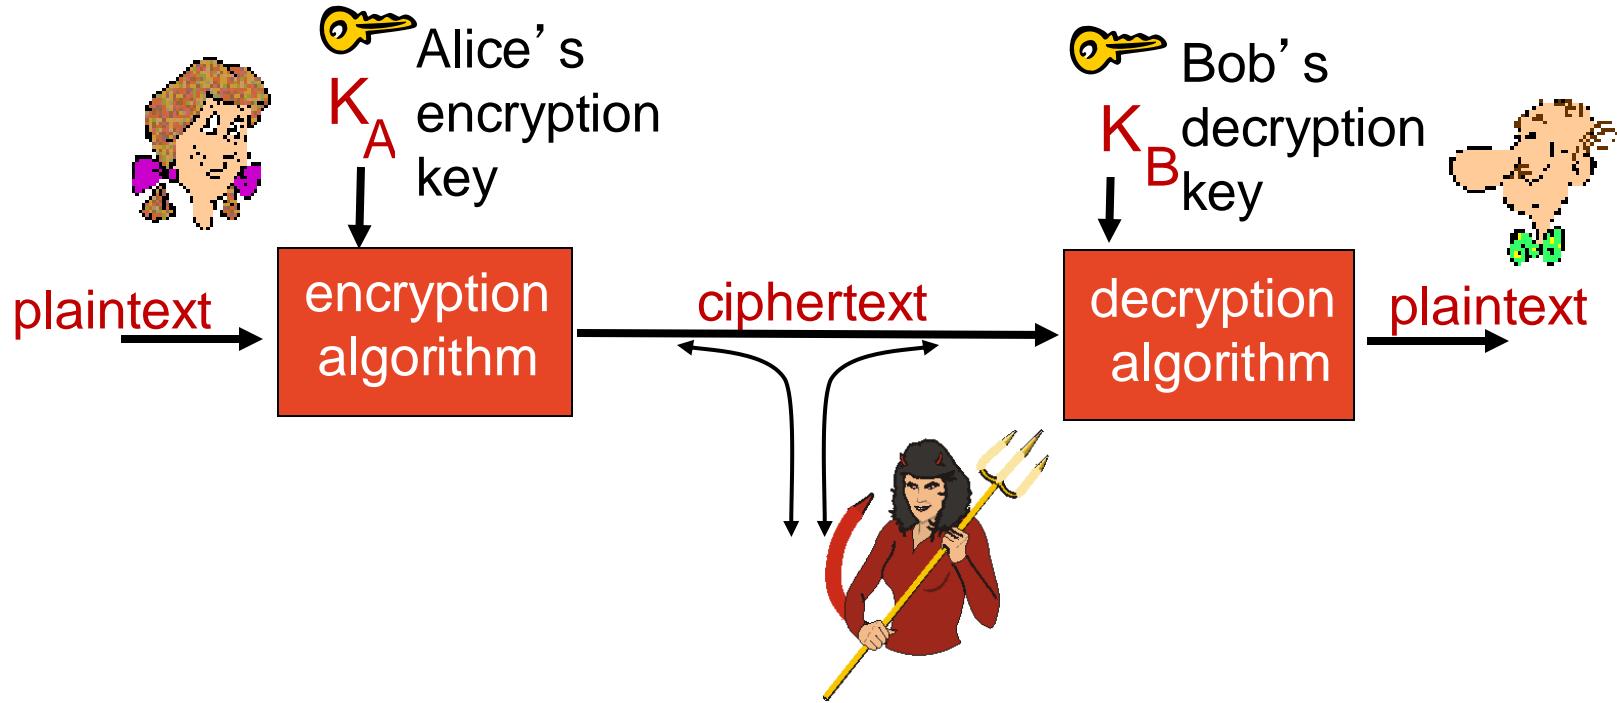
Q: What can a “bad guy” do?

A: A lot!

- **eavesdrop**: intercept messages
- actively **insert** messages into connection
- **impersonation**: can fake (spoof) source address in packet (or any field in packet) *钓鱼 web*
- **hijacking**: “take over” ongoing connection by removing sender or receiver, inserting himself in place
- **denial of service**: prevent service from being used by others (e.g., by overloading resources)

Principles of cryptography

The language of cryptography



m plaintext message

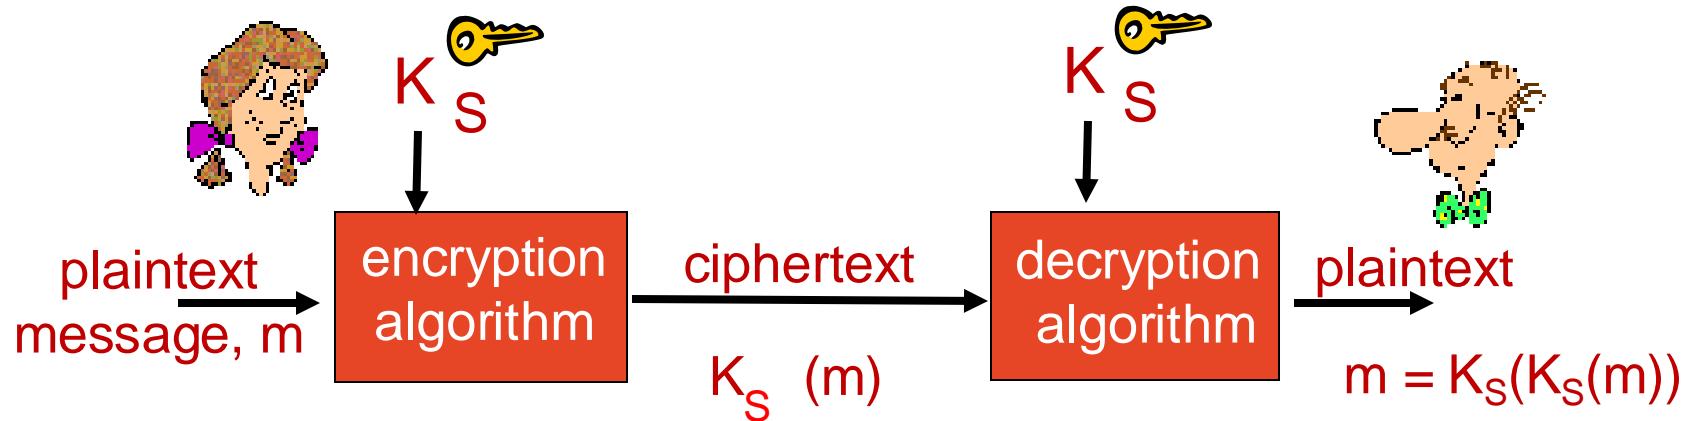
$K_A(m)$ ciphertext, encrypted with key K_A

$m = K_B(K_A(m))$

Breaking an encryption scheme

- **cipher-text only attack:**
Trudy has ciphertext she can analyze
- **two approaches:**
 - brute force: search through all keys
 - statistical analysis
- **known-plaintext attack:**
Trudy has plaintext corresponding to ciphertext
 - e.g., in monoalphabetic cipher, Trudy determines pairings for a,l,i,c,e,b,o,
- **chosen-plaintext attack:**
Trudy can get ciphertext for chosen plaintext

Symmetric key cryptography



symmetric key crypto: Bob and Alice share same (symmetric) key: K

- e.g., key is knowing substitution pattern in mono alphabetic substitution cipher

Q: how do Bob and Alice agree on key value?

Simple encryption scheme

substitution cipher: substituting one thing for another

- monoalphabetic cipher: substitute one letter for another

plaintext: abcdefghijklmnopqrstuvwxyz
ciphertext: mnbvvcxzdasdfghjklpoiuytrewq

The diagram illustrates a monoalphabetic substitution cipher. It shows two rows of letters. The top row is labeled "plaintext" and contains the letters "abcdefghijklmnopqrstuvwxyz". The bottom row is labeled "ciphertext" and contains the letters "mnbvvcxzdasdfghjklpoiuytrewq". Red arrows point vertically from each letter in the plaintext to its corresponding letter in the ciphertext, demonstrating the mapping rule.

e.g.: Plaintext: bob. i love you. alice

ciphertext: nkn. s gktc wky. mgsbc



Encryption key: mapping from set of 26 letters
to set of 26 letters

A more sophisticated encryption approach

- n substitution ciphers, M_1, M_2, \dots, M_n
- cycling pattern:
 - e.g., $n=4: M_1, M_3, M_4, M_3, M_2; M_1, M_3, M_4, M_3, M_2; \dots$
- for each new plaintext symbol, use subsequent substitution pattern in cyclic pattern
 - dog: d from M_1 , o from M_3 , g from M_4



Encryption key: n substitution ciphers, and cyclic pattern

- key need not be just n-bit pattern

A more sophisticated encryption approach

plaintext: abcdefghijklmnopqrstuvwxyz

M1 : asdfghmnbvcxzjklpoiuytrewq

M2 : oiuymnbvcxzasdfghjklptrewq

M3 : jklpoiuytrewqmnbvcxzasdfgh

M4 : iuybvcmnfghxzasdrewqjklpo



Symmetric key crypto: DES

DES: Data Encryption Standard

- US encryption standard [NIST 1993]
- 56-bit symmetric key, 64-bit plaintext input
- block cipher
- how secure is DES?
 - DES Challenge: 56-bit-key-encrypted phrase decrypted (brute force) in less than a day
 - no known good analytic attack
- making DES more secure:
 - 3DES: encrypt 3 times with 3 different keys

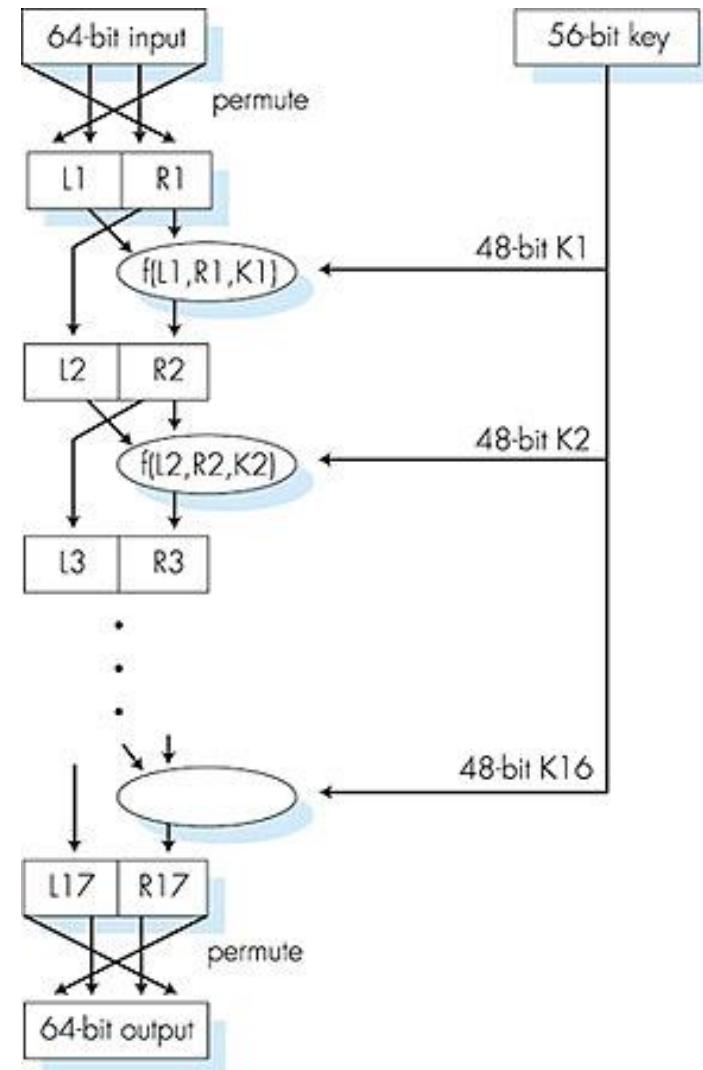
Symmetric key crypto: DES

DES operation

initial permutation

16 identical “rounds” of function application,
each using different 48 bits of key

final permutation



AES: Advanced Encryption Standard

- symmetric-key NIST standard, replaced DES (Nov 2001)
- processes data in 128 bit blocks
- 128, 192, or 256 bit keys
- brute force decryption (try each key) taking 1 sec on DES, takes 149 trillion years for AES

Public Key Cryptography

symmetric key crypto

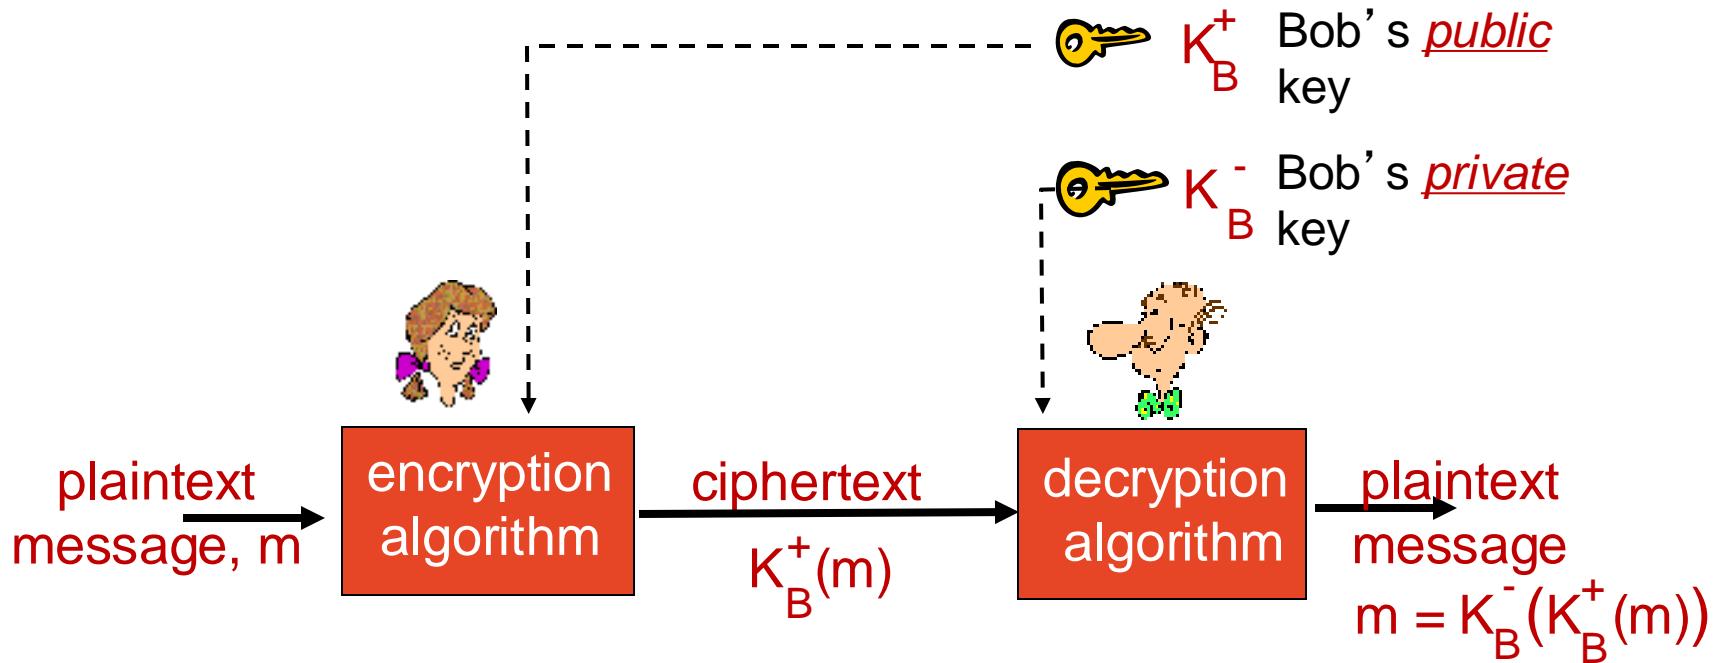
- requires sender, receiver know shared secret key
- Q: how to agree on key in first place (particularly if never “met”)?

public key crypto

- radically different approach [Diffie-Hellman76, RSA78]
- sender, receiver do *not* share secret key
- *public* encryption key known to *all*
- *private* decryption key known only to receiver



Public key cryptography



Public key encryption algorithms

requirements:

- ① need $K_B^+(\cdot)$ and $K_B^-(\cdot)$ such that

$$K_B^-(K_B^+(m)) = m$$

- ② given public key K_B^+ , it should be impossible to compute private key K_B^-

RSA: Rivest, Shamir, Adelson algorithm

Prerequisite: modular arithmetic

- $x \bmod n$ = remainder of x when divide by n
- facts:
 - $[(a \bmod n) + (b \bmod n)] \bmod n = (a+b) \bmod n$
 - $[(a \bmod n) - (b \bmod n)] \bmod n = (a-b) \bmod n$
 - $[(a \bmod n) * (b \bmod n)] \bmod n = (a*b) \bmod n$
- thus
 - $(a \bmod n)^d \bmod n = a^d \bmod n$
- example: $x=14$, $n=10$, $d=2$:
 - $(x \bmod n)^d \bmod n = 4^2 \bmod 10 = 6$
 - $x^d = 14^2 = 196 \quad x^d \bmod 10 = 6$

RSA: getting ready

- message: just a bit pattern
- bit pattern can be uniquely represented by an integer number
- thus, encrypting a message is equivalent to encrypting a number

example:

- $m = 10010001$. This message is uniquely represented by the decimal number 145.
- to encrypt m , we encrypt the corresponding number, which gives a new number (the ciphertext).

RSA: Creating public/private key pair

1. choose two large prime numbers p, q .
(e.g., 1024 bits each)
2. compute $n = pq$, $z = (p-1)(q-1)$
3. choose e (with $e < n$) that has no common factors with z (e, z are “relatively prime”).
4. choose d such that $ed - 1$ is exactly divisible by z .
(in other words: $ed \bmod z = 1$).
5. public key is $\underbrace{(n,e)}_{K_B^+}$. private key is $\underbrace{(n,d)}_{K_B^-}$.

RSA: encryption, decryption

0. given (n,e) and (n,d) as computed above

I. to encrypt message $m (< n)$, compute

$$c = m^e \bmod n$$

2. to decrypt received bit pattern, c , compute

$$m = c^d \bmod n$$

magic happens! $m = \underbrace{(m^e \bmod n)^d}_{c} \bmod n$

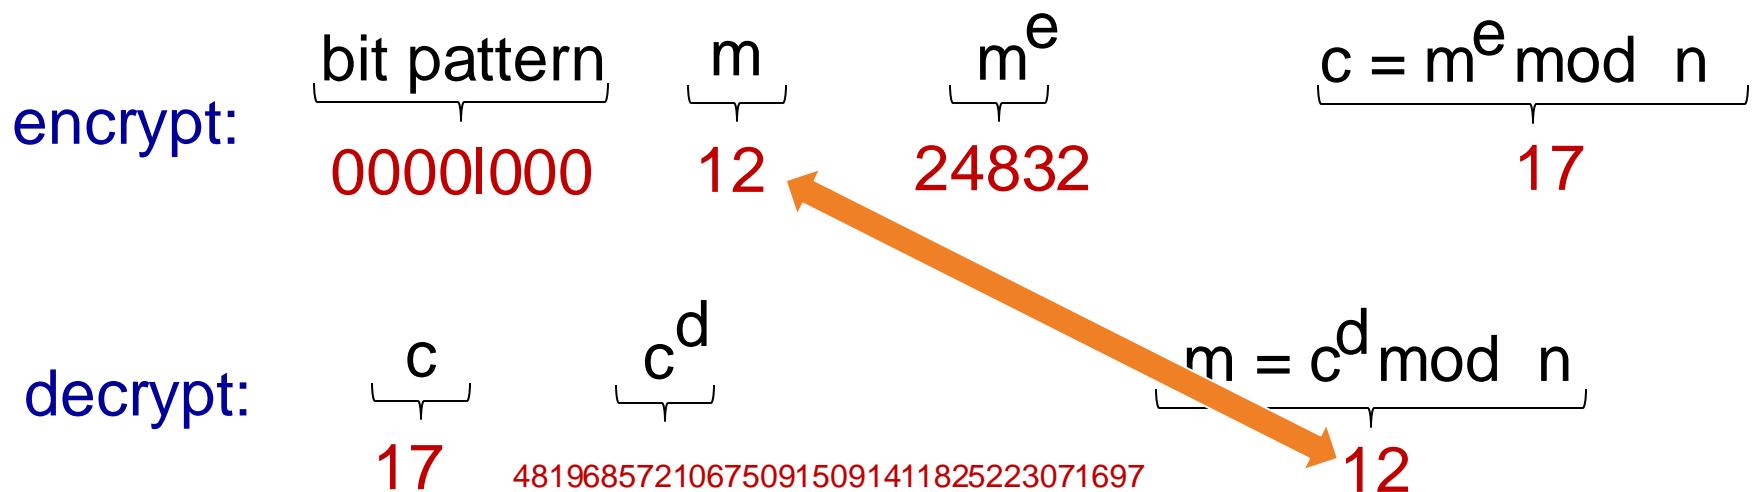
RSA example:

Bob chooses $p=5$, $q=7$. Then $n=35$, $z=24$.

$e=5$ (so e , z relatively prime).

$d=29$ (so $ed-1$ exactly divisible by z).

encrypting 8-bit messages.



Why does RSA work?

- must show that $c^d \bmod n = m$
where $c = m^e \bmod n$
- fact: for any x and y : $x^y \bmod n = x^{(y \bmod z)} \bmod n$
 - where $n = pq$ and $z = (p-1)(q-1)$
- thus,
$$\begin{aligned} c^d \bmod n &= (m^e \bmod n)^d \bmod n \\ &= m^{ed} \bmod n \\ &= m^{(ed \bmod z)} \bmod n \\ &= m^l \bmod n \\ &= m \end{aligned}$$

RSA: another important property

The following property will be *very* useful later:

$$\underbrace{K_B^-(K_B^+(m))}_{\text{use public key first,}} = m = \underbrace{K_B^+(K_B^-(m))}_{\text{use private key first, followed by public key}}$$

use public key first,
followed by
private key

use private key
first, followed by
public key

result is the same!

RSA: another important property

follows directly from modular arithmetic:

$$\begin{aligned}(m^e \bmod n)^d \bmod n &= m^{ed} \bmod n \\&= m^{de} \bmod n \\&= (m^d \bmod n)^e \bmod n\end{aligned}$$

Why is RSA secure?

- suppose you know Bob's public key (n, e) . How hard is it to determine d ?
- essentially need to find factors of n without knowing the two factors p and q
 - fact: factoring a big number is hard

RSA in practice: session keys

- exponentiation in RSA is computationally intensive
- DES is at least 100 times faster than RSA
- use public key crypto to establish secure connection, then establish second key – symmetric session key – for encrypting data

session key, K_s

- Bob and Alice use RSA to exchange a symmetric key K_s
- once both have K_s , they use symmetric key cryptography