

COMP9121: Design of Networks and Distributed Systems

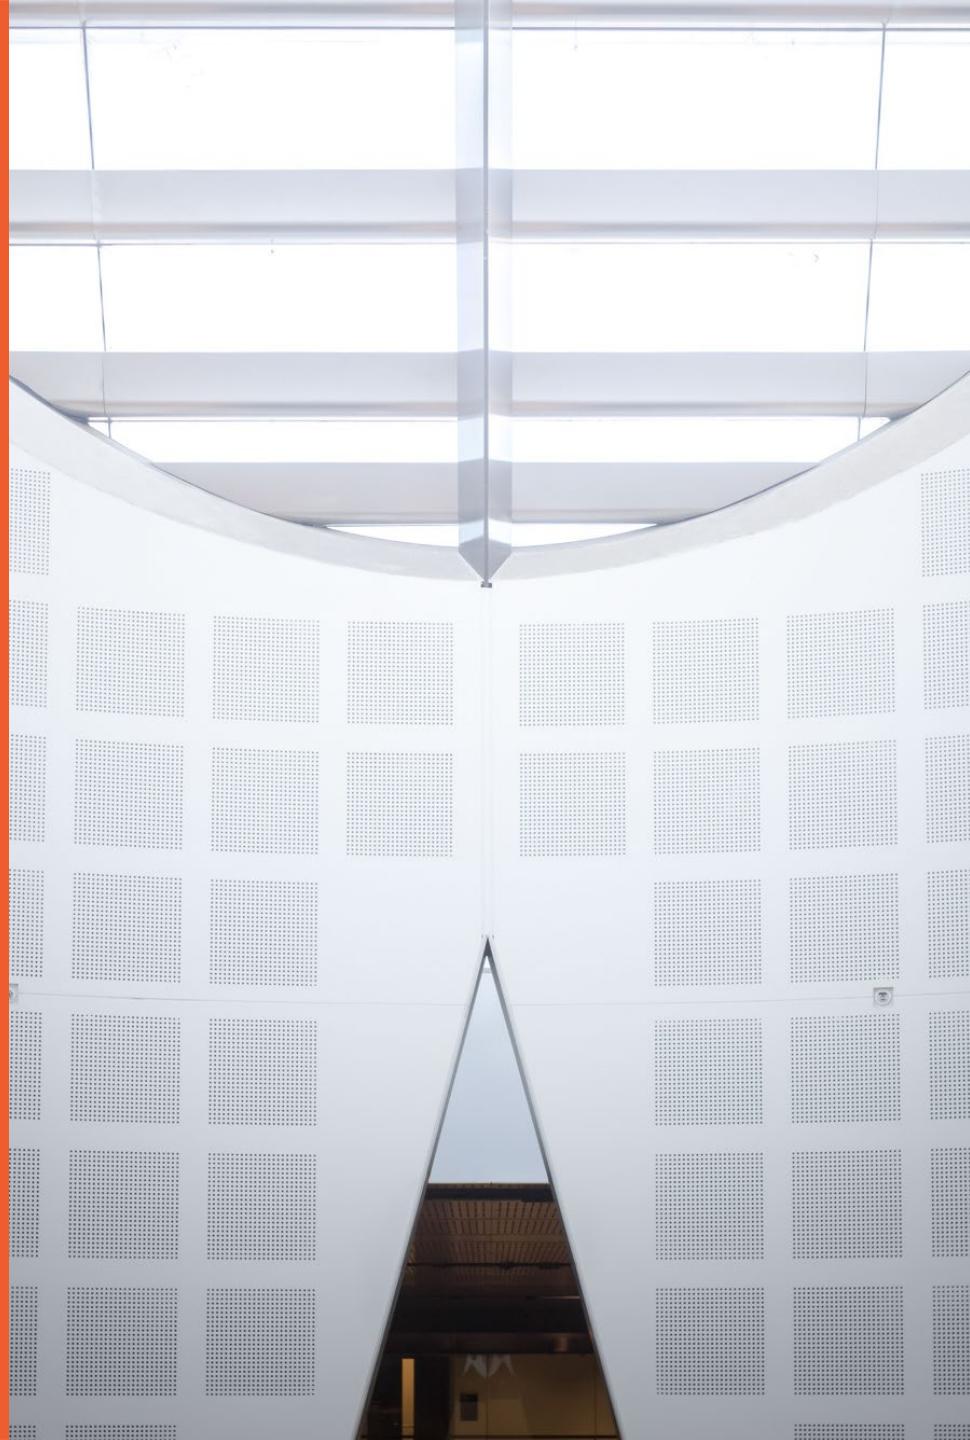
Week 11: Network Security 1

Wei Bao

School of Computer Science



THE UNIVERSITY OF
SYDNEY



Network Security

Chapter goals:

- understand principles of network security:
 - cryptography and its *many* uses beyond “confidentiality”
 - authentication
 - message integrity
- security in practice:
 - firewalls and intrusion detection systems
 - security in application, transport, network, link layers

What is network security?

confidentiality: only sender, intended receiver should “understand” message contents

- sender encrypts message
- receiver decrypts message

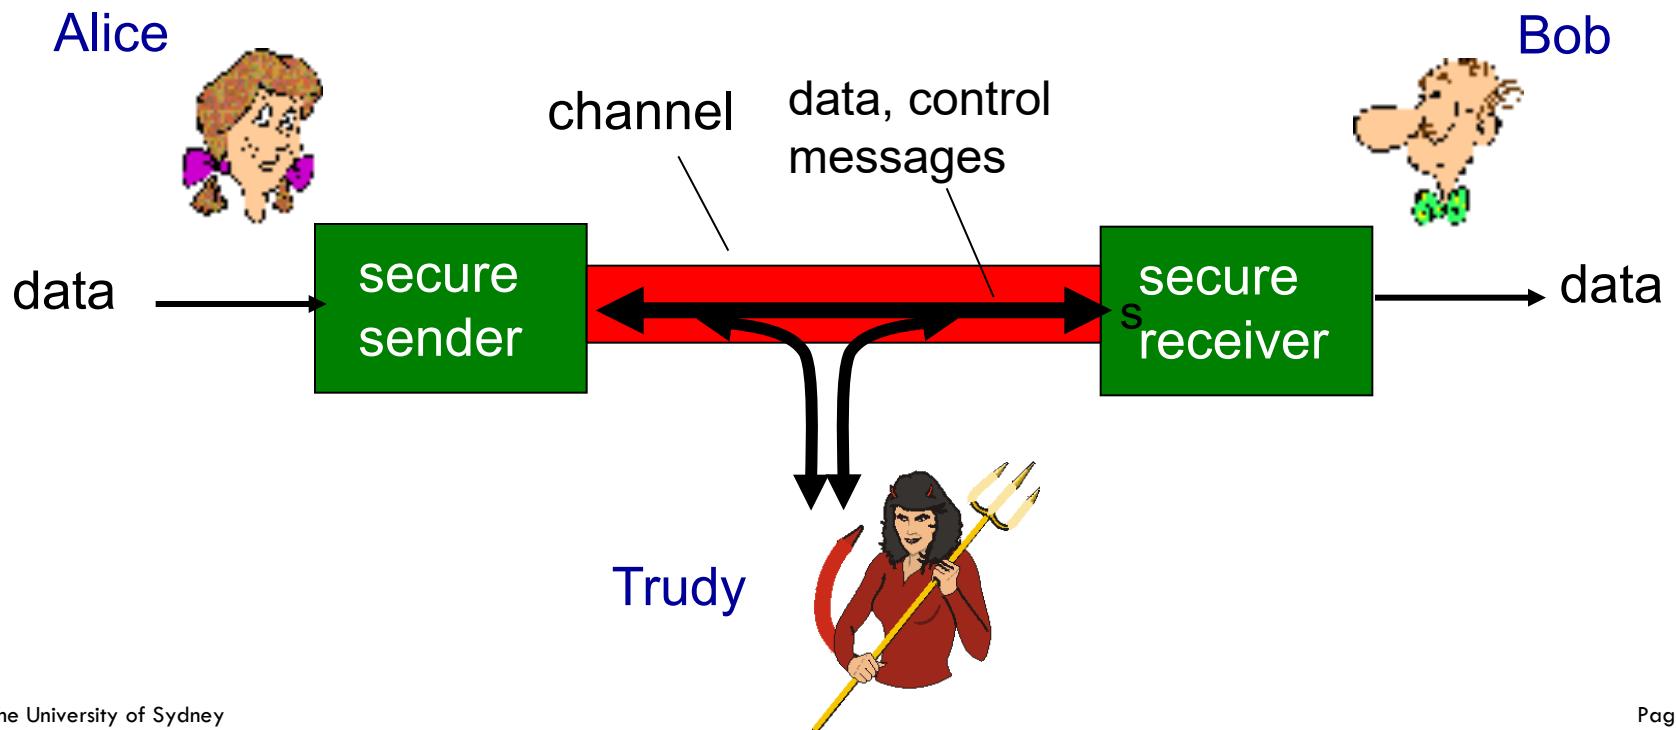
authentication: sender, receiver want to confirm identity of each other

message integrity: sender, receiver want to ensure message not altered (in transit, or afterwards) without detection

access and availability: services must be accessible and available to users

Friends and enemies: Alice, Bob, Trudy

- well-known in network security world
- Bob, Alice (lovers!) want to communicate “securely”
- Trudy (intruder) may intercept, delete, add messages



Who might Bob, Alice be?

- ... well, *real-life* Bobs and Alices!
- Web browser/server for electronic transactions (e.g., on-line purchases)
- on-line banking client/server
- DNS servers
- routers exchanging routing table updates

There are bad guys (and girls) out there!

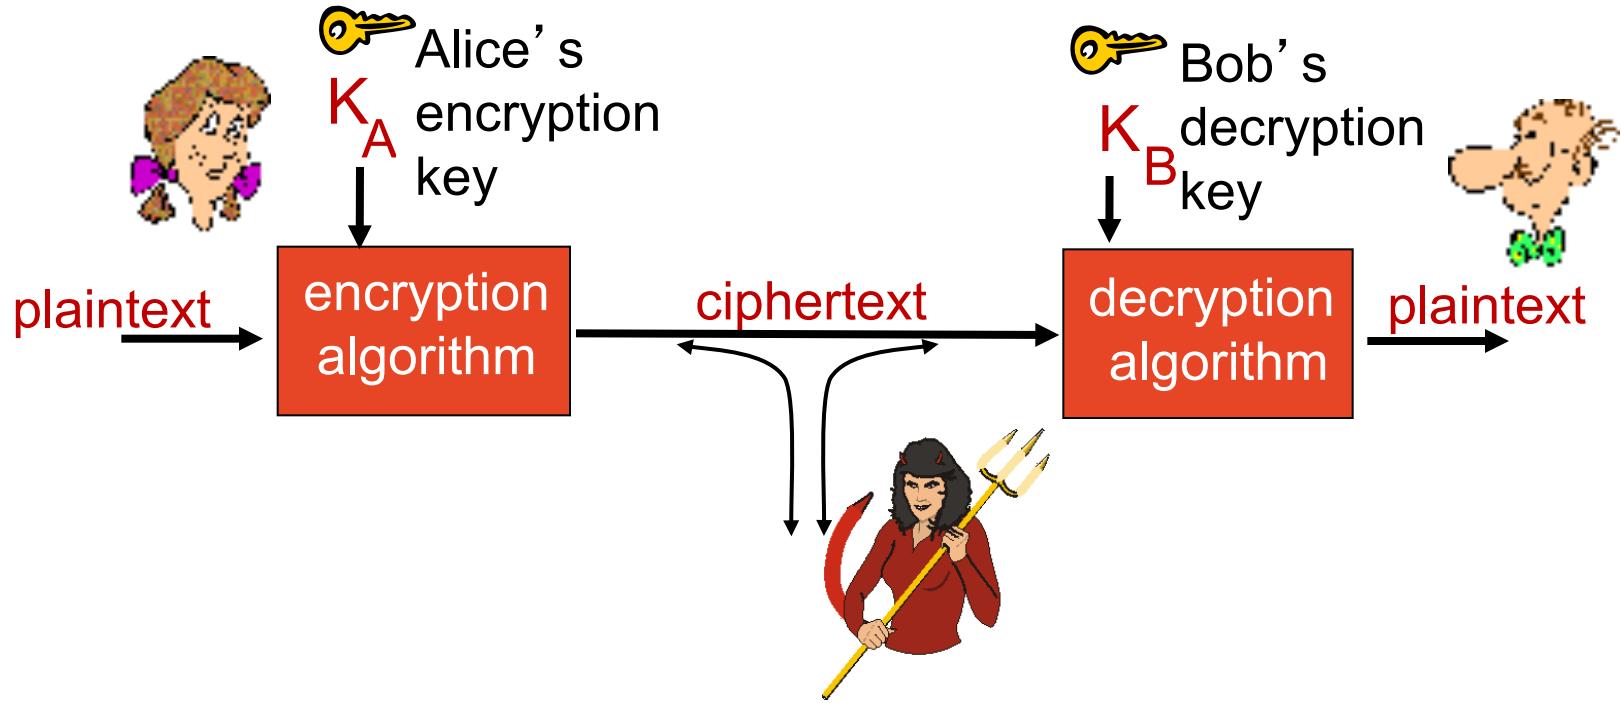
Q: What can a “bad guy” do?

A: A lot!

- 攻击 {
- **eavesdrop**: intercept messages
 - actively **insert** messages into connection
 - **impersonation**: can fake (spoof) source address in packet (or any field in packet)
 - **hijacking**: “take over” ongoing connection by removing sender or receiver, inserting himself in place
 - **denial of service**: prevent service from being used by others (e.g., by overloading resources)

Principles of cryptography

The language of cryptography



m plaintext message

$K_A(m)$ ciphertext, encrypted with key K_A

$m = K_B(K_A(m))$

Breaking an encryption scheme

know only cipher-text

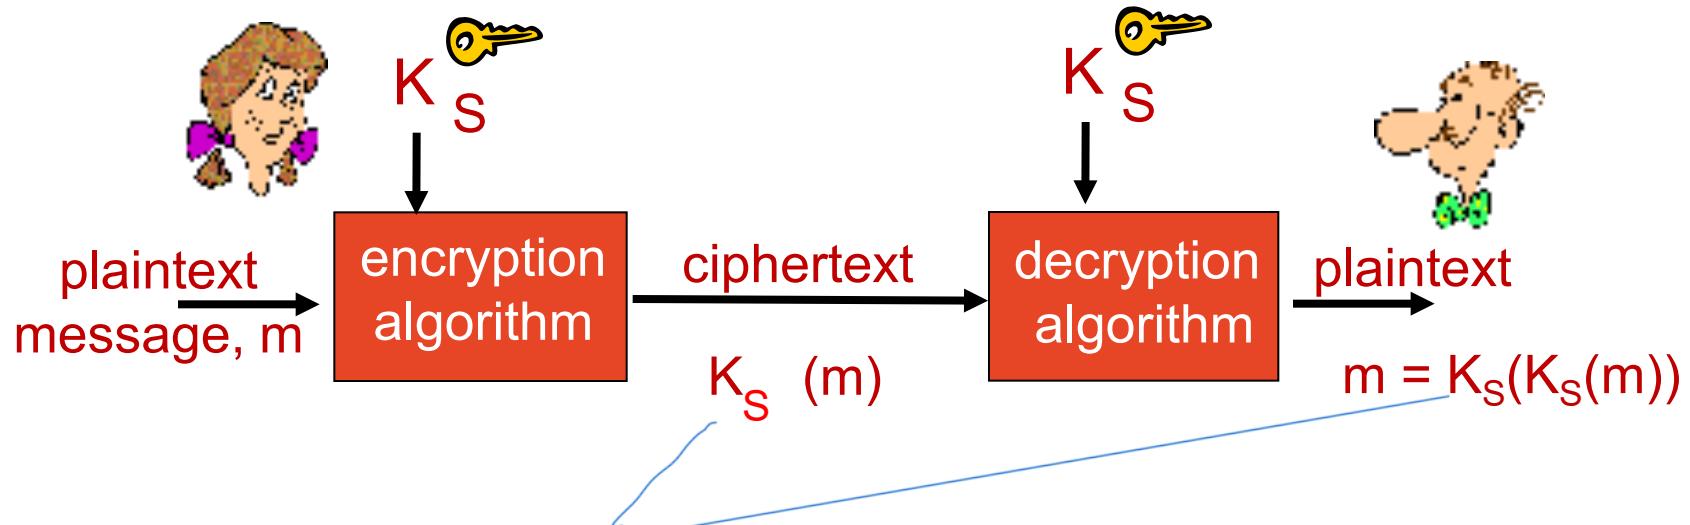
- **cipher-text only attack:**
Trudy has ciphertext she can analyze
- two approaches:
 - brute force: search through all keys
 - statistical analysis

↓
[如果明文出现频率较高，则很容易判断出对应的 cipher-text]

know some plaintext with its corresponding ciphertext

- **known-plaintext attack:**
Trudy has plaintext corresponding to ciphertext
 - e.g., in monoalphabetic cipher, Trudy determines pairings for a,l,i,c,e,b,o,
- **chosen-plaintext attack:**
Trudy can get ciphertext for chosen plaintext

Symmetric key cryptography



symmetric key crypto: Bob and Alice share same (symmetric) key: K

- e.g., key is knowing substitution pattern in mono alphabetic substitution cipher

Q: how do Bob and Alice agree on key value?

Simple encryption scheme

substitution cipher: substituting one thing for another

- monoalphabetic cipher: substitute one letter for another

plaintext: abcdefghijklmnopqrstuvwxyz

ciphertext: mnbvcxzasdfghjklpoiuytrewq

e.g.: Plaintext: bob. i love you. alice

ciphertext: nkn. s gktc wky. mgsbc



Encryption key: mapping from set of 26 letters
to set of 26 letters

A more sophisticated encryption approach

- n substitution ciphers, M_1, M_2, \dots, M_n
- **cycling pattern:**
 - e.g., $n=4$: $\underline{M_1, M_3, M_4, M_3, M_2}$; M_1, M_3, M_4, M_3, M_2 ; ..
- for each new plaintext symbol, use subsequent substitution pattern in cyclic pattern
 - dog: d from M_1 , o from M_3 , g from M_4



Encryption key: n substitution ciphers, and cyclic pattern

- key need not be just n-bit pattern

A more sophisticated encryption approach

plaintext: abcdefghijklmnopqrstuvwxyz

M1 : asdfghmnbcxzjklpoiuytrewq

M2 : oiuymnbvcxzasdfghjklptrewq

M3 : jklpoiuytrewqmnbnvcxzasdfgh

M4 : iuybvcmnfgfxzasdtrewqjklpo



Symmetric key crypto: DES

DES: Data Encryption Standard

- US encryption standard [NIST 1993]
 - 56-bit symmetric key, 64-bit plaintext input
 - block cipher
 - how secure is DES?
 - DES Challenge: 56-bit-key-encrypted phrase decrypted (brute force) in less than a day
 - no known good analytic attack
 - making DES more secure:
 - 3DES: encrypt 3 times with 3 different keys
-

Symmetric key crypto: DES

DES operation

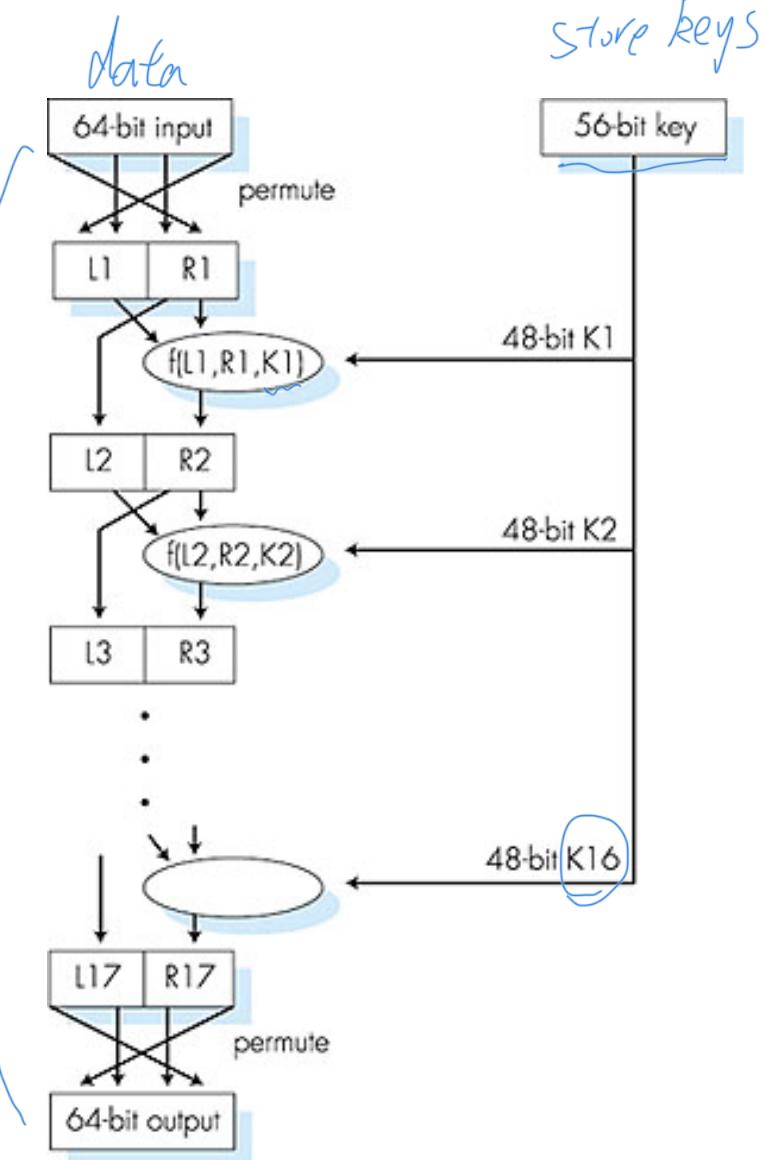
initial permutation

16 identical “rounds” of
function application,
each using different 48
bits of key

final permutation

不是一变全变的，
只变一个 bit

会导致最
下面改
变很多



AES: Advanced Encryption Standard

- symmetric-key NIST standard, replaced DES (Nov 2001)
- processes data in 128 bit blocks
- 128, 192, or 256 bit keys
- brute force decryption (try each key) taking 1 sec on DES, takes 149 trillion years for AES

Public Key Cryptography

symmetric key crypto

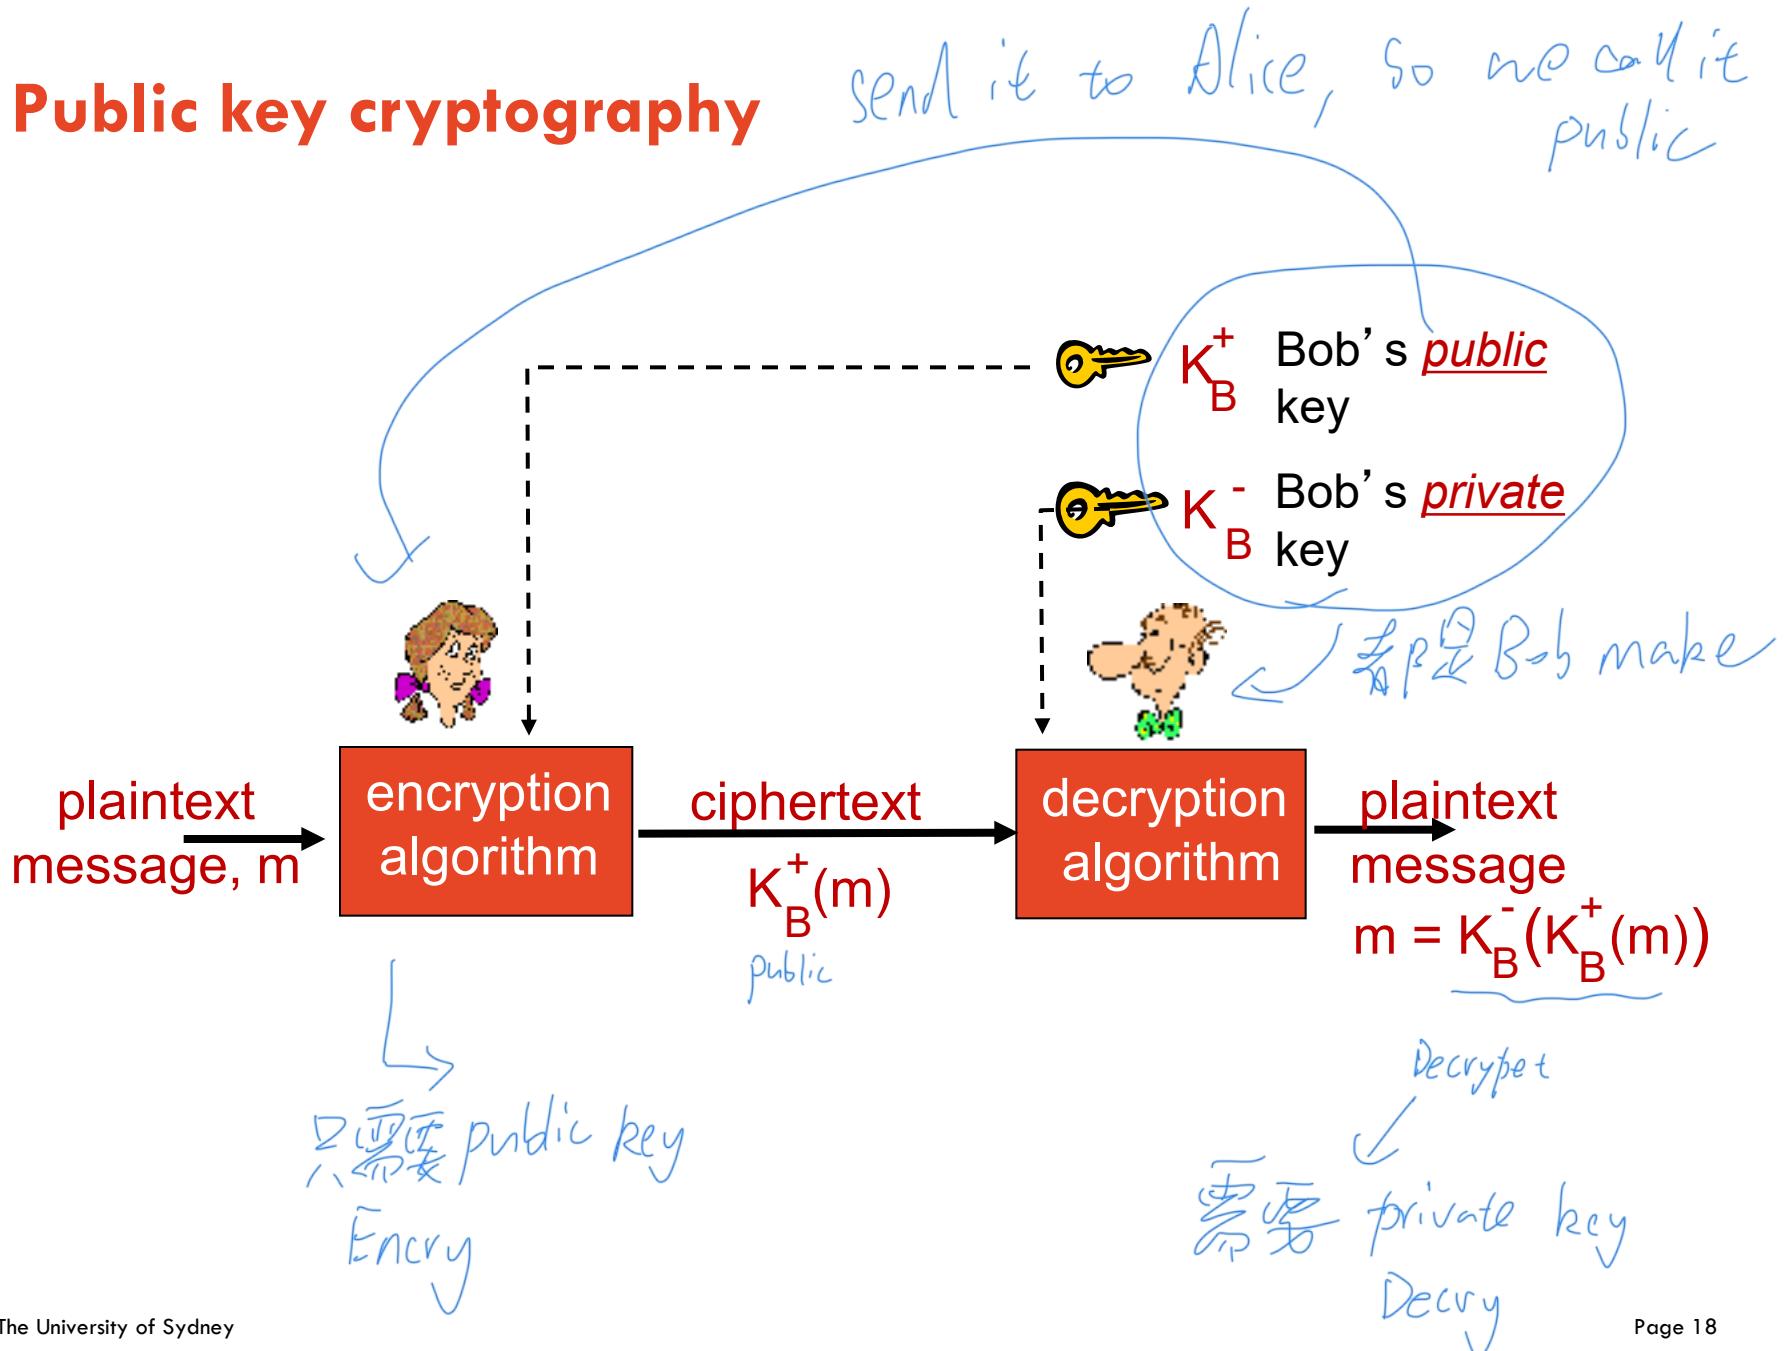
- requires sender, receiver know shared secret key
- Q: how to agree on key in first place (particularly if never “met”)?

public key crypto

- radically different approach [Diffie-Hellman76, RSA78]
- sender, receiver do *not* share secret key
- *public* encryption key known to *all*
- *private* decryption key known only to receiver



Public key cryptography



Public key encryption algorithms

requirements:

- ① need $K_B^+(\cdot)$ and $K_B^-(\cdot)$ such that

$$K_B^-(K_B^+(m)) = m$$

- ② given public key K_B^+ , it should be impossible to compute private key K_B^-

RSA: Rivest, Shamir, Adelson algorithm

Prerequisite: modular arithmetic

- $x \bmod n$ = remainder of x when divide by n
- facts:
 - $[(a \bmod n) + (b \bmod n)] \bmod n = (a+b) \bmod n$
 - $[(a \bmod n) - (b \bmod n)] \bmod n = (a-b) \bmod n$
 - $[(a \bmod n) * (b \bmod n)] \bmod n = (a*b) \bmod n$
- thus
 - $(a \bmod n)^d \bmod n = a^d \bmod n$
- example: $x=14$, $n=10$, $d=2$:
 - $(x \bmod n)^d \bmod n = 4^2 \bmod 10 = 6$
 - $x^d = 14^2 = 196 \quad x^d \bmod 10 = 6$

RSA: getting ready

- message: just a bit pattern
- bit pattern can be uniquely represented by an integer number
- thus, encrypting a message is equivalent to encrypting a number

example:

- $m = 10010001$. This message is uniquely represented by the decimal number 145.
- to encrypt m , we encrypt the corresponding number, which gives a new number (the ciphertext).

只有产生 key 的人需要做下面的步骤

RSA: Creating public/private key pair

因为很大，所以基本不可能在知道

public key 的情况下
找到 private
key

1. choose two large prime numbers p, q .
(e.g., 1024 bits each)
2. compute $n = pq$, $z = (p-1)(q-1)$
3. choose e (with $e < n$) that has no common factors with z (e, z are “relatively prime”).
4. choose d such that $\frac{ed-1}{z}$ is exactly divisible by z .
(in other words: $ed \bmod z = 1$).
5. public key is $\underbrace{(n,e)}_{K_B^+}$. private key is $\underbrace{(n,d)}_{K_B^-}$.

RSA: encryption, decryption

0. given (n,e) and (n,d) as computed above

I. to **encrypt** message $m (< n)$, compute

$$c = m^e \text{ mod } n \quad m \text{ means plain text}$$

2. to **decrypt** received bit pattern, c , compute

$$m = c^d \text{ mod } n \quad c \text{ means cipher text}$$

magic happens! $m = \underbrace{(m^e \text{ mod } n)^d}_{c} \text{ mod } n$



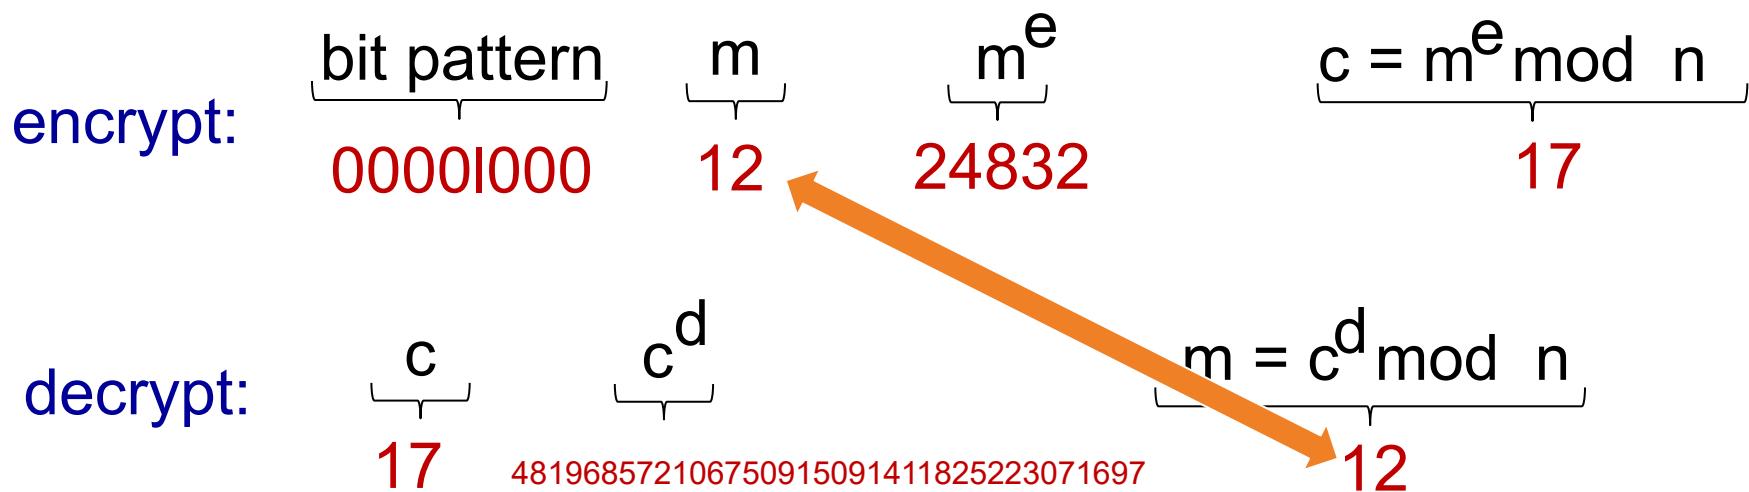
RSA example:

Bob chooses $p=5$, $q=7$. Then $n=35$, $z=24$.

$e=5$ (so e , z relatively prime).

$d=29$ (so $ed-1$ exactly divisible by z).

encrypting 8-bit messages.



Why does RSA work?

just need to
know

- must show that $c^d \bmod n = m$
where $c = m^e \bmod n$
- fact: for any x and y : $x^y \bmod n = x^{(y \bmod z)} \bmod n$
 - where $n = pq$ and $z = (p-1)(q-1)$
- thus,
$$\begin{aligned} \underline{c^d \bmod n} &= (\underline{m^e \bmod n})^d \bmod n \\ &= m^{ed} \bmod n \\ &= m^{(ed \bmod z)} \bmod n \\ &= m^l \bmod n \\ &= m \end{aligned}$$

RSA: another important property

The following property will be *very* useful later:

$$\underbrace{K_B^-(K_B^+(m))}_{\text{use public key first,}} = m = \underbrace{K_B^+(K_B^-(m))}_{\text{use private key first, followed by private key}}$$



use public key first,
followed by
private key

use private key
first, followed by
public key

result is the same!

RSA: another important property

follows directly from modular arithmetic:

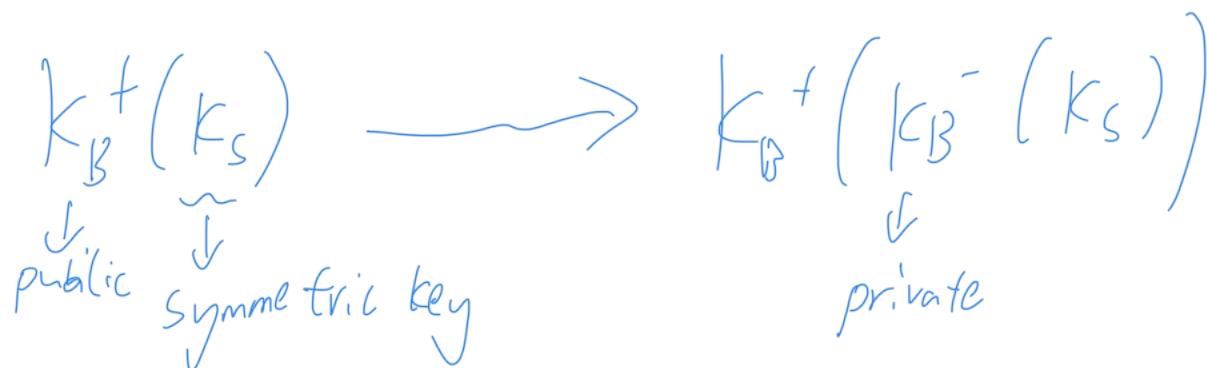
$$\begin{aligned}(m^e \bmod n)^d \bmod n &= m^{ed} \bmod n \\&= m^{de} \bmod n \\&= (m^d \bmod n)^e \bmod n\end{aligned}$$

Why is RSA secure?

- suppose you know Bob's public key (n, e) . How hard is it to determine d ?
- essentially need to find factors of n without knowing the two factors p and q
 - fact: factoring a big number is hard

RSA in practice: session keys

- exponentiation in RSA is computationally intensive
- DES is at least 100 times faster than RSA
- ① use public key crypto to establish secure connection, then establish second key – symmetric session key – for encrypting data
*为了较少计算时间, 这样使用 Symmetric session key
比如 2222^4 和 $2222 + 2222$ 的计算时间不同*
- Bob and Alice use RSA to exchange a symmetric key K_s
- once both have K_s , they use symmetric key cryptography



3. 混合加密的流程

为了解决公钥加密慢、对称加密密钥管理难的问题，常用的方案是混合加密，具体步骤如下：

1. **建立安全连接**：使用非对称加密（如RSA）来安全地传递一个会话密钥（即对称加密的密钥）。这个过程相对较短，因为只需要加密少量的数据（如密钥本身）。
2. **传递会话密钥**：一旦接收方获取到了会话密钥，双方就可以使用这个对称会话密钥来加密接下来的数据通信。
3. **加密数据传输**：之后的数据传输过程都使用对称加密算法（如DES或AES），因为对称加密速度更快，适合大数据量的传输。

4. 为什么用两种加密方式？

- **非对称加密**（如RSA）用于传递对称密钥，是因为它能保证密钥的安全性，即使攻击者截获了传输的数据，也无法解密，因为没有对应的私钥。
- **对称加密**（如DES或AES）用于后续数据传输，是因为它速度更快，适合加密大量数据。虽然它的密钥必须保密，但在整个通信过程只需用一次，并且可以通过非对称加密安全传输。



Message integrity and authentication

Authentication

Goal: Bob wants Alice to “prove” her identity to him

Protocol ap1.0: Alice says “I am Alice”



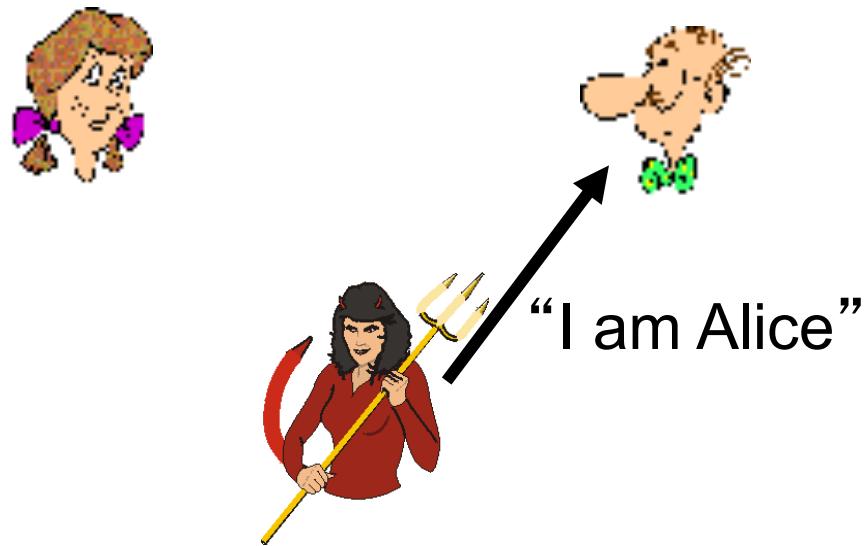
Failure scenario??



Authentication

Goal: Bob wants Alice to “prove” her identity to him

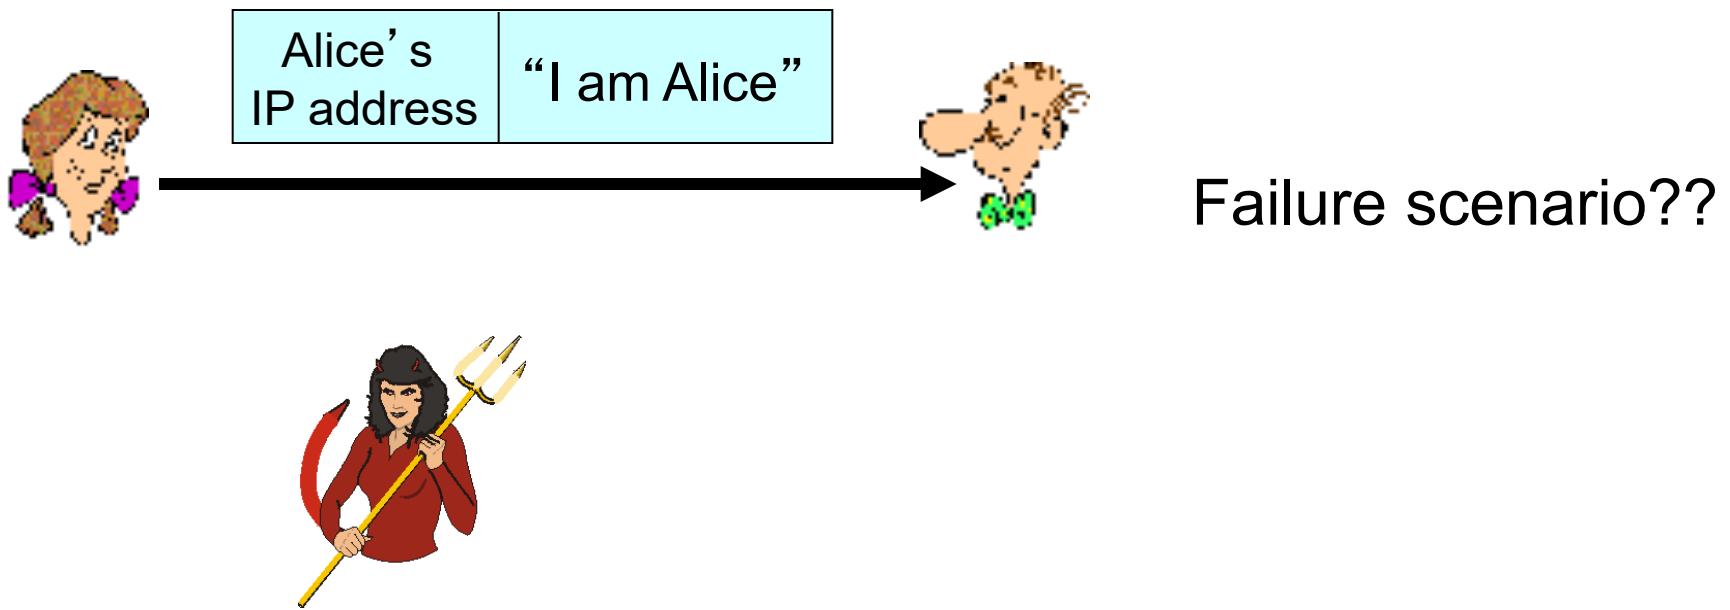
Protocol 1.0: Alice says “I am Alice”



in a network,
Bob can not “see” Alice,
so Trudy simply declares
herself to be Alice

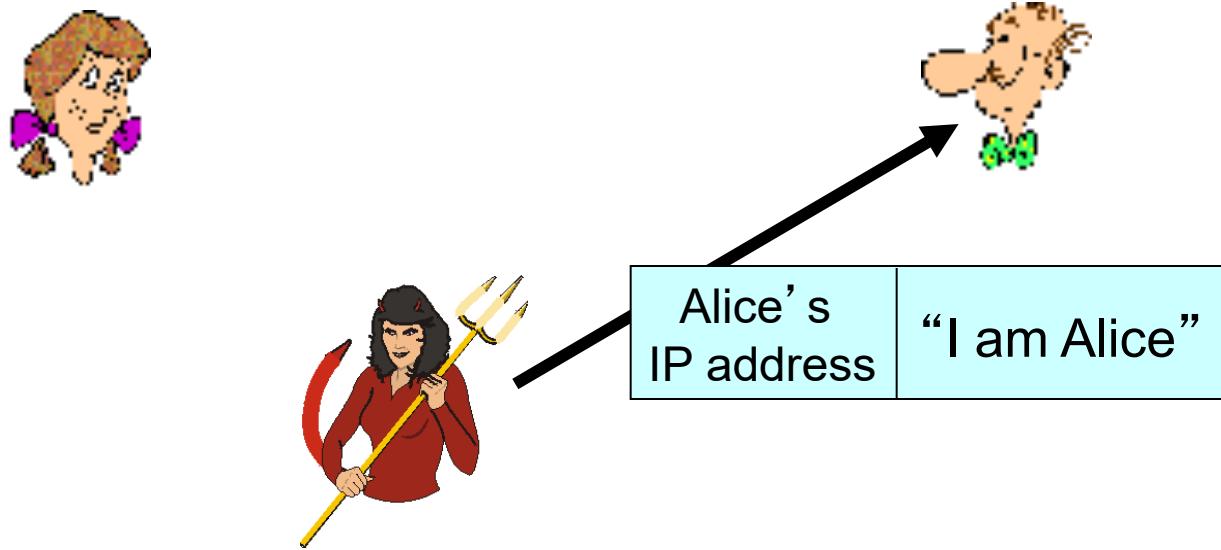
Authentication: another try

Protocol ap2.0: Alice says “I am Alice” in an IP packet containing her source IP address



Authentication: another try

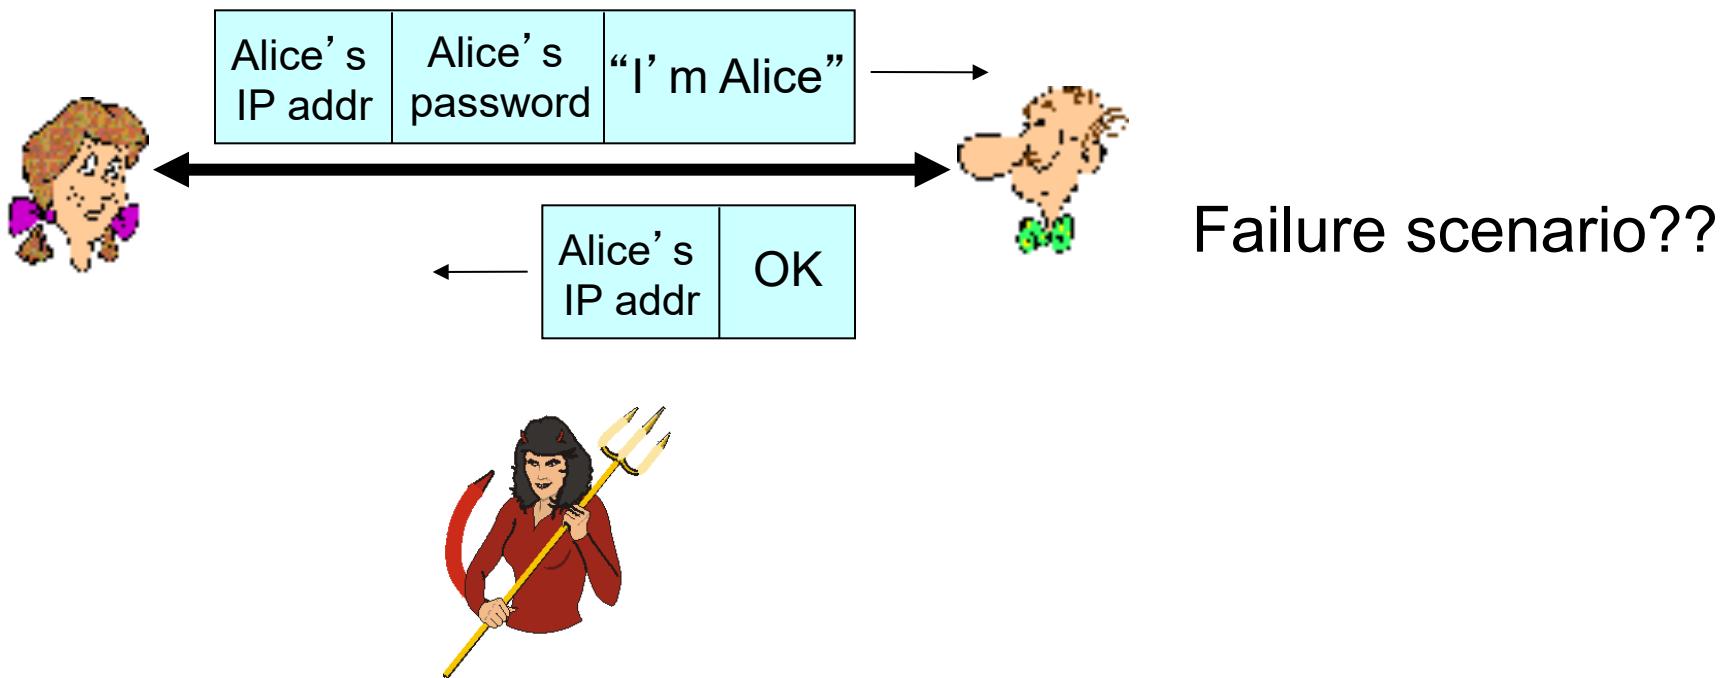
Protocol ap2.0: Alice says “I am Alice” in an IP packet containing her source IP address



Trudy can create
a packet
“spoofing”
Alice’s address

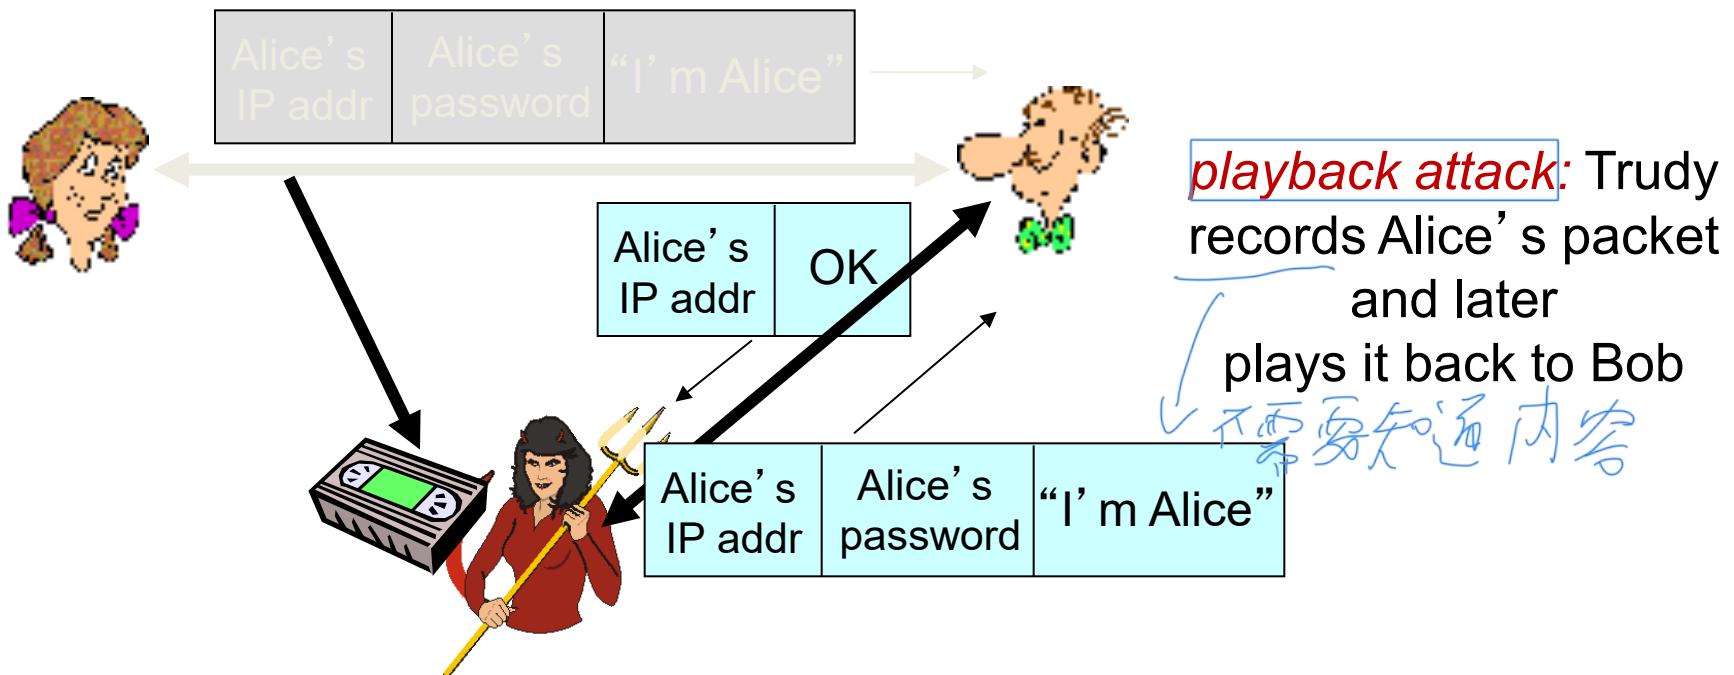
Authentication: another try

Protocol ap3.0: Alice says “I am Alice” and sends her secret password to “prove” it.



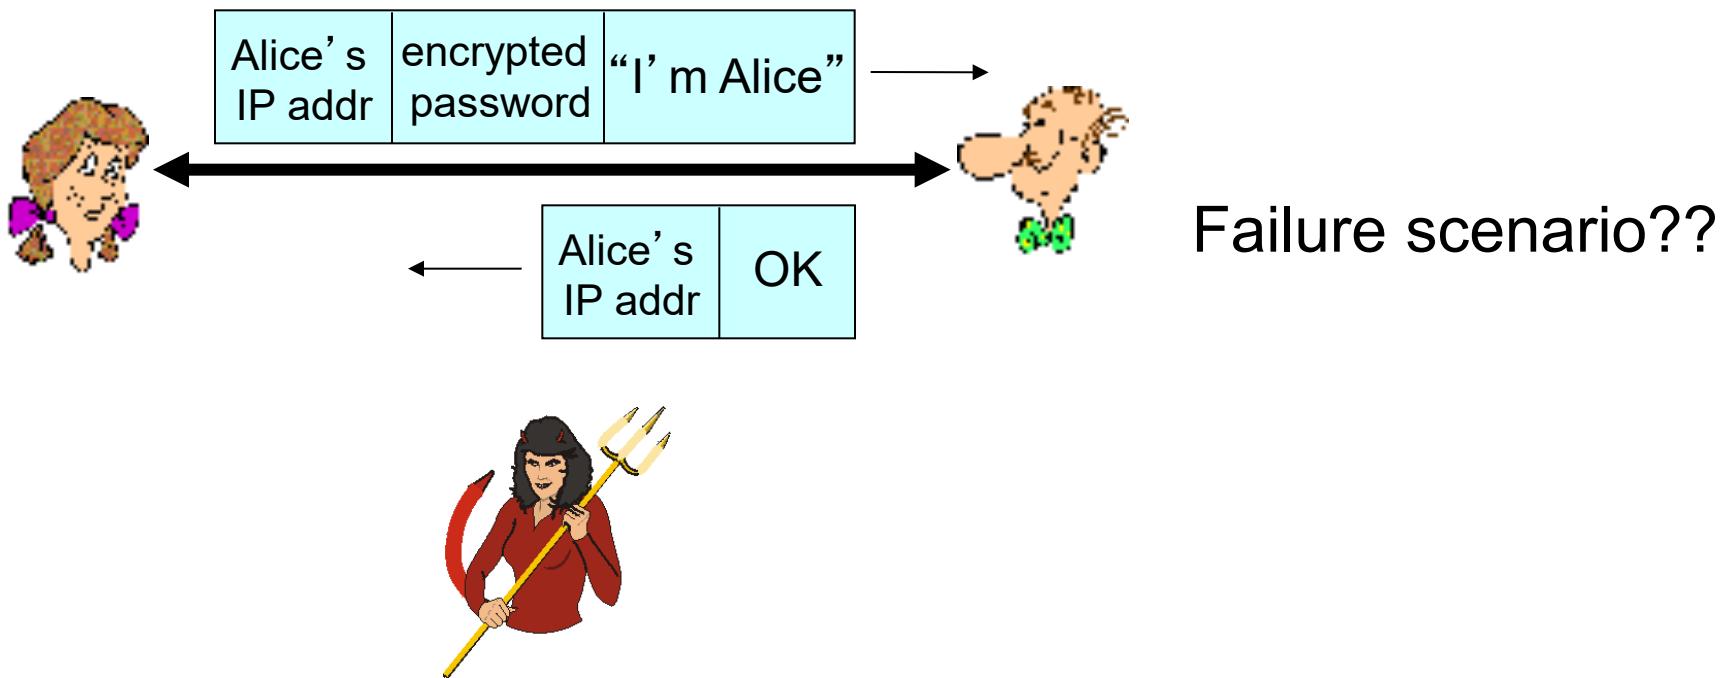
Authentication: another try

Protocol ap3.0: Alice says “I am Alice” and sends her secret password to “prove” it.



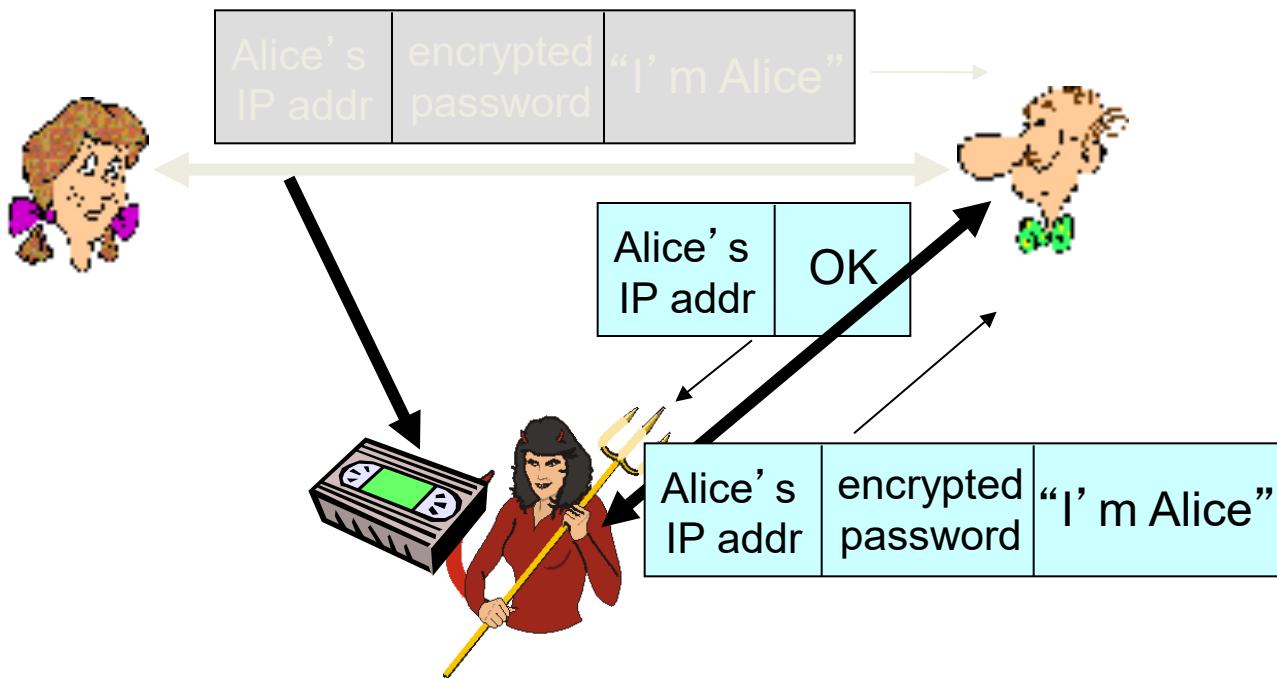
Authentication: yet another try

Protocol ap3.1: Alice says “I am Alice” and sends her **encrypted** secret password to “prove” it.



Authentication: yet another try

Protocol ap3.1: Alice says “I am Alice” and sends her **encrypted** secret password to “prove” it.



record
and
playback
still works!

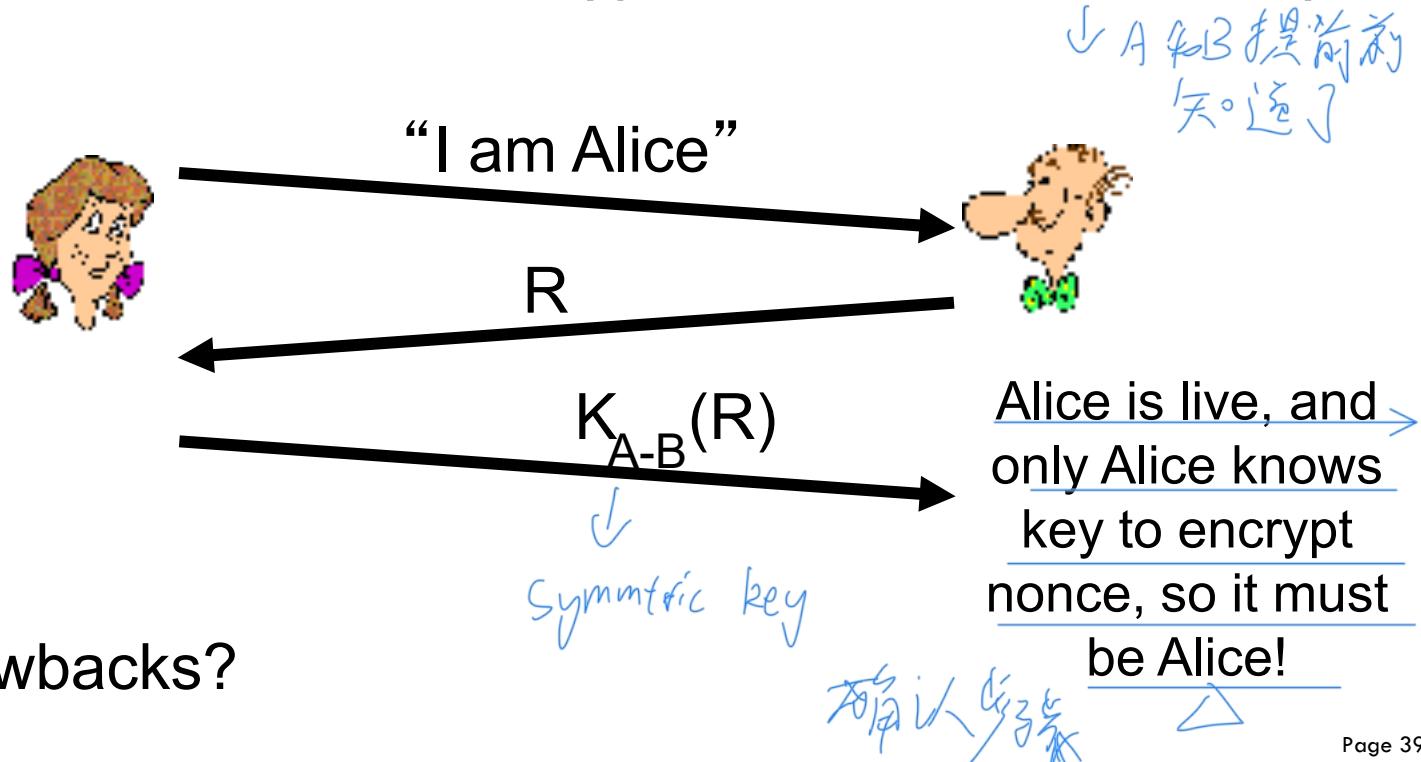
Authentication: yet another try

Goal: avoid playback attack

nonce: number (R) used only once-in-a-lifetime

ap4.0: to prove Alice “live”, Bob sends Alice **nonce**, R. Alice must return R, encrypted with shared secret key

因为下次 Bob 不会使用相同的 R
↑ So it's play back attack 不会发生



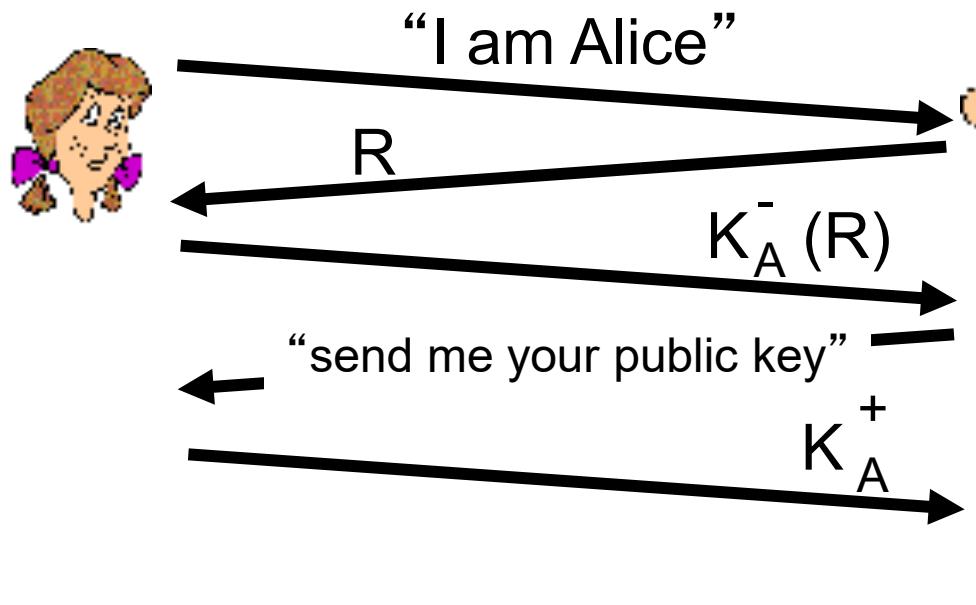
Drawbacks?

Authentication: ap5.0

ap4.0 requires shared symmetric key

- can we authenticate using public key techniques?

ap5.0: use nonce, public key cryptography



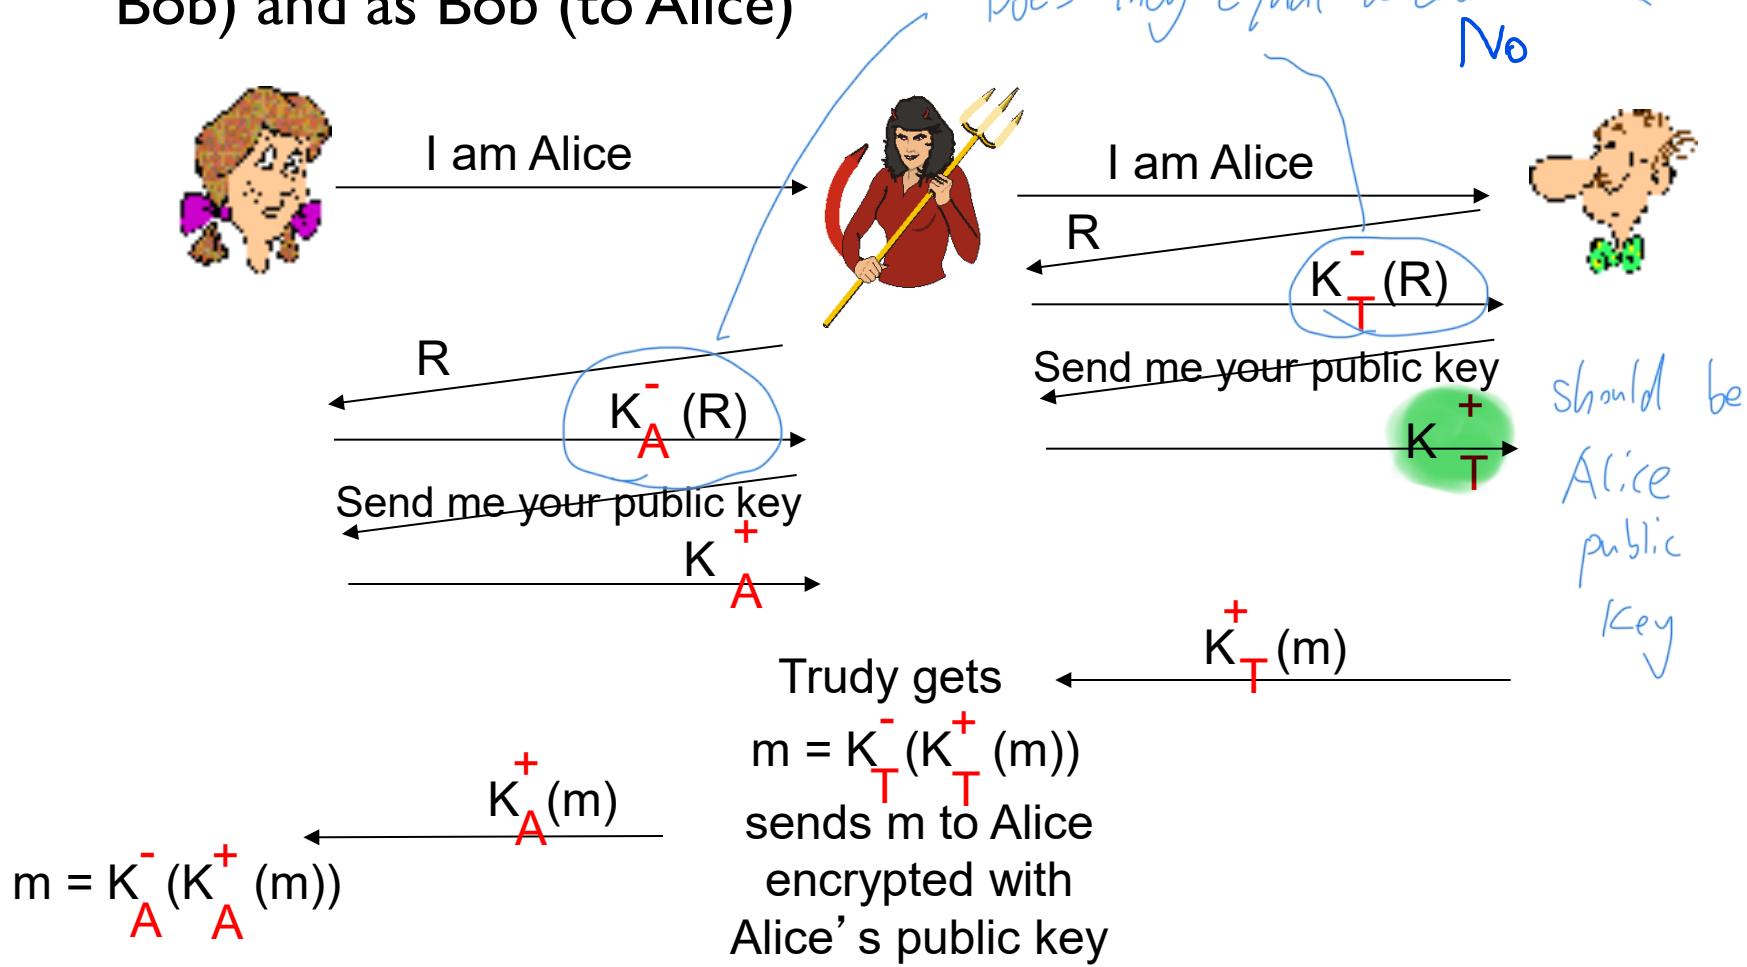
Bob computes
 $K_A^+(K_A^-(R)) = R$
and knows only Alice could have the private key, that encrypted R such that
 $K_A^+(K_A^-(R)) = R$

ap5.0: security hole

symmetric
key \rightarrow K_T

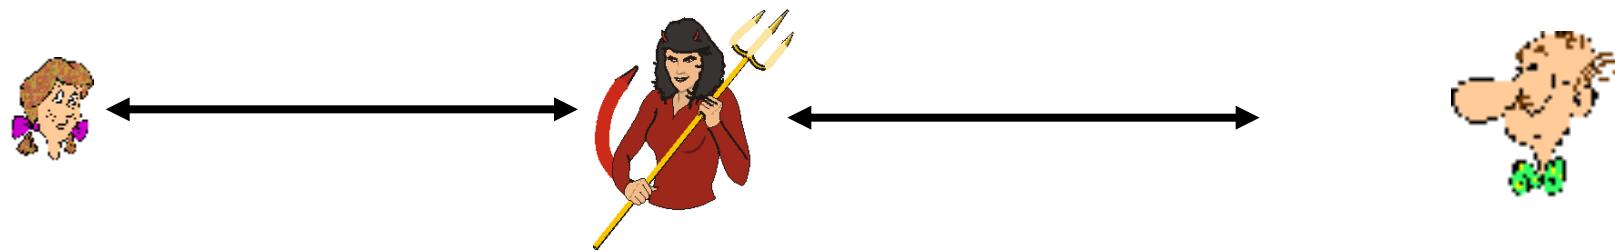
man (or woman) in the middle attack: Trudy poses as Alice (to Bob) and as Bob (to Alice)

Does they equal to each other?
No



ap5.0: security hole

man (or woman) in the middle attack: Trudy poses as Alice (to Bob) and as Bob (to Alice)



difficult to detect:

- Bob receives everything that Alice sends, and vice versa.
(e.g., so Bob, Alice can meet one week later and recall conversation! Nothing wrong is detected!)
- problem is that Trudy receives all messages as well!

κ_T^+ (m) is Trudy's public key. How does bob verify?

Digital signatures

solve above problem

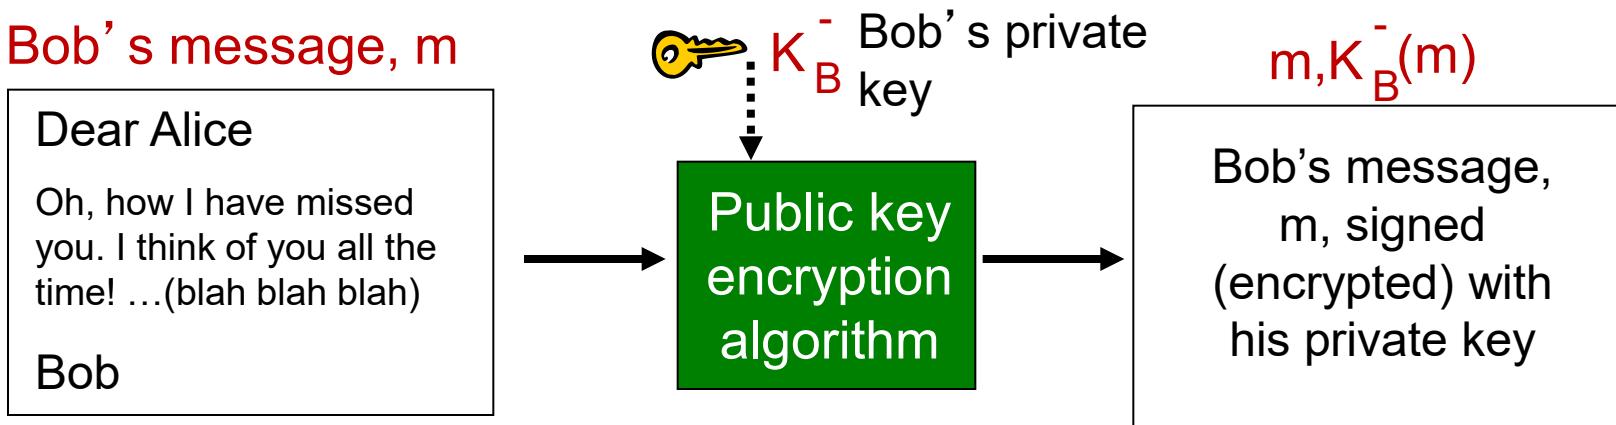
cryptographic technique analogous to hand-written signatures:

- sender (Bob) digitally signs document, establishing he is document owner/creator.
- *verifiable, nonforgeable*: recipient (Alice) can prove to someone that Bob, and no one else (including Alice), must have signed document

Digital signatures

simple digital signature for message m :

- Bob signs m by encrypting with his private key K_B^- , creating “signed” message, $K_B^-(m)$



Digital signatures

- suppose Alice receives msg m , with signature: $m, K_B^-(m)$
- Alice verifies m signed by Bob by applying Bob's public key K_B to $K_B^+()$ then checks $K_B^+(K_B^-(m)) = m$. $K_B^+(K_B^-(m)) = K_B^-(K_B^+(m))$
- If $K_B^+(K_B^-(m)) = m$, whoever signed m must have used Bob's private key.

Alice thus verifies that:

- Bob signed m
- no one else signed m
- Bob signed m and not m'

non-repudiation:

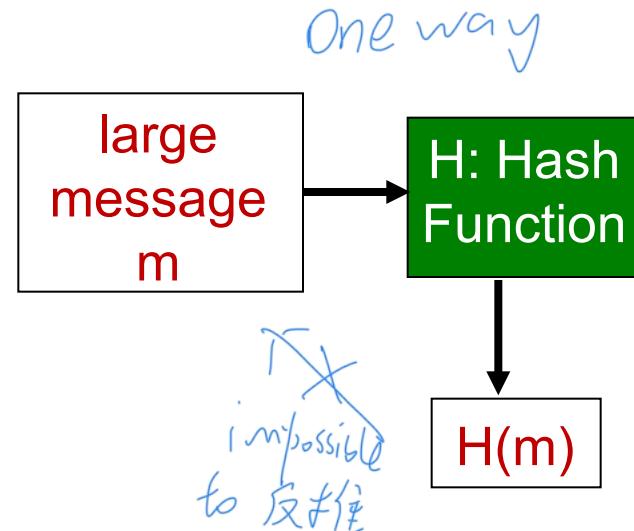
- ✓ Alice can take m , and signature $K_B^-(m)$ to court and prove that Bob signed m

Message digests

computationally expensive to public-key-encrypt long messages

goal: fixed-length, easy-to-compute digital “fingerprint”

- apply hash function H to m , get fixed size message digest, $H(m)$.



Hash function properties:

- many-to-1
- produces fixed-size msg digest (fingerprint)
- given message digest x , computationally infeasible to find m such that $x = H(m)$

Internet checksum: poor crypto hash function

Internet checksum has some properties of hash function:

- produces fixed length digest (16-bit sum) of message
- is many-to-one

But given message with given hash value, it is easy to find another message with same hash value:

<u>message</u>	<u>ASCII format</u>
I O U 1	49 4F 55 31
0 0 . 9	30 30 2E 39
9 B O B	39 42 D2 42
<hr/>	
B2 C1 D2 AC	

<u>message</u>	<u>ASCII format</u>
I O U <u>9</u>	49 4F 55 <u>39</u>
0 0 . <u>1</u>	30 30 2E <u>31</u>
9 B O B	39 42 D2 42
<hr/>	
B2 C1 D2 AC	

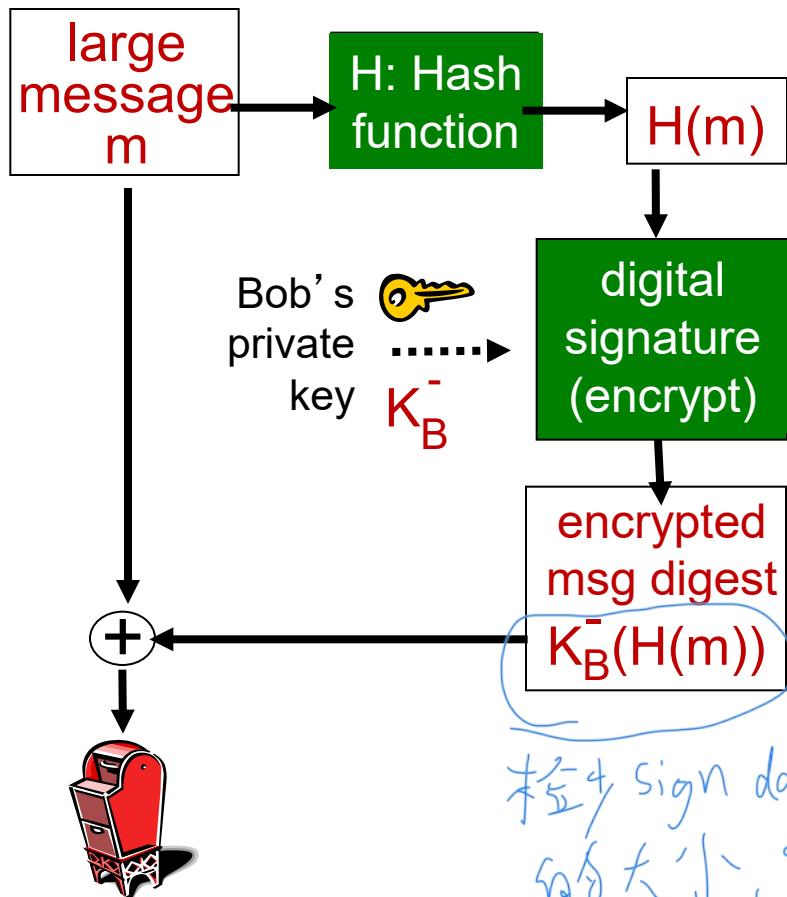
different messages
but identical checksums!

Hash function algorithms

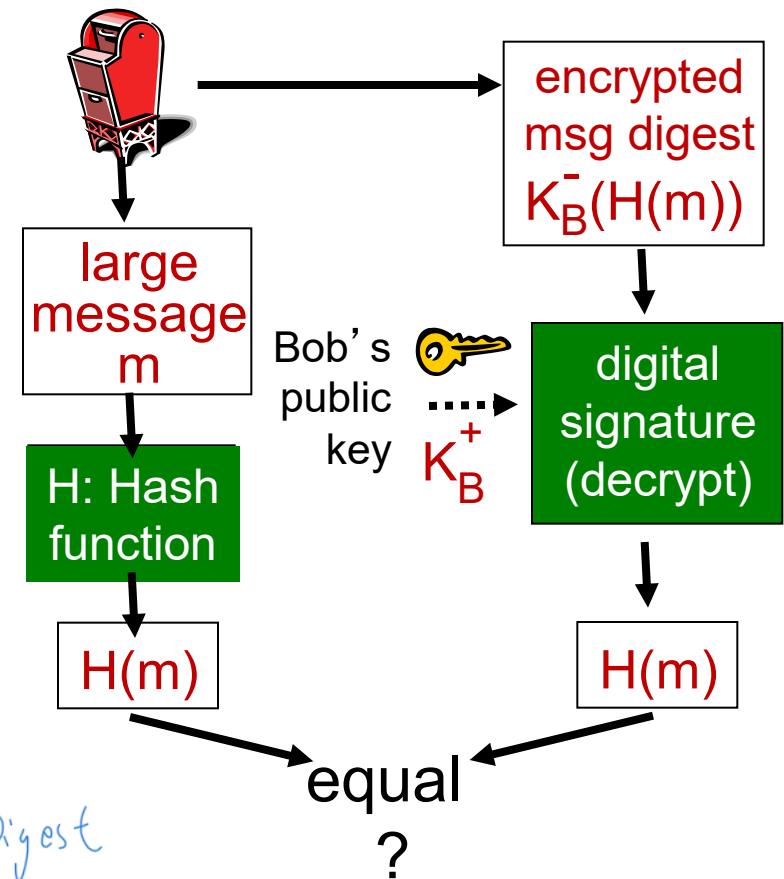
- MD5 hash function widely used (RFC 1321)
 - computes 128-bit message digest in 4-step process.
 - arbitrary 128-bit string x , appears difficult to construct msg m whose MD5 hash is equal to x
- SHA-1 is also used
 - US standard [NIST, FIPS PUB 180-1]
 - 160-bit message digest

Digital signature = signed message digest

Bob sends digitally signed message:



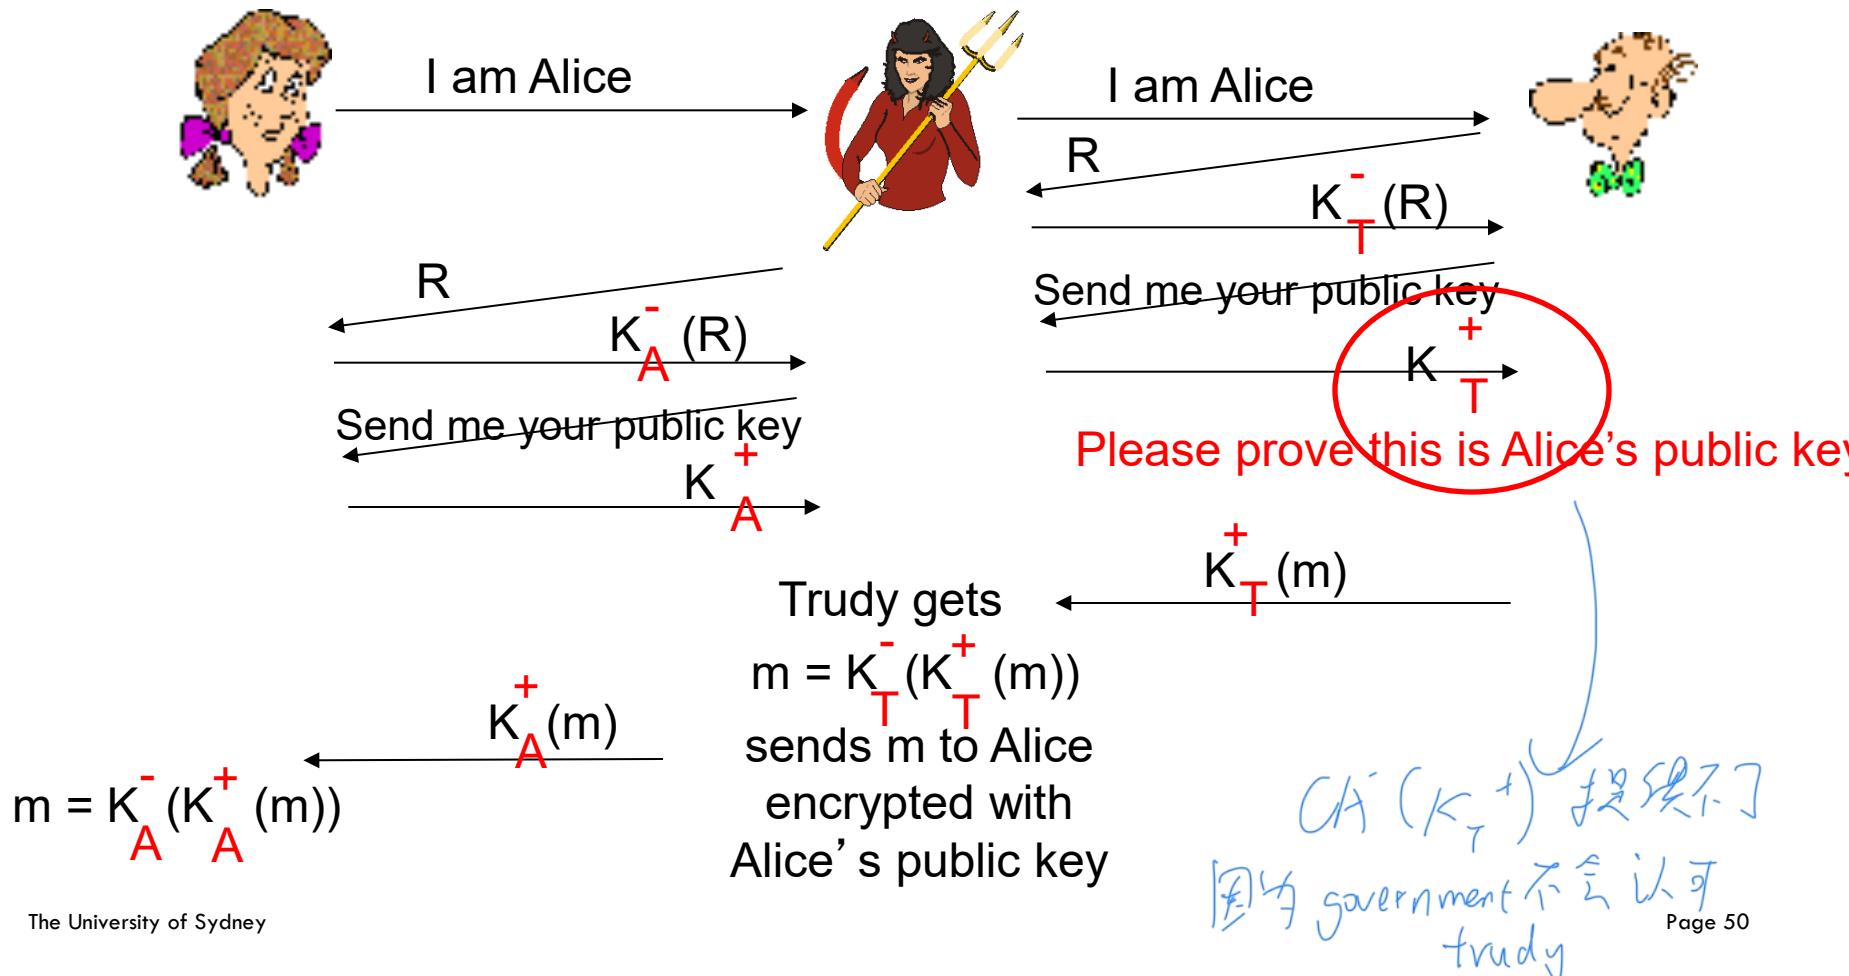
Alice verifies signature, integrity of digitally signed message:



检+ sign data
是否大小, ? sign Digest

Recall: ap5.0 security hole

man (or woman) in the middle attack: Trudy poses as Alice (to Bob) and as Bob (to Alice)

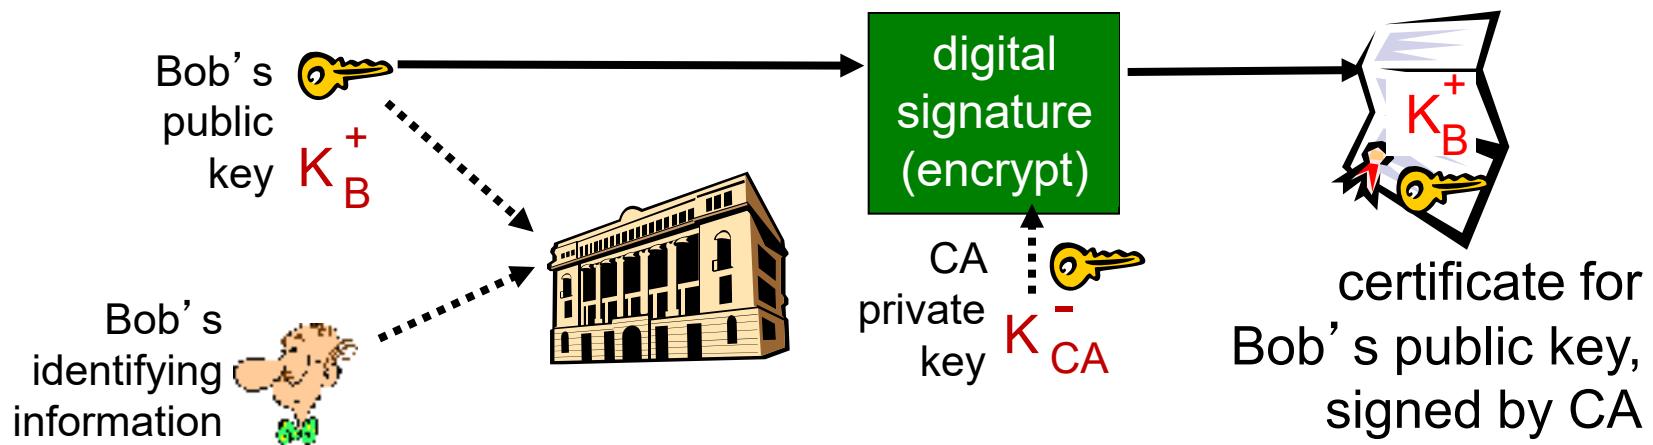


Certification authorities

- certification authority (CA): binds public key to particular entity, E.
- E (person, router) registers its public key with CA.
 - E provides “proof of identity” to CA.
 - CA creates certificate binding E to its public key.
 - certificate containing E’s public key digitally signed by CA – CA says “this is E’s public key”

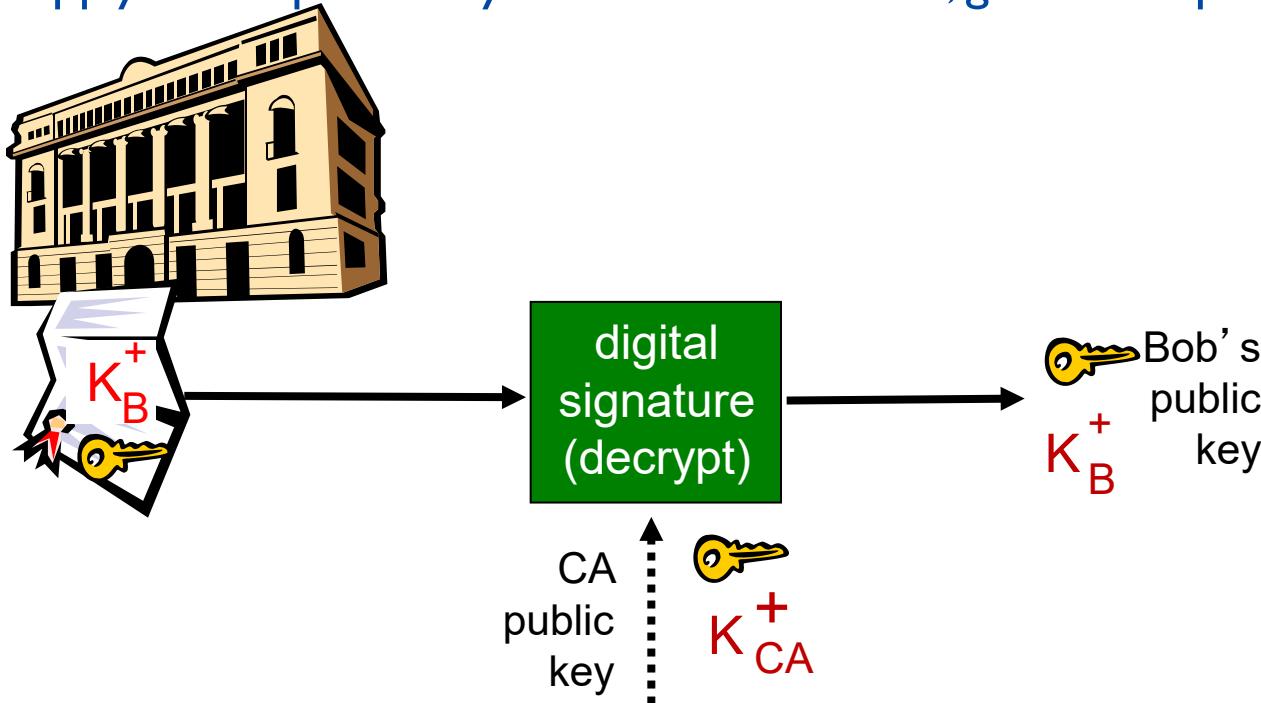
Government sign on Bob's public key

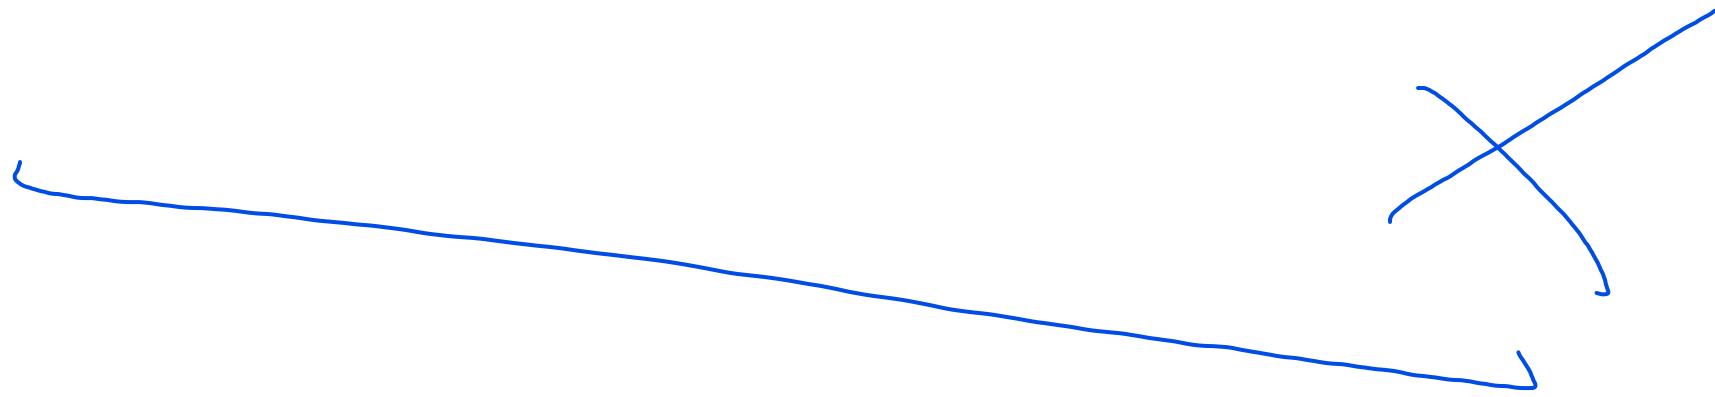
K_{CA}^- (K_B^+ is bob's public key)



Certification authorities

- when Alice wants Bob's public key:
 - gets Bob's certificate (Bob or elsewhere).
 - apply CA's public key to Bob's certificate, get Bob's public key

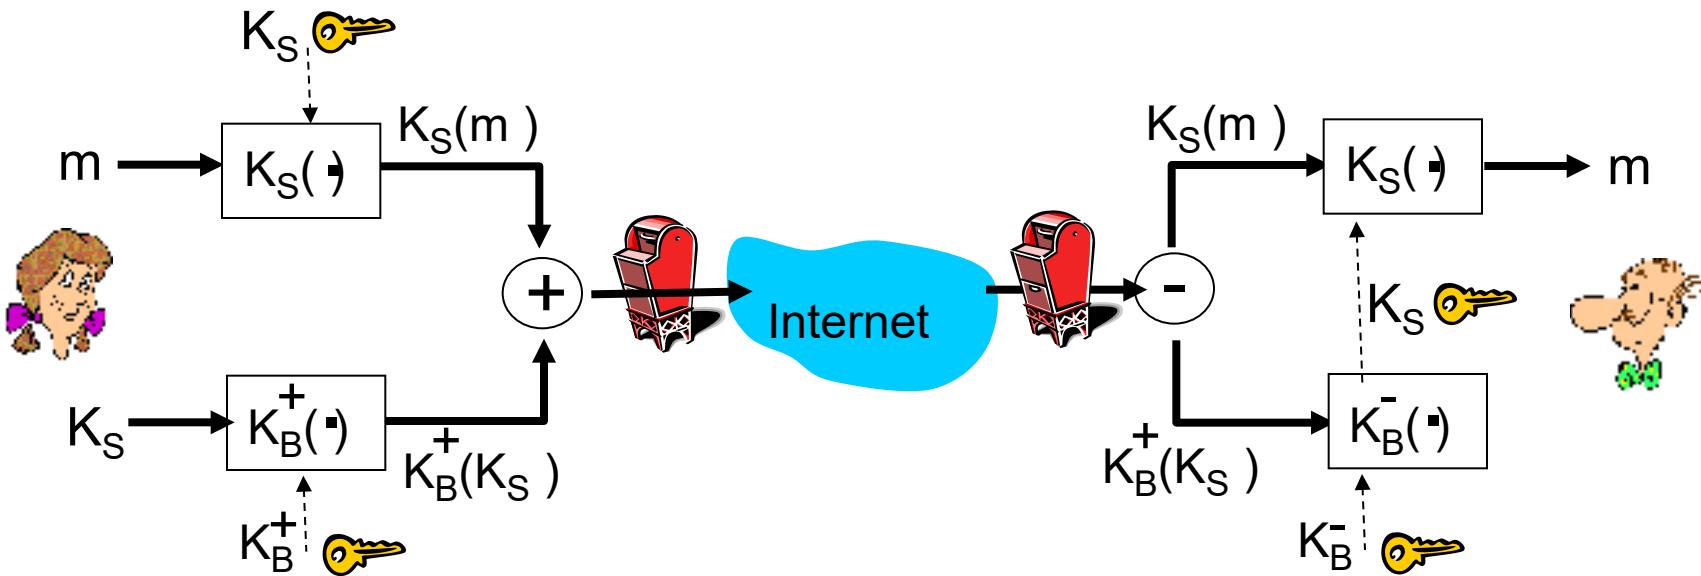




Securing e-mail

Secure e-mail, Pretty Good Privacy (PGP)

Alice wants to send **confidential** e-mail, m , to Bob.

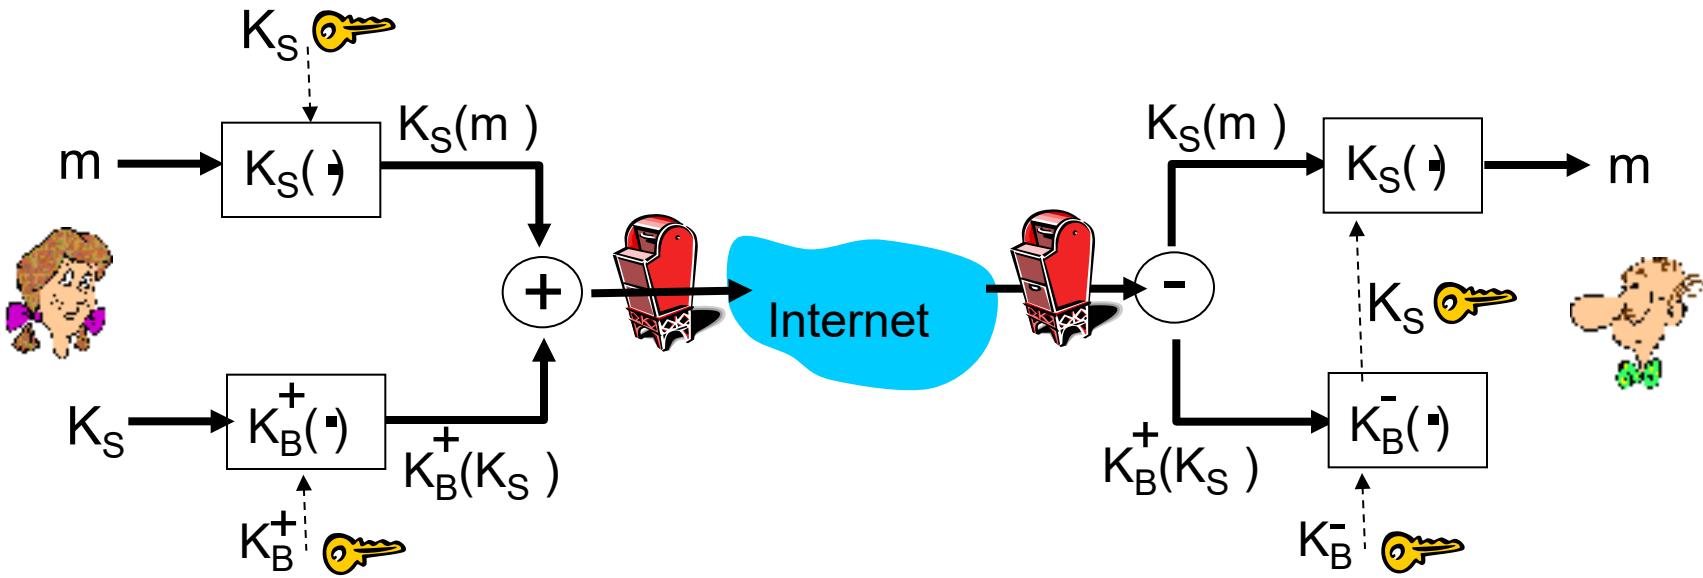


Alice:

- generates random *symmetric* private key, K_S
- encrypts message with K_S (for efficiency)
- also encrypts K_S with Bob's public key
- sends both $K_S(m)$ and $K_B(K_S)$ to Bob

Secure e-mail

Alice wants to send **confidential** e-mail, m , to Bob.

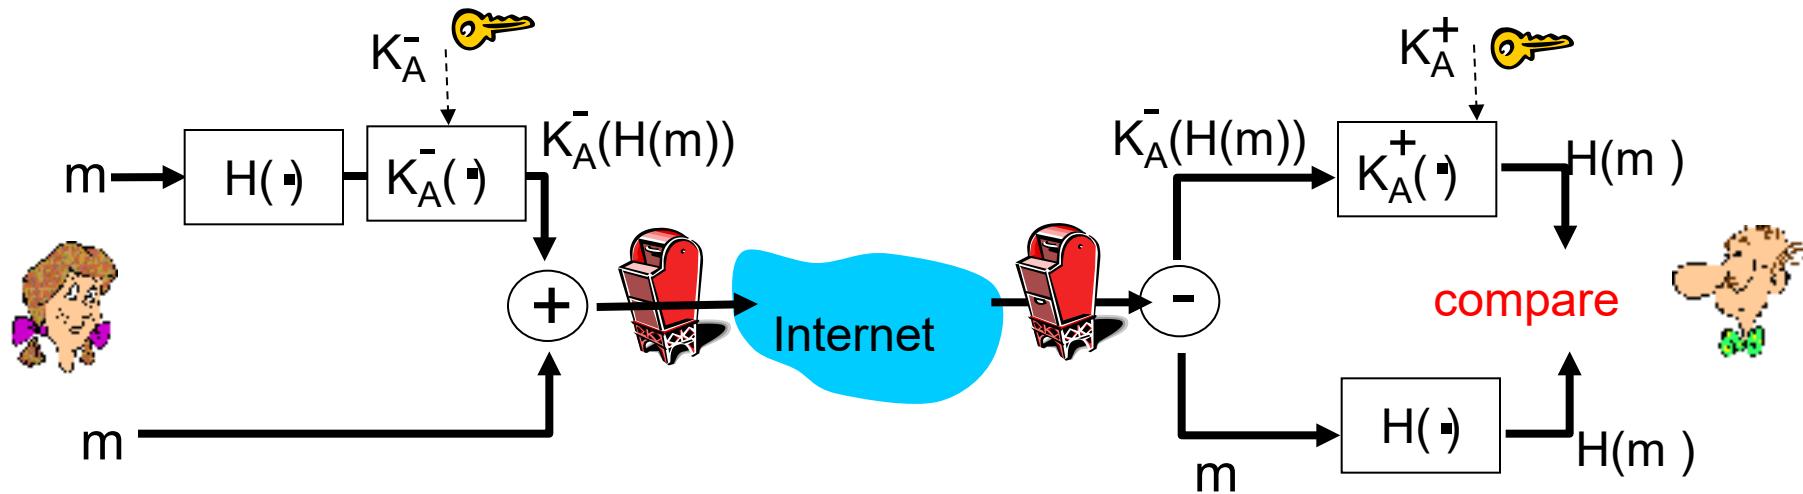


Bob:

- uses his private key to decrypt and recover K_S
- uses K_S to decrypt $K_S(m)$ to recover m

Secure e-mail (continued)

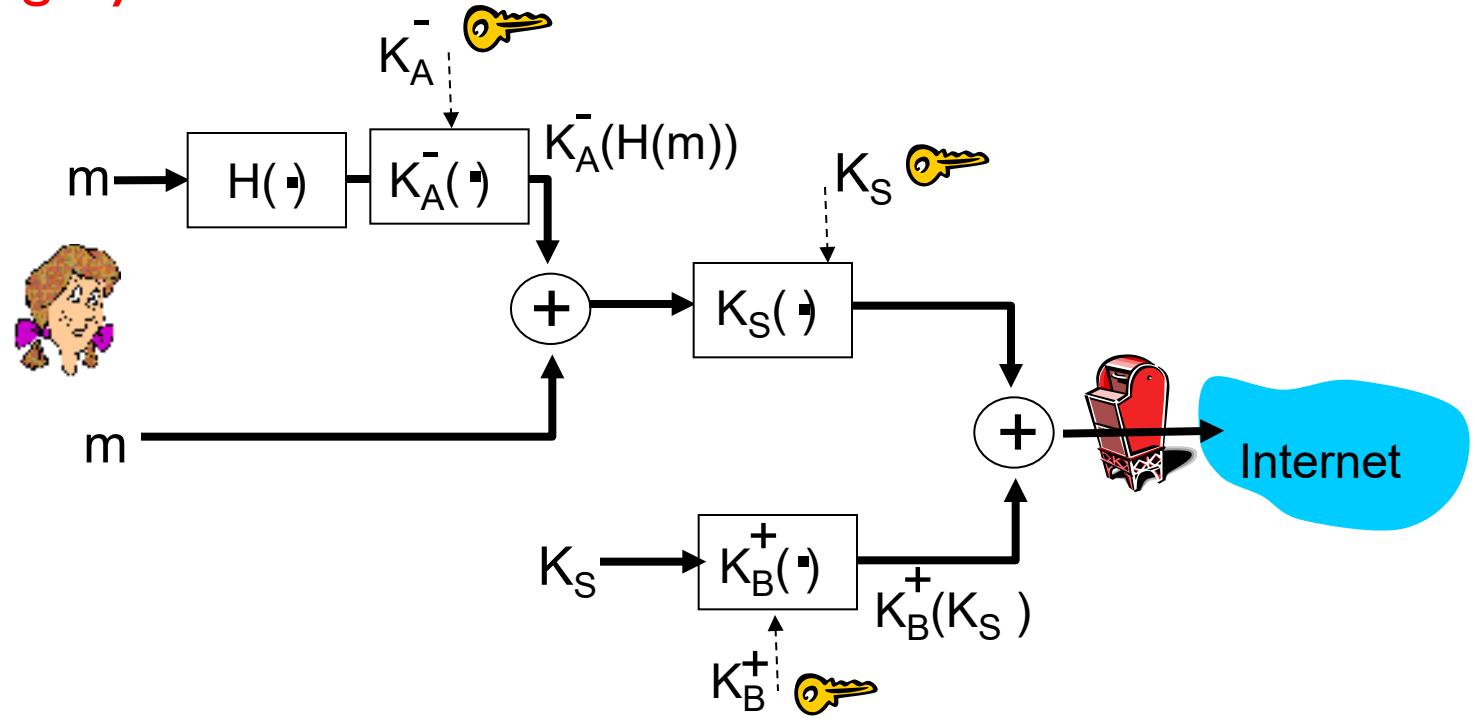
Alice wants to provide sender **authentication message integrity**



- Alice digitally signs message
- sends both message (in the clear) and digital signature

Secure e-mail (continued)

Alice wants to provide **secrecy**, **sender authentication**, **message integrity**.



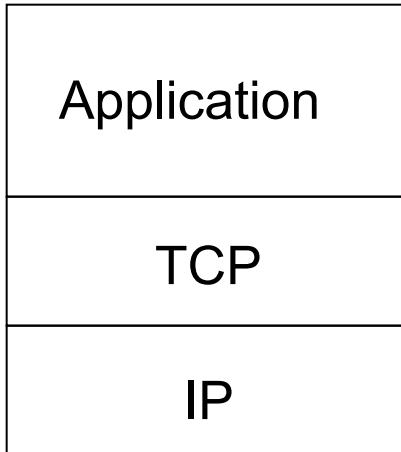
Alice uses three keys: her private key, Bob's public key, newly created symmetric key

Securing TCP connections: SSL

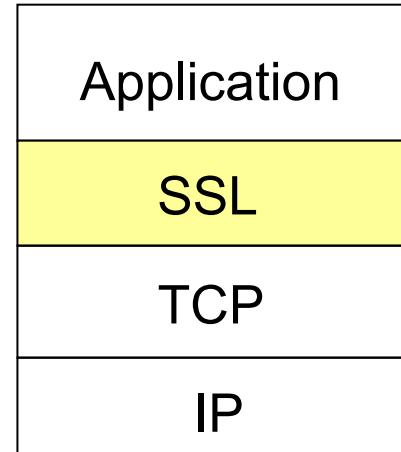
SSL: Secure Sockets Layer

- widely deployed security protocol
 - supported by almost all browsers, web servers
 - https
 - billions \$/year over SSL
- mechanisms: [Woo 1994], implementation: Netscape
- **New name -TLS: transport layer security, RFC 2246**
- provides
 - *confidentiality*
 - *integrity*
 - *authentication*
- original goals:
 - Web e-commerce transactions
 - encryption (especially credit-card numbers)
 - Web-server authentication
 - optional client authentication
 - minimum hassle in doing business with new merchant
 - available to all TCP applications
 - secure socket interface

SSL and TCP/IP



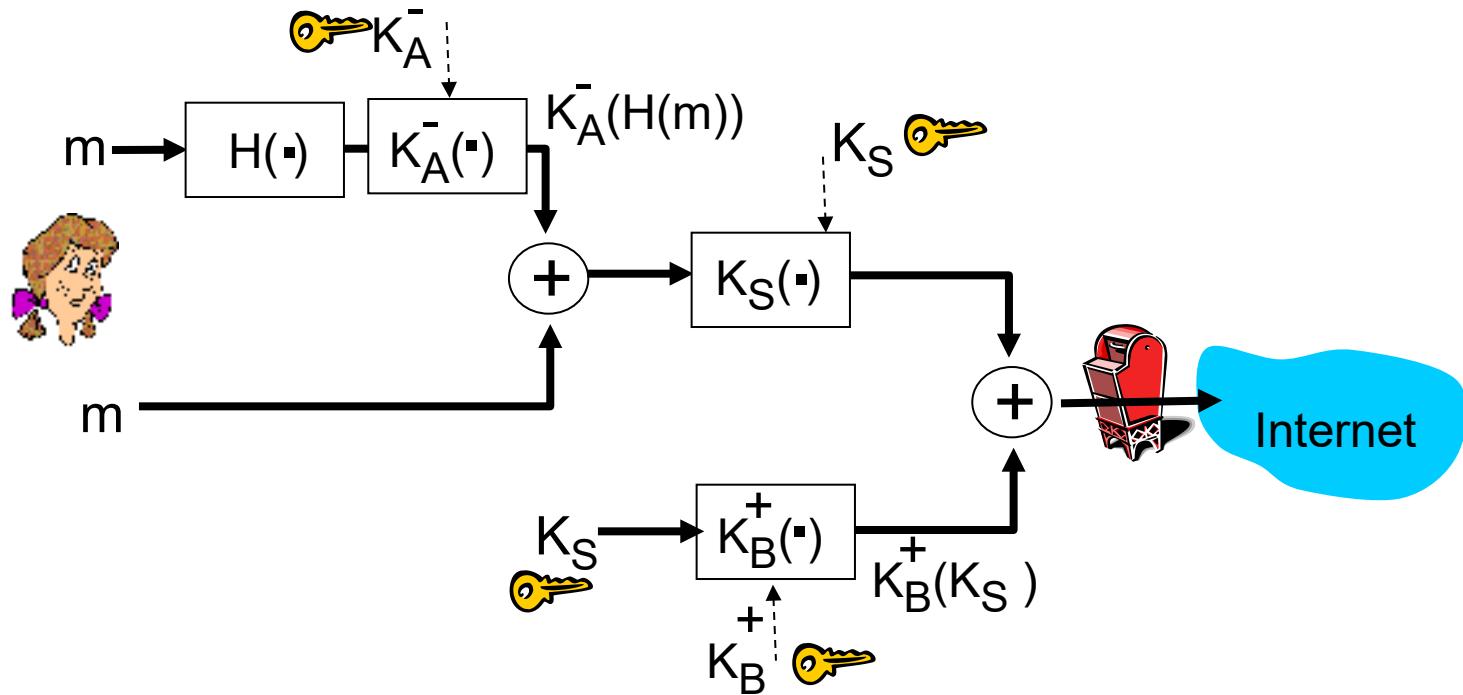
normal application



application with SSL

- SSL provides application programming interface (API) to applications
- C and Java SSL libraries/classes readily available

Could do something like PGP:

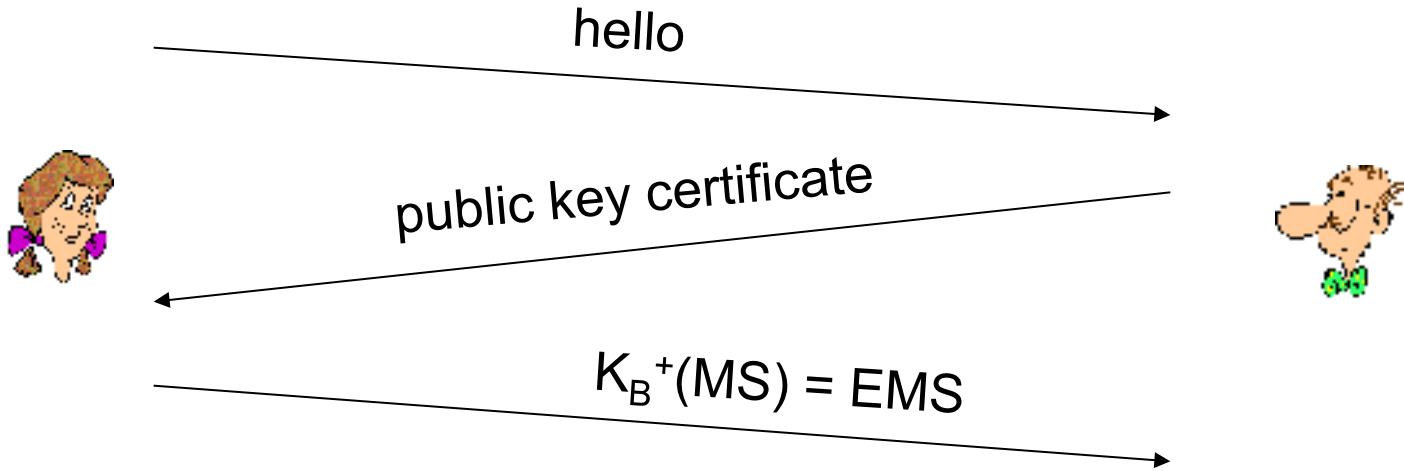


- but want to send byte streams & interactive data
- want set of secret keys for entire connection
- want certificate exchange as part of protocol: handshake phase

Toy SSL: a simple secure channel

- *handshake*: Alice and Bob use their certificates, private keys to authenticate each other and exchange shared secret
- *key derivation*: Alice and Bob use shared secret to derive set of keys
- *data transfer*: data to be transferred is broken up into series of records
- *connection closure*: special messages to securely close connection

Toy: a simple handshake



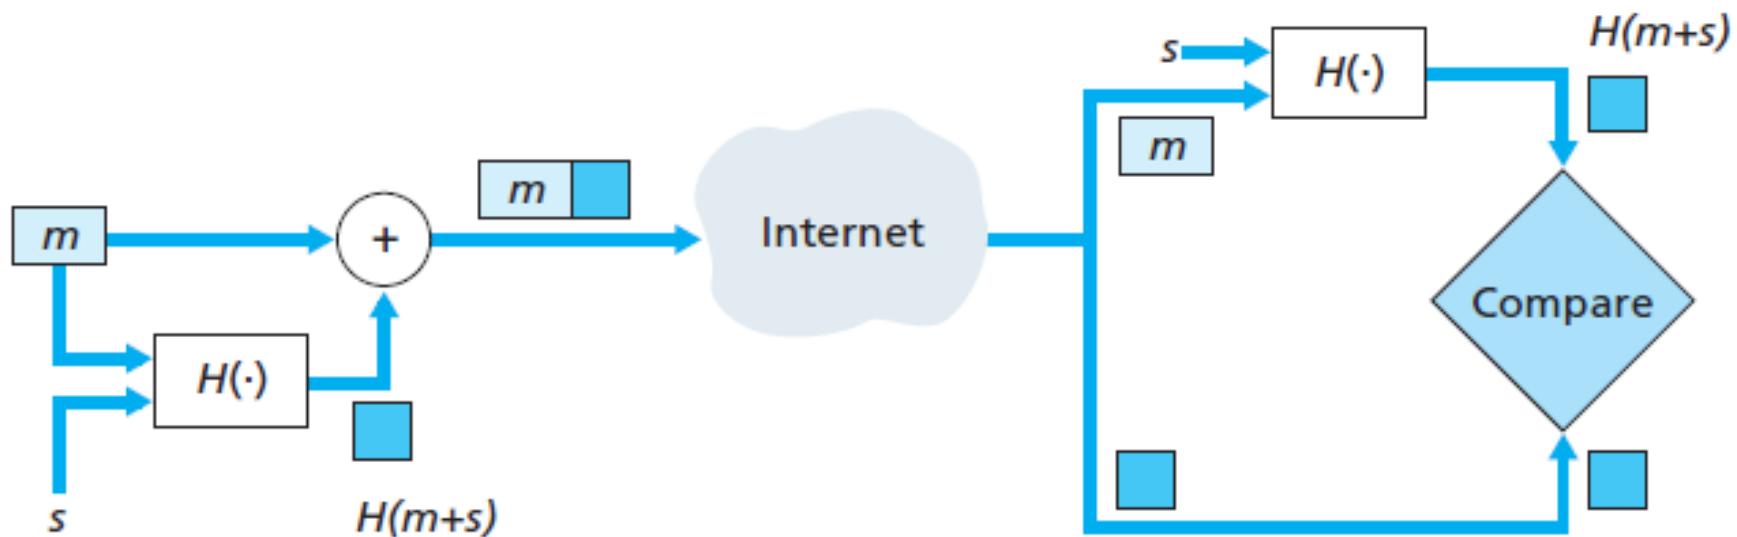
MS: master secret

EMS: encrypted master secret

Toy: key derivation

- considered bad to use same key for more than one cryptographic operation
 - use different keys for message authentication code (MAC) and encryption
- four keys:
 - K_c = encryption key for data sent from client to server
 - M_c = session MAC key for data sent from client to server
 - K_s = encryption key for data sent from server to client
 - M_s = session MAC key for data sent from server to client
- keys derived from key derivation function (KDF)
 - takes master secret and (possibly) some additional random data and creates the keys

Message authentication code (MAC)



Key:

- m = Message
- s = Shared secret

Toy: data records

- why not encrypt data in constant stream as we write it to TCP?
 - where would we put the MAC? If at end, no message integrity until all data processed.
 - e.g., with instant messaging, how can we do integrity check over all bytes sent before displaying?
- instead, break stream in series of records
 - each record carries a MAC
 - receiver can act on each record as it arrives
- issue: in record, receiver needs to distinguish MAC from data
 - want to use variable-length records



Toy: sequence numbers

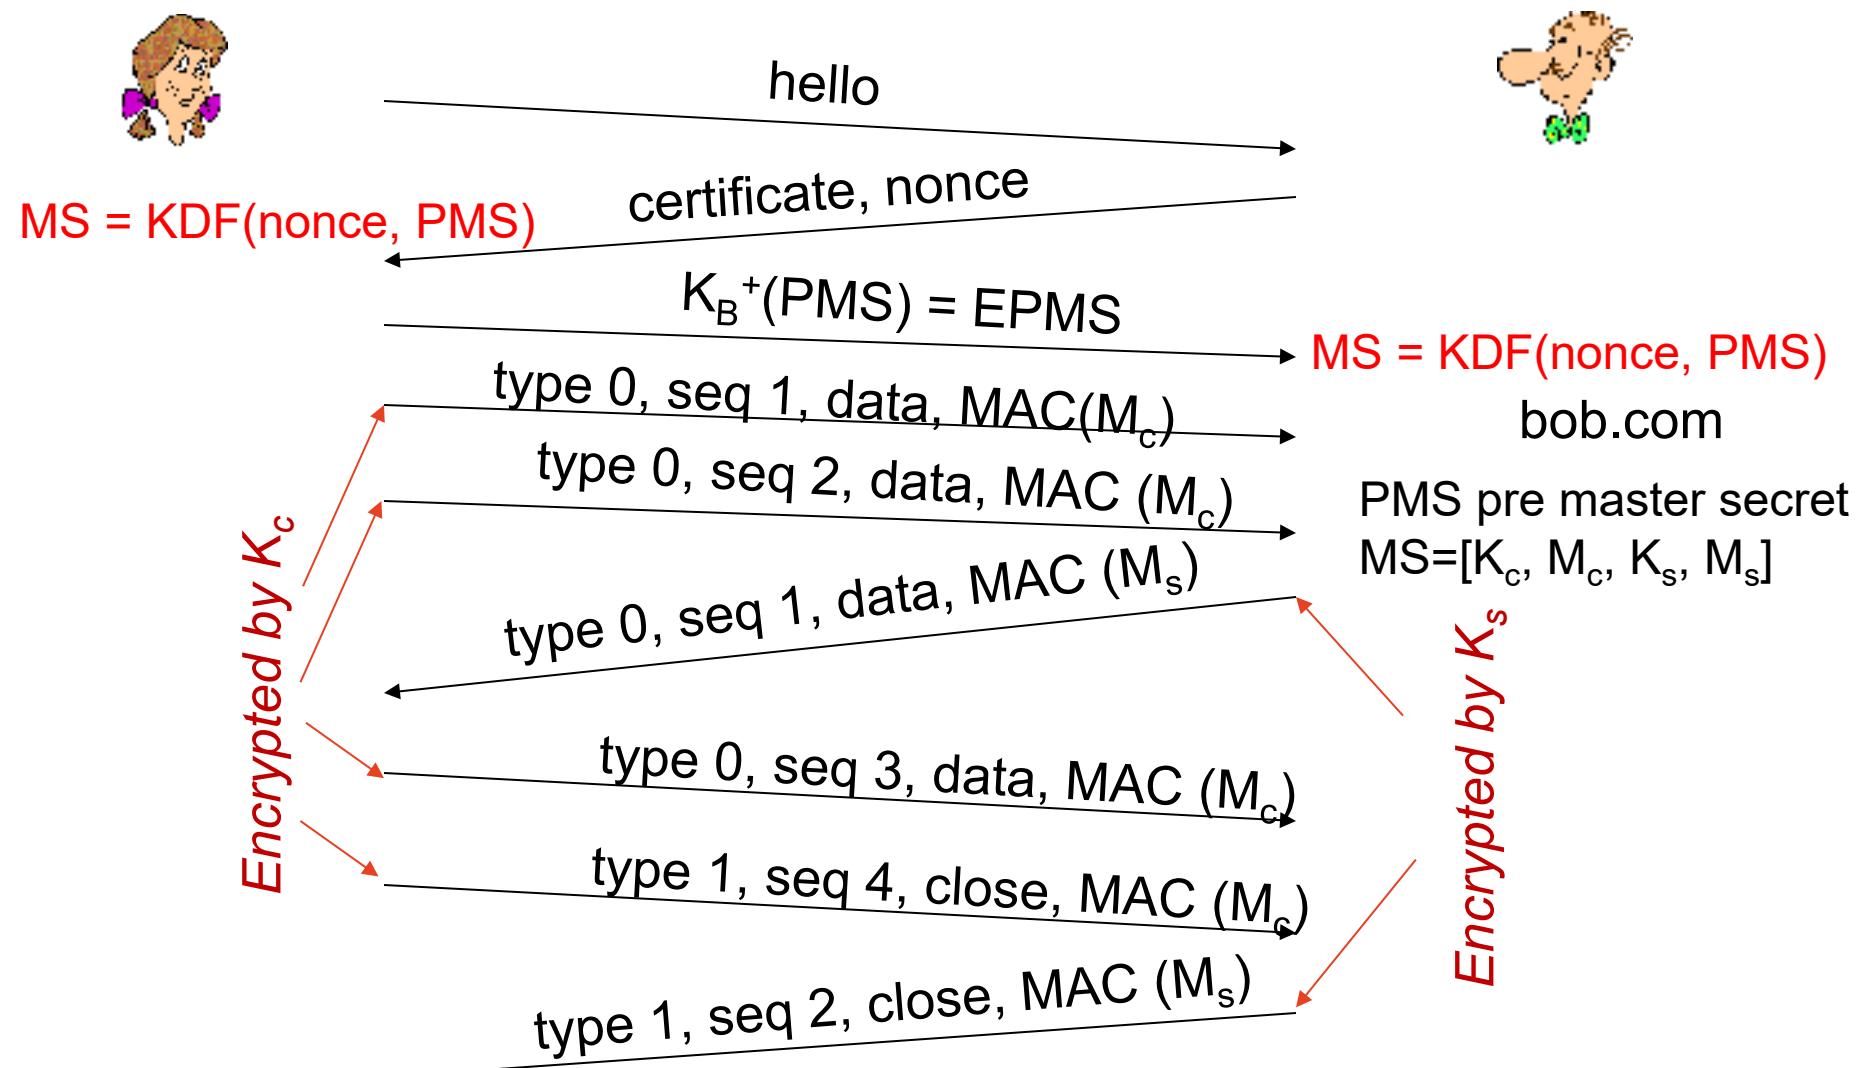
- **problem:** attacker can capture and replay record or re-order records
- **solution:** put sequence number into MAC:
 - $\text{MAC} = \text{MAC}(M_x, \text{sequence} \parallel \text{data})$
- **problem:** attacker could replay all records
- **solution:** use nonce

Toy: control information

- *problem:* truncation attack:
 - attacker forges TCP connection close segment
 - one or both sides thinks there is less data than there actually is.
- *solution:* record types, with one type for closure
 - type 0 for data; type 1 for closure
- $\text{MAC} = \text{MAC}(M_x, \text{sequence} \parallel \text{type} \parallel \text{data})$



Toy SSL: summary



SSL cipher suite

- cipher suite
 - public-key algorithm
 - symmetric encryption algorithm
 - MAC algorithm
- SSL supports several cipher suites
- negotiation: client, server agree on cipher suite
 - client offers choice
 - server picks one

common SSL symmetric ciphers

- DES – Data Encryption Standard: block
- 3DES – Triple strength: block
- RC2 – Rivest Cipher 2: block
- RC4 – Rivest Cipher 4: stream

SSL Public key encryption

- RSA

Real SSL: handshake (I)

Purpose

1. server authentication
2. negotiation: agree on crypto algorithms
3. establish keys
4. client authentication (optional)

Real SSL: handshake (2)

1. client sends list of algorithms it supports, along with client nonce
2. server chooses algorithms from list; sends back: choice + certificate + server nonce
3. client verifies certificate, extracts server's public key, generates pre_master_secret, encrypts with server's public key, sends to server
4. client and server independently compute encryption and MAC keys from pre_master_secret and nonces
5. client sends a MAC of all the handshake messages
6. server sends a MAC of all the handshake messages

Real SSL: handshaking (3)

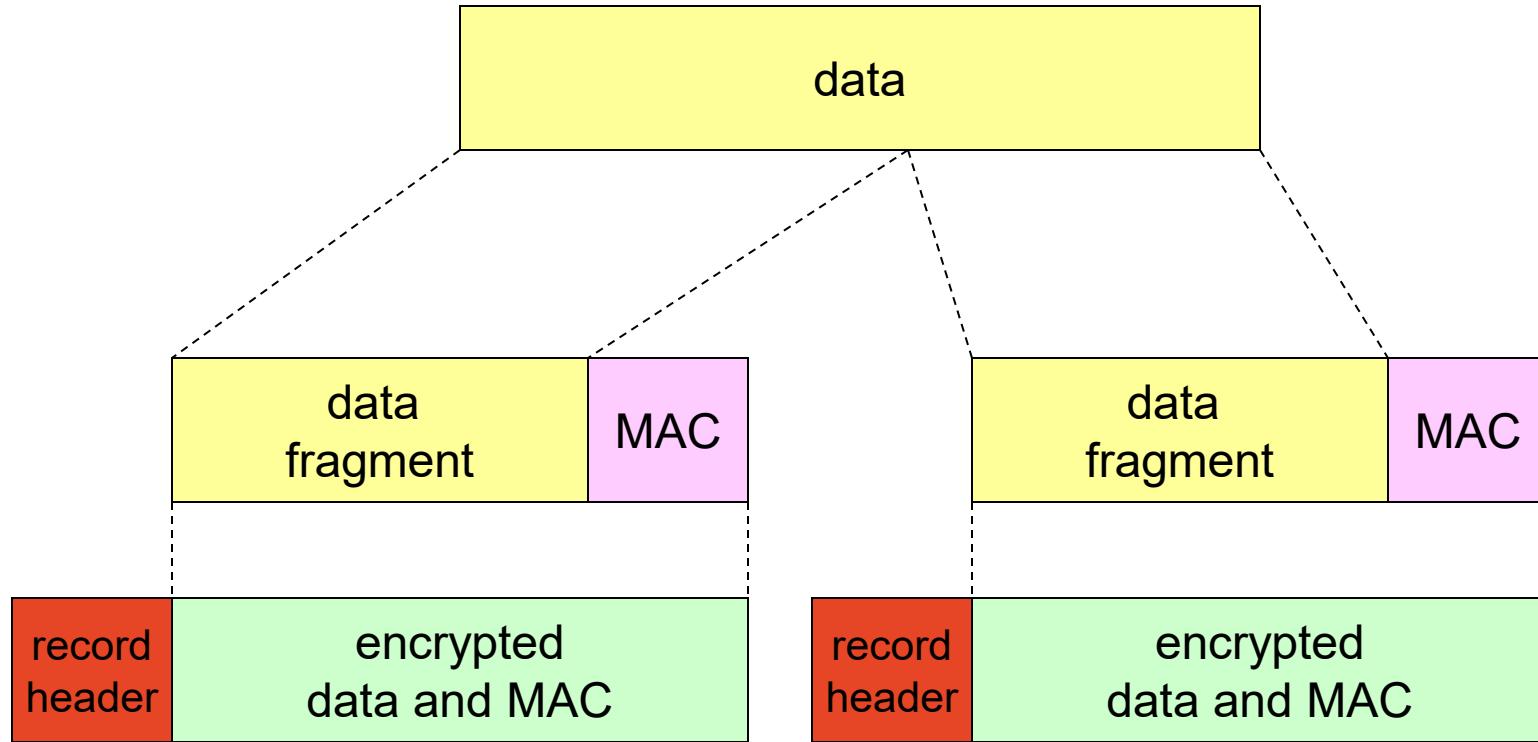
last 2 steps protect handshake from tampering

- client typically offers range of algorithms, some strong, some weak
- man-in-the middle could delete stronger algorithms from list
- last 2 steps prevent this
 - last two messages are encrypted

Real SSL: handshaking (4)

- why two random nonces?
- suppose Trudy sniffs all messages between Alice & Bob
- next day, Trudy sets up TCP connection with Bob, sends exact same sequence of records
 - Bob (Amazon) thinks Alice made two separate orders for the same thing
 - solution: Bob sends different random nonce for each connection. This causes encryption keys to be different on the two days
 - Trudy's messages will fail Bob's integrity check

SSL record protocol

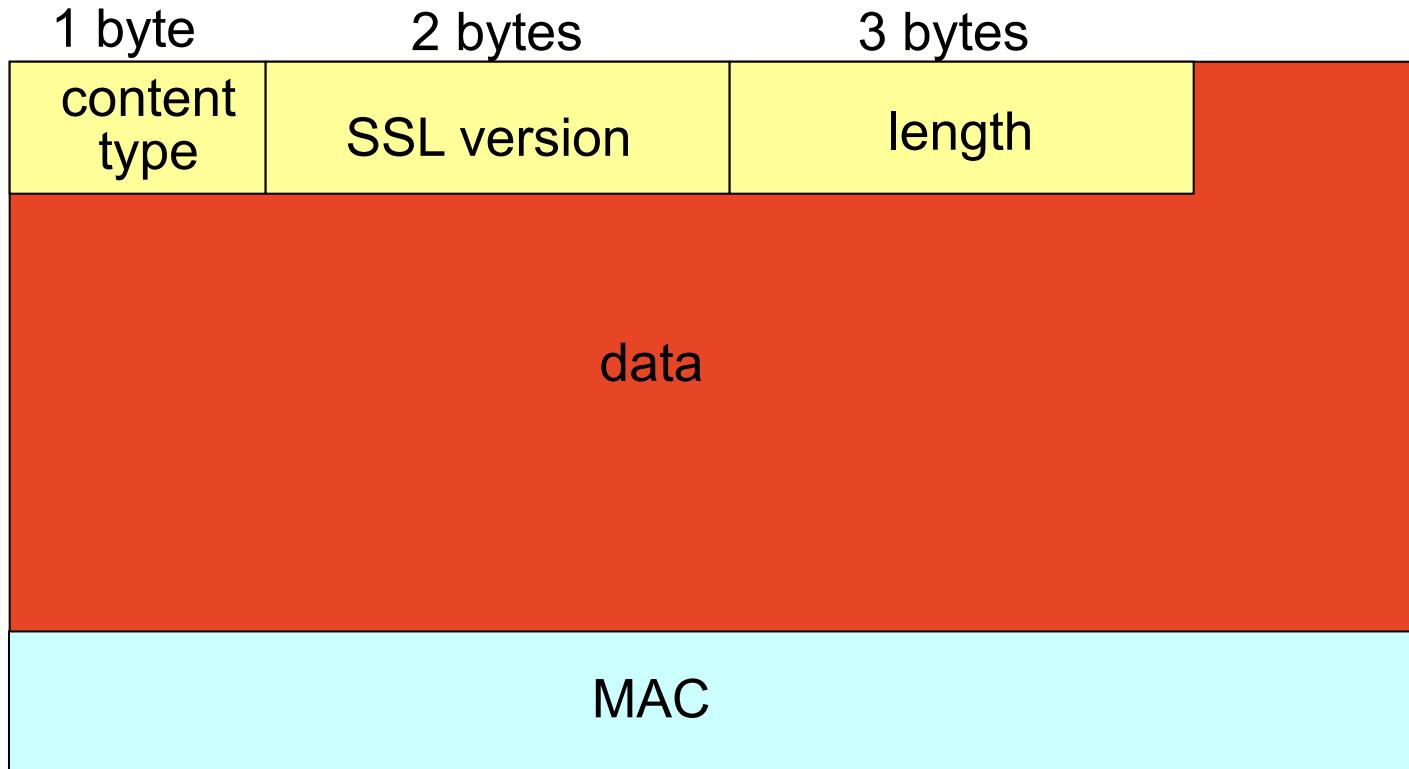


record header: content type; version; length

MAC: includes sequence number, MAC key M_x

fragment: each SSL fragment 2^{14} bytes (~ 16 Kbytes)

SSL record format



data and MAC encrypted (symmetric algorithm)

Nonce 1, Available cipher suite

handshake: ClientHello

N_1



Pre-master secret
(Random number)

N_3

$K_B^+(N_3)$

May or may not
check client's certificate

Master secret=KDF(N_1, N_2, N_3)

handshake: ServerHello

Nonce 2, Choose cipher suite

N_2

handshake: Certificate K_B^+

handshake: ServerHelloDone



K_A^+ handshake: ClientKeyExchange
ChangeCipherSpec

Master secret=KDF(N_1, N_2, N_3)

handshake: Finished

Session key and MAC key
are generated from master
secret

ChangeCipherSpec

handshake: Finished

application_data

application_data

Alert: warning, close_notify

TCP FIN follows