

**Push:** Source will send data to Target. Apache Kafka, MQTT. **Pull:** file download or database querying. **Poll:** Target periodically check the Source for change, if it changes, then read again. Web crawling.

**Bounded Data:** discrete datasets or batches; **Unbounded Data:** continuous stream

**Throughput, Latency and Scalability:** One-off tasks不需要考虑 latency;在 unbounded data stream 时要考慮 latency;当 data volume grow 时必须考虑 throughput 和 scalability

**ETL:**

加载前转换,适用于**结构化数据**,数据仓库, Data analysts excluded from ETL processes; 结构化数据

**ELT:** 加载后转换,适用于**大数据**,数据湖,非结构化数据

**Orchestration:** Manages the flow of data, automating extraction, transformation, and loading (ETL).

**Data Quality:** Believable(subjective), Value-added, Relevant, Accurate(objective), Interpretable

**Data Cleaning:** Missing(removed or imputed), Default, Incorrect(remove or adjust), Inconsistent(reconcile)

**csv.reader** reads rows into arrays

**csv.DictReader** reads rows into dictionaries

**Bar Charts** for Categorical/Nominal Data

Histograms for Numerical Data

Scatter Plots for comparing two continuing variables

自动生成 Schema: 约束弱, 没有 FK 和 Integrity Constraints; **自定义 Schema:** 约束强, 速度慢

特性 <sup>2</sup>	Operational Systems <sup>2</sup>	Analytical Systems <sup>2</sup>
访问	OLTP (Online Transaction Processing) <sup>2</sup>	OLAP (Online Analytical Processing) <sup>2</sup>
核心任务 <sup>2</sup>	事务处理 (Transaction) <sup>2</sup>	分析与统计 (Analysis & Aggregation) <sup>2</sup>
操作	插入 / 更新 / 删除 (Insert / Update / Delete) <sup>2</sup>	通常只读 (Read-only) <sup>2</sup>
用户 <sup>2</sup>	外部用户 (External users) <sup>2</sup>	业务分析师 / 数据科学家 <sup>2</sup>
查询	查少量记录 → Point Query <sup>2</sup>	扫描大量记录 → Aggregate Query (count, sum, avg, min, max, group by) <sup>2</sup>
方式 <sup>2</sup>	小、实时 <sup>2</sup>	大、历史数据全量 <sup>2</sup>

OLAP: Roll-up:向上汇总数据; Drill-down:小粒度; Slice: select a single dimension from a cube; Dice: create a smaller sub-cube from an OLAP cube. Pivot: 改变观察角度, 行列互换

Feature <sup>2</sup>	Database (OLTP) <sup>2</sup>	Data Warehouse (OLAP) <sup>2</sup>
Main read pattern <sup>2</sup>	Point queries (find individual records by key) <sup>2</sup>	Aggregations over a large number of records <sup>2</sup>
Main write pattern <sup>2</sup>	Create, update, and delete individual records <sup>2</sup>	Bulk import (ETL) or event stream 聚合、统计、查询 <sup>2</sup>
Typical human user <sup>2</sup>	End user of web/mobile application <sup>2</sup>	Internal analyst (decision support) <sup>2</sup>
Machine use	Checking if an action is authorized; 用户系统, 银行系统 <sup>2</sup>	Detecting fraud/abuse patterns; BI 和大数据分析 <sup>2</sup>
Data	Lates: state of data (current point in time) <sup>2</sup>	Historical events accumulated over time <sup>2</sup>
representatives <sup>2</sup>		
Dataset size <sup>2</sup>	Gigabytes to terabytes <sup>2</sup>	Terabytes to petabytes <sup>2</sup>

Feature <sup>2</sup>	Data Lake (ETL) <sup>2</sup>	Data Warehouse (ETL) <sup>2</sup>
Data type <sup>2</sup>	Structured, semi-structured (JSON, XML), unstructured (images, audio, logs) <sup>2</sup>	Primarily Structured data (tables, SQL) <sup>2</sup>
Schema <sup>2</sup>	Schema-on-read (defined at query time) <sup>2</sup>	Schema-on-write (defined before writing) <sup>2</sup>
Storage & Processing <sup>2</sup>	Distributed cheap storage (HDFS, S3). Specialized storage, high performance, low cost <sup>2</sup>	Raw data stored; flexible processing for analytics, MLL and exploration <sup>2</sup>
Typical use cases <sup>2</sup>	Big data analysis, machine learning, real-time processing <sup>2</sup>	Cleaned, transformed, indexed, optimized for analytics <sup>2</sup>
Agility <sup>2</sup>	More agile as it accepts raw data without a predefined structure <sup>2</sup>	Less agile due to predefined schema <sup>2</sup>

特性 <sup>2</sup>	Star Schema (星型模式) <sup>2</sup>	Snowflake Schema (雪花模式) <sup>2</sup>
结构 <sup>2</sup>	一个事实表 + 多个非规范化维度表 (维度拆分维度表) <sup>2</sup>	一个事实表 + 多个规范化维度表 (维度拆分维度表) <sup>2</sup>
维度表	宽大、包含完整信息 <sup>2</sup>	层次复杂、减少冗余 <sup>2</sup>
适用场景	BI 报表、实时看板、非技术用户查询 <sup>2</sup>	数据量极大、维度层次复杂、对存储和数据一致性要求高 <sup>2</sup>
查询性能	相对低, 需要多次 JOIN <sup>2</sup>	高。现代云数仓对宽优化好 <sup>2</sup>
维护难度	易于理解与维护 <sup>2</sup>	较复杂, 理解和维护难度大 <sup>2</sup>
存储占比	占用较多 <sup>2</sup>	节省存储空间 <sup>2</sup>
设计理念	以空间换时间 → 提升查询速度 <sup>2</sup>	以时间换空间 → 提升存储效率和数据整洁 <sup>2</sup>

**Data Warehouse:** contains a read-only copy of data from various OLTP; Can be both source or sink; 有 Data Silo; **Fact Table** 通常包含 **度量** 和 **外键** 维度, 通常都是数字, 可以进行聚合; **Dimension Table**.

特性 <sup>2</sup>	Web Scraping <sup>2</sup>	Web API <sup>2</sup>
返回数据类型	HTML / 网页内容 (可能还是 JSON / XML / 半结构化数据)	JSON / XML / 半结构化数据
数据结构	半结构化半结构化, 需要解析 DOM <sup>2</sup>	结构化或半结构化, 直接可用
稳定性	网站可能随时会坏掉坏掉 <sup>2</sup>	网站稳定, 只要 API 不变
速度	相当慢, 基本都解析半天 <sup>2</sup>	快, 直接调用 API <sup>2</sup>
依赖技术	解析 HTML / DOM / XPath / CSS	HTTP 请求 + JSON/XML 解析
应用场景	无 API 或网页需要抓取 <sup>2</sup>	官方提供 API, 可直接获取数据

**Web Scraping: Reconnaissance** (Check terms-of-service and robots.txt; if in doubt, contact owner); **Webpage Retrieval; Data Extraction; Data Cleaning transformation; Data Storage, Analysis URL:** protocol://site/path\_to\_resource

**Scrapy:** 可编程的爬虫框架

**Selenium:** 可编程的浏览器

**HTTP:** Head 部分用户看不到, Body 部分是网页内容. DOM 是一个 element tree, Sibling tags 有从 left to right! 方法有下面 2 个

请求方式 <sup>2</sup>	参数位置 <sup>2</sup>	常用场景 <sup>2</sup>
GET <sup>2</sup>	URI (key=value)	获取数据 (搜索、下载网页)
POST <sup>2</sup>	请求体 (body)	提交数据 (登录、上传表单)

**RESTful APIs:** Stateless, 使用上面 2 个 HTTP 方法; **XML Web Services:** 基于 SOAP, 也是一种 Web API

**Semi-structured data:** Nesting, Heterogeneous collections; HTML, XML and JSON belong to it.

**HTML:** 主要是用于网页设计和展示, pre-defined tags, 宽松语法规, 浏览器容错率高,即使写错也能渲染

**XML:** 主要用于数据交换和存储, user-defined tags, 严格语法规, 必须成对闭合</title></title>, 区分大小写, Empty elements</AUD/>

**Well-Formed XML:** 文档符合 XML 的基本语法规则, **Valid XML:** 文档不仅格式正确 (well-formed), 还要符合一个预定义的约束规则

**DTD:** 定义 XML 文档的语法规则, 70 或 10\*或多或少, 至少 1 次属性 genre 在<>里面, 子元素在<>之外.

<book> 元素由 title, author, price 组成. 'title' 必须出现 1 次; 'author' 表示至少一个作者; 'price' 表示价格可选 (0 或 1 次)

<ELEMENT book (title,author+,price?)>

<!--ATTLIST book genre CDATA #REQUIRED-->

<!--ATTLIST book genre CDATA #REQUIRED-->

**Feature<sup>2</sup>**

**Transactional data streams:** log interactions between entities: 记录实体之间发生的“事件” / 交互”的数据流,比如 credit card purchases by consumers from merchants

**Measurement data streams:** monitor evolution of entity states: 持续监测某个实体“状态随时间变化”的数据流,IP network

Database Systems (DBMS) <sup>3</sup>	Data Stream Systems (DSMS) <sup>3</sup>
<b>Persistent relations<sup>3</sup></b>	<b>Transient streams<sup>3</sup></b>
set/bag of tuples <sup>3</sup>	Sequence of tuples 有序 <sup>3</sup>
Bounded by stored data, 静态 <sup>3</sup>	Unbound data, 动态 <sup>3</sup>
modifications <sup>3</sup>	appends <sup>3</sup>
transient / one-time <sup>3</sup>	continuous / persistent <sup>3</sup>
exact <sup>3</sup>	approximate <sup>3</sup>

由 查询处理器和物理数据流设计 决定<sup>3</sup> 不可预测, 因为数据特性和到达模式不确定  
第一种 DSMS: Publish/Subcribe Messaging System  
ability to consume constant stream of events, with QoS guarantees.  
布者不直接把消息发给具体接收者, 而是按“主题”(topic)“发布”到 Broker: 订阅者只从 broker 收到自己订阅的主题的消息。常见的有: Apache Kafka(Topics in Kafka are multi-subscribe), MQTT. A broker mediates communication between producers and consumers, providing decoupling, scalability, and reliable message delivery in distributed systems

能力 <sup>3</sup>	解法 <sup>3</sup>
1. Publish and subscribe to streams of records <sup>3</sup>	支持应用序 发布消息到主题 (topic), 以及 订阅自己感兴趣的消 息。实现数据分发。 <sup>4</sup>
2. Can store streams of records in a fault-tolerant way <sup>3</sup>	可以容错持久化消息 保证在系统故障时消息不丢失。典型实现: 消息在 Broker 中持久化 (如 Kafka 的 log 机制). <sup>4</sup>
3. Lets apps process streams of records as they occur <sup>3</sup>	应用可以实时处理到达的数据流 无需等待数据全部存储再分析, 支持事件驱动处理和延迟响应。 <sup>4</sup>

**MQTT:** 第一个 is subscriber, 第二个 is publisher  
subscribe.callback(print\_msg, "MyTopic", hostname=broker)  
print\_msg: 回调函数 = 系统帮你调用的函数, 不是你手动调用  
MyTopic: 订阅 MyTopic 这个主题, 当有新消息到来时, 用  
print\_msg 这个函数来处理消息。Topics 可以是 sensitive  
hostname=broker: 是从 broker 订阅, 而不是 Publisher  
publish.single("MyTopic", "TheMessage", hostname=broker)  
By default, MQTT broker does not store any message. 只有在线的订  
阅者 才能收到消息, 但是可以设置为保留

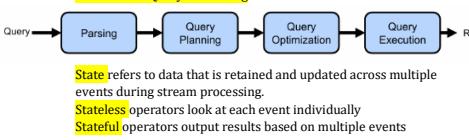
Apache Kafka (Distributed Streaming Platform) <sup>3</sup>	MQTT (IoT Messaging Protocol) <sup>3</sup>
For large-scale stream data storage and real-time processing <sup>3</sup>	For resource-constrained devices with low-bandwidth, real-time communication <sup>3</sup>
Distributed commit log (append-only log); 保 证 topic 不会丢失 <sup>3</sup>	Broker-based message forwarding model <sup>3</sup>

Persistent by default data stored on disk.	Transient by default messages usually deleted after forwarding (optional retained messages or offline cache) <sup>3</sup>
Push-based: consumers fetch messages from the queue in order, control their own offset <sup>3</sup>	Push-based: broker actively pushes messages to subscribers for real-time delivery <sup>3</sup>
Strictly guaranteed order within a partition <sup>3</sup>	Depends on network and QoS level, global ordering hard to guarantee at scale <sup>3</sup>

Very high data replicated across multiple nodes, single-node failures do not affect safety<sup>3</sup>

大规模、可靠、持久化、企业级流处理;<sup>4</sup>  
金融交易风控, 网站日志分析<sup>3</sup>

**Quality of Service**  
At most once (QoS 0: Default): every message is processed at most once, but no guarantee: 尽力而为, 发了就不管了, 最快  
At least once (QoS 1): no message will get ignored; 一定送到, 但可能送两次  
Exactly once (QoS 2): system guarantees that every message is processed exactly once; 一定送到, 而且只送一次, 最慢  
第二种 DSMS: Data Stream Processors  
ability to execute query processing on stream data with clear semantics.  
**Data Stream Query Processing**



Stateful Applications <sup>3</sup>	Stateless Applications <sup>3</sup>	Speed-Up <sup>3</sup>	Scale-Up <sup>3</sup>
Retain user session information across requests <sup>3</sup>	Do not retain user state between requests <sup>3</sup>	数据量不变, 资源增加, 运行时间应按比例减少 <sup>3</sup>	数据量和资源同时按比例增加, 运行时间保持不变 <sup>3</sup>
Less scalable, require complex load balancing and session management <sup>3</sup>	Highly scalable, each request is independent <sup>3</sup>	固定 <sup>3</sup>	随着资源增加而增加 <sup>3</sup>
Lower Fault Tolerance, server failure may cause session loss unless replication is used <sup>3</sup>	Higher Fault Tolerance, server failure does not affect sessions <sup>3</sup>	增加计算资源 (CPU/节点/内存) <sup>3</sup>	增加计算资源与数据量成比例 <sup>3</sup>
More Resource Need, require memory and processing for session handling <sup>3</sup>	Less, no session data to manage <sup>3</sup>	减少处理时间 <sup>3</sup>	响应时间不变 <sup>3</sup>
More complex, need careful session and state management <sup>3</sup>	Simpler, no need to manage state across requests <sup>3</sup>	时间按比例减少, 实际收 overhead 影响	响应时间保持恒定 <sup>3</sup>

**Aggregation** – need first to define a window on the data stream to process.

**Agglomerative Window:** 增量地 aggregates incoming data and updates the result continuously without storing all events in the window.

**Sliding Window:** fixed size, moves forward at a fixed interval, overlap.

**Tumbling Window:** fixed-size, non-overlapping, 是特殊的 Sliding Window.

**Tuple-Count-Based Windows:** 包含 Sliding Window 和 Tumbling Window.

基于 tuple 的数量进行 window 的划分, 有个问题就是 Tie, 即假设 3 条数据时间戳都是 10:00:01, 但窗口大小刚好在它们之间中断。系统这次可能把 A 放入窗口, 下次可能把 B 放入窗口, 解决办法为在时间戳之外, 增加一个唯一且稳定的字段 tie-breaker.

**Punctuation-Based Windows:** 基于 punctuation 划分窗口, 窗口长度不固定。

**Event time:** 事实真正发生的时间。有个 Watermarks 的机制可以确保这个, 它是一个时间戳  $T$ , 代表系统声明: “所有 Event Time  $\leq T$  的数据都到了齐了”。因此可以通过窗口的 last timestamp  $\leq$  watermark 来判断是否输出结果。

**Ingestion time:** 数据进入系统边界的时刻。

**Processing time:** 数据被算法算子执行的时刻。



**Conventional Table Joins:** 将两个已经存在的、静态的静态表 (Table A 和 Table B) 通过公共键 (Join Key) 合并。

**Enrichment:** 将实时的“事实层”与相对静态的数据通过一个 window 进行匹配。这个 window 确保了即使维度表在不断变化 (Stored Data updates), 流数据也能找到它“应有的”历史镜像。因为查询数据库会导致延迟, 有 2 种办法解决 In-memory Storage: 通过把常用的存储在内存中提升效率。External data Storage: 通过 index 等提升查询效率。

**Stream-to-stream Joining:** 将两个实时产生的数据流, 在某个时间窗口内进行匹配。

**数据流处理器 Recall:** Kafka or MQTT are not data stream processors.

**Apache Spark<sup>3</sup>**

**Apache Flink<sup>3</sup>**

Set-oriented data, Lazy<sup>3</sup>

**evaluation principle:** Plan gets actually executed by Flink only when `env.execute()` is called<sup>3</sup>

Processing at Separate stages<sup>3</sup>

Optimizer is SparkSQL<sup>3</sup>

RDD<sup>3</sup>

**Micro-batching (DStream)<sup>3</sup>**

**Higher Latency<sup>3</sup>**

Use Case: Micro-batch<sup>3</sup>

fixed, sliding, tumbling<sup>3</sup>

**Agglomerative<sup>3</sup>**

**Scale-Up:** To scale with increasing load, buy more powerful, larger hardware

**Scale-Out:** need to scale-out to a cluster of multiple servers (nodes). shared-nothing architecture

Transformations by iterating over collections with pipelining<sup>3</sup>

“透明的” One NameNode, Multiple DataNodes. Each block is replicated across multiple nodes default 3.

**MapReduce** 是实现这些目标的经典处理模型. Map: filtering and sorting students by first name into queues. Reduce: summary operation counting the number of students in each queue. Hadoop = HDFS+YARN+MapReduce

**Loosely-Coupled System:** Separation of computer and storage nodes. The more nodes, the higher the probability of some failure. Latency can increase.

**Data Architecture:** Architecture first, technology second.

**Operational Architecture:** What needs to be done?

**Technical Architecture:** How data is ingested, stored, transformed, and served.

**Batch-driven:** 延时处理, 大数据架构的一种。

**Data Lake:** Dump all of your data, structured and unstructured, into a central location. 容易导致 data swamp, 缺乏治理、元数据和质量控制, 数据湖变得混乱、难以理解、几乎无法使用的状态。

**Feature Store Architecture:** data pipeline server for training and serving machine learning models. Transformation, Storage, Serving, Monitoring, Feature Registry. **Offline Serving:** Access historical data; **Online Serving:** Provide fresh feature. **Batch Transform** 静态数据 **Streaming Transform** 动态数据

**On-Demand Transform** 只在预测的时候用. 替代方案: **Custom ETL Pipelines:** High flexibility: Increased complexity, lacks consistency. **Data Warehouses & Lakes:** Easy Integrated; Not suit real time ML. **In-House Feature Management Systems:** Tailored;扩展性差.

**Lambda Architecture:** Address latency by creating batch/cold layer 存入 serving layer 确保了准确性 and speed/hot layer 直接被 analytics 使用确保低 latency. 用户也可以通过 serving layer 结合两者。

**Event-driven:** 时时处理, 大数据架构的一种. **Data Stream Processing Architectures:** 例子 **Kappa Architecture**, Only one path.

Scope	Event-driven (Event-driven)
设计哲学	效率优先, 追求整体吞吐量
数据视图	静态 (针对已保存的数据块)
典型代表	ETL 报表, 工资发放, 财务对账
系统解耦	弱 (通常依赖中心化数据库)
复杂度	较低, 错误容易重跑
Continuous Integration (CI):	Automatically testing and integrating new code into a shared repository
Continuous Deployment (CD):	Automatically deploying tested code production environments.
DataOps	is a set of collaborative data management practices designed to speed up data delivery, maintain quality, and foster collaboration.
-Key goal:	Break Silo by Unifying Data, Collaborative Framework
-Key Functions:	<ul style="list-style-type: none"> <li>Pipeline Orchestration: 负责统一调度和管理数据与 ML 流水线</li> <li>Data Quality Monitoring: 持续监控数据和特征的完整性、准确性与一致性, 及时发现缺失值、异常分布</li> <li>Governance and Security: 确保数据使用符合组织规范与安全要求</li> <li>Self-Service Data Access: 通过标准化接口、目录和文档, 使数据科学家和业务人员能够自助发现、理解并使用数据与特征, 减少对数据工程团队的依赖</li> </ul>
-Lifecycle:	Plan, Dev, Integrate, Test, Deploy, Monitor
-Data Curation:	Automate data cleansing, transformation, and standardization to ensure high-quality data.
-Master Data Management:	确保数据一致性。
-The Five Pillars:	Freshness, Distribution(acceptable ranges), Volume(Monitor Missing), Schema(Track structural changes), Lineage(data move/transform across system)

**3 种 Scaling ML 的方法如下**

**Feature Store:** TFIDF 是一种特征提取方法.

**Data to ML/Data to Computation:** 适合小数据, 不好扩展. **Scale-up** ML to Data to Computation: 适合大数据, 好扩展. 典型代表是 **MADlib**(In-database analytic), 它还可以通过 ARIMA 支持 time series(只有 point) 使用 Greenplum 则是 a shared-nothing database using postgres per node 或 PostgreSQL. **Scale-out**

PII: Information that, when used alone or with other relevant data, can identify an individual.

**Sensitive information:** is personal information that includes

information or an opinion about an individual. 比如信仰, 性取向。

**Data Minimalism:** best way. Avoid collecting unless necessary. 不需要的历史数据要删除.

**Least Privilege:** grant human/machine only the necessary access. 视图可达

**组件加密 (At Rest)**

静止 (停在磁盘上)  $\Rightarrow$  流动 (在网络中传输)  $\Rightarrow$

主要威胁: 磁盘被窃、数据未被双录、跨网、中间人攻击 (MITM)、非法劫持<sup>3</sup>

核心技术: AES-256, RSA, KMS, TDE<sup>3</sup>

性能损耗: 主要在数据插入 / 读取时的磁盘 I/O<sup>3</sup>

主要挑战: 证书颁发机构 (CA) 信任链的安全性<sup>3</sup>

**3 - 2 - 1 rule:** At least 3 copies On 2 different media At least 1 off-premise

**RPO:** how much data you can afford to lose; **RTO:** how fast you need to recover.

An RPO of 1 hour means data backups must be made at least every hour

An RTO of 2 hours means the system must be fully restored within that time.

**On-Premise Backup:** Local storage devices like NAS.

**Cloud backups:** may take longer to restore than local backups.

**Best Practice:** Combining local replication with remote cloud backups ensures both high availability and disaster recovery readiness.

**Crash Recovery:** Using log. **Disaster Recovery:** using log and backup.

**Choosing Technologies across the Data Engineering Lifecycle by using MOBIST:** Monolith/Modular, Optimizing cost, Build vs Buy, Interoperability 互操作性 API, Location 本地/云, Immutability, Speed to market, Team size

**Monolith:**

自包含系统<sup>3</sup> 采用最佳技术, 通过 API 通信, 各组件各司其职<sup>3</sup>

简单, 一切在同一处 (更少移动部件, 减少组件替换, 技术更新/升级, 适合变化较小工具)<sup>3</sup>

脆弱, 移动困难<sup>3</sup> 系统数据量众多, 维护复杂, 更新/发布耗时

不可靠, 难以维护<sup>3</sup>

UI, 业务层、数据接入全部在同一系统<sup>3</sup> 各层可独立拆分, 模块化实现<sup>3</sup>

Scope	Event-driven (Event-driven)
设计哲学	效率优先, 追求整体吞吐量
数据视图	速度优先, 追求响应时效
典型代表	ETL 报表、工资发放、财务对账
系统解耦	强 (通过消息中间件如 Kafka 实现异步)
复杂度	较高, 需处理乱序和补偿机制
Continuous Integration (CI):	Automatically testing and integrating new code into a shared repository
Continuous Deployment (CD):	Automatically deploying tested code production environments.
DataOps	is a set of collaborative data management practices designed to speed up data delivery, maintain quality, and foster collaboration.
-Key goal:	Break Silo by Unifying Data, Collaborative Framework
-Key Functions:	<ul style="list-style-type: none"> <li>Pipeline Orchestration: 负责统一调度和管理数据与 ML 流水线</li> <li>Data Quality Monitoring: 持续监控数据和特征的完整性、准确性与一致性, 及时发现缺失值、异常分布</li> <li>Governance and Security: 确保数据使用符合组织规范与安全要求</li> <li>Self-Service Data Access: 通过标准化接口、目录和文档, 使数据科学家和业务人员能够自助发现、理解并使用数据与特征, 减少对数据工程团队的依赖</li> </ul>
-Lifecycle:	Plan, Dev, Integrate, Test, Deploy, Monitor
-Data Curation:	Automate data cleansing, transformation, and standardization to ensure high-quality data.
-Master Data Management:	确保数据一致性。
-The Five Pillars:	Freshness, Distribution(acceptable ranges), Volume(Monitor Missing), Schema(Track structural changes), Lineage(data move/transform across system)

Scope	Stream-first architecture with batch mode
设计哲学	Large-scale storage (HDFS) & processing
数据视图	MapReduce jobs across a cluster
典型代表	Apache Hadoop
系统解耦	弱 (通常依赖中心化数据库)
复杂度	较高, 需处理乱序和补偿机制
Continuous Integration (CI):	Automatically testing and integrating new code into a shared repository
Continuous Deployment (CD):	Automatically deploying tested code production environments.
DataOps	is a set of collaborative data management practices designed to speed up data delivery, maintain quality, and foster collaboration.
-Key goal:	Break Silo by Unifying Data, Collaborative Framework
-Key Functions:	<ul style="list-style-type: none"> <li>Pipeline Orchestration: 负责统一调度和管理数据与 ML 流水线</li> <li>Data Quality Monitoring: 持续监控数据和特征的完整性、准确性与一致性, 及时发现缺失值、异常分布</li> <li>Governance and Security: 确保数据使用符合组织规范与安全要求</li> <li>Self-Service Data Access: 通过标准化接口、目录和文档, 使数据科学家和业务人员能够自助发现、理解并使用数据与特征, 减少对数据工程团队的依赖</li> </ul>
-Lifecycle:	Plan, Dev, Integrate, Test, Deploy, Monitor
-Data Curation:	Automate data cleansing, transformation, and standardization to ensure high-quality data.
-Master Data Management:	确保数据一致性。
-The Five Pillars:	Freshness, Distribution(acceptable ranges), Volume(Monitor Missing), Schema(Track structural changes), Lineage(data move/transform across system)

Scope	Stream processing
设计哲学	Runs MapReduce jobs for both batch and stream processing
数据视图	In-memory processing for real-time data
典型代表	Apache Flink
系统解耦	强 (通过消息中间件如 Kafka 实现异步)
复杂度	高, 需处理乱序和补偿机制
Continuous Integration (CI):	Automatically testing and integrating new code into a shared repository
Continuous Deployment (CD):	Automatically deploying tested code production environments.
DataOps	is a set of collaborative data management practices designed to speed up data delivery, maintain quality, and foster collaboration.
-Key goal:	Break Silo by Unifying Data, Collaborative Framework
-Key Functions:	<ul style="list-style-type: none"> <li>Pipeline Orchestration: 负责统一调度和管理数据与 ML 流水线</li> <li>Data Quality Monitoring: 持续监控数据和特征的完整性、准确性与一致性, 及时发现缺失值、异常分布</li> <li>Governance and Security: 确保数据使用符合组织规范与安全要求</li> <li>Self-Service Data Access: 通过标准化接口、目录和文档, 使数据科学家和业务人员能够自助发现、理解并使用数据与特征, 减少对数据工程团队的依赖</li> </ul>
-Lifecycle:	Plan, Dev, Integrate, Test, Deploy, Monitor
-Data Curation:	Automate data cleansing, transformation, and standardization to ensure high-quality data.
-Master Data Management:	确保数据一致性。
-The Five Pillars:	Freshness, Distribution(acceptable ranges), Volume(Monitor Missing), Schema(Track structural changes), Lineage(data move/transform across system)

Scope	Streaming & Real-time Processing





<tbl\_r cells="2" ix="5" maxcspan="1" maxrspan="1