

COMP5339: Data Engineering

W7: Temporal Data Engineering

Presented by

Uwe Roehm

School of Computer Science



THE UNIVERSITY OF
SYDNEY



Temporal Data is Ubiquitous

- Almost all data is marked with time (period or point)
 - Web stores
 - Data warehousing
 - Medical records, loans, ...
 - Sensor data and time series
 - Transport information
 - ...
- How do we represent and analyse temporal data & time series data?

Some Examples

Weather Observations



```
obsdate, city, temp,wind,wind_dir
2017-01-1,Adelaide, 21.1, 33, SW
2017-01-1,Brisbane, 31.6, 13, NE
2017-01-1,Hobart, 17.3, 9, SSW
2017-01-1,Melbourne, 20.0, 11, SSW
2017-01-1,Perth, 27.9, 19, SW
2017-01-1,Sydney, 24.9, 19, E
2017-01-2,Adelaide, 24.7, 28, SSE
2017-01-2,Brisbane, 30.4, 15, NE
2017-01-2,Melbourne, 16.3, 19, S
```

Water Measurements

Station	Date	Level (m)	Discharge (ml/d)	Temp (C)	EC @ 25C (us/cm)
409017	4/04/2009	2.472	7817.866	20.569	49.823
219018	4/04/2009	0.166	0.371		
409204	5/04/2009	0.637	2358.659	20.633	51.75
409017	5/04/2009	2.389	7744.308		52.604
409204	6/04/2009	0.637	2390.783	20.125	51.5
409017	6/04/2009	2.403	7861.138	20.715	52.688
219018	6/04/2009	2.403	0.239		
409204	7/04/2009	0.637	2358.659	19.157	50.24
409017	7/04/2009	2.39	7649.435	20.281	44.875
219018	7/04/2009	0.166	0.119		
409204	4/04/2009	0.653	2409.639	20.181	55.906
409204	8/04/2009	0.628	2271.601	18.583	49.417
409017	8/04/2009	2.368	7299.264	19.925	40.771
219018	8/04/2009	0.162	0.091		
409204	9/04/2009	0.615	2185.795	18.604	49.094

Big Data also means Historic Data

- Businesses can enhance their customer relationships by embracing many of Arbesman’s “long data” ideas [Arbesman: *Stop Hyping Big Data and Start Paying Attention to ‘Long Data’* (Wired 2013)]:
 - Big data is more about taking a slice in time across many different channels. But long data involves **looking at information on a much longer timescale**.
 - Many companies have data that goes back **10, 20, or 30 years**. A longer-term view can provide information that businesses might miss if they examine data that only goes back five years or less.
 - There’s no international standard definition of big data. But the most popular description – Gartner’s 3V of high volume, high velocity, and high variety – doesn’t explicitly include **historical information as a key component of big data**.

[Jeff Bertolucci, InformationWeek 2013]

Temporal Databases



THE UNIVERSITY OF
SYDNEY

Temporal Database Management System

- A temporal database management system (DBMS) is a DBMS that provides built-in support for the time dimension
 - special facilities for storing, querying, and updating data with respect to time.
- A temporal DBMS can distinguish between historical data, current data, and data that will be in effect in the future.
- The intent of a temporal database management system is to reason with time.
 - A temporal DBMS provides a temporal version of SQL

Temporal Databases

- A temporal database stores data relating to time periods and time instances.
 - It offers temporal data types and stores information relating to past, present and future time.
 - For example, it stores the history of a stock or the movement of employees within an organization.
- Differences between a temporal database and a conventional database
 - A temporal database maintains data with respect to time and allows time-based reasoning
 - A conventional database captures only a current snapshot of reality.
 - cannot directly support historical queries about past status and cannot represent inherently retroactive or proactive changes.
- Examples: TimescaleDB, InfluxDB

Concepts in Temporal Databases

理解每个terminology的意思

– Temporal data types

- Instant | Point
- Interval
- Period

something happening at an instant of time

a length of time; a duration

an anchored duration of time (a recurring interval of time, e.g., calendar year)

– Kinds of time

- User-defined time
- Valid time
- Transaction time

an uninterpreted time value 用户自定义时间, 数据库不解释其意义

when a fact was true in the modelled reality 有效时间, 事实在现实中成立的时间

when a fact was stored in the database 事务时间, 事实被记录入数据库的时间

– Temporal statements

- Current
- Sequenced
- Nonsequenced

now

at a specific instant or period of time 针对某个时间点或时间段的查询
ignoring time

Cf. https://docs.teradata.com/r/Enterprise_IntelliFlex_VMware/Temporal-Table-Support

Temporal Data Types

- Instant data types:
 - DATE
 - SQL-92: day, month and year of a time instant (from year 1 to 9999)
 - Postgresql: date (no time of day) from 4713 BC to 5874897 AD
 - TIMESTAMP
 - SQL-92: date + time with variable resolution of fractions of a second (default: 1ms)
 - Postgresql: date + time of same range as DATE with 1 ms resolution; optional time zone
 - TIME
 - SQL-92: hours, minutes, seconds and optional fractional digits of second
 - not really a time instant (no date!); in PostgreSQL with 1ms resolution
- Interval data types:
 - Various specification options, eg. Year-Month Intervals: 1-2 (1 year 2 months)

Cf. <https://www.postgresql.org/docs/current/static/datatype-datetime.html>

Temporal Data Types

- Period data types:
 - It has a beginning bound (defined by the value of a beginning element) and an ending bound.
 - Beginning and ending elements can be DATE, TIME, or TIMESTAMP types, but both must be the same type.
 - The duration that a period represents starts from the beginning bound and extends up to, but does not include, the ending bound.

```
CREATE TABLE Policy (
    Policy_ID      INTEGER,
    Customer_ID    INTEGER,
    Policy_Type    CHAR(2) NOT NULL,
    Policy_Details CHAR(40),
    Validity       PERIOD(DATE)
)
```

```
INSERT INTO Policy VALUES (
    541008,
    246824626,
    'AU',
    'STD-CH-345-NXY-00',
    PERIOD(DATE '2009-10-01', DATE '2009-10-31')
);
```

- SQL supports time *instants* and *intervals* (but no periods)


PostgreSQL不支持perio

Cf. https://docs.teradata.com/r/Enterprise_IntelliFlex_VMware/Temporal-Table-Support

Kinds of Time

- **User-defined time:** an *uninterpreted time value*
 - E.g. a birthdate or a publication time
- **Valid time** records the time when a fact is true in the real world.
 - Can move forward and backward
- **Transaction time** records the history of database activity.
 - Only moves forward (as you cannot go back in history and change things – alas!)
 - Therefore, allows rollback (very useful for auditing)

Two Interpretations of Time



- Database time (or **Transaction time**) only moves forward.
- Fact time (or **Valid time**) can move forward or backward
 - Are we changing history here?
No, more our ‘perception’ of history...

Temporal Tables

- Temporal tables store and maintain information with respect to time, and can process statements and queries that include time-based reasoning.
- Temporal tables include one or two special columns, which store time information:
 - A transaction-time column
 - A valid-time column
 - Each of these time dimensions is represented by a column with a period data type.

```
CREATE TABLE Policy (
    Policy_ID          INTEGER,
    Customer_ID        INTEGER,
    Policy_Type        CHAR(2) NOT NULL,
    Policy_Details     CHAR(40),
    Validity           PERIOD(DATE) NOT NULL AS VALIDTIME
)
```

Temporal Tables

- A table with a transaction-time column is called a **transaction-time table**.
- A table with a valid-time column is called a **valid-time table**.
- A table with both a transaction-time and a valid-time column is called a **bitemporal table**.
- A table with neither a transaction-time nor a valid-time column is a **nontemporal table**.

Handling of NOW or UC (Until Changed)

- What should the timestamp be for current data? (valid now until changed)
 - Minimum possible timestamp (*min-timestamp approach*)
 - E.g. '00-00-0000' - quite counter-intuitive
 - NULL value
 - Complicates queries: Comparisons with NULL return FALSE in SQL...
 - Use of NULL as 'unknown' is no longer possible...
 - Maximum possible timestamp (*max-timestamp approach*)
 - E.g. '31-12-9999'
 - Simulates 'the end of time' or 'forever'
 - **Point representation**
 - Set `start_time = end_time`
- **Max-timestamp approach is extensively used**
- Warning: PostgreSQL supports a 'now' keyword which refers to the transaction's start time, not to 'valid until changed...'; instead use 'infinity' in PostgreSQL

3种办法设置 数据库中“当前有效数据”的时间范围

把 `end_time` 填成最小可能值，表示数据仍然有效

结束时间为 空，表示数据仍然有效

给当前数据的结束时间设置最大值，表示数据仍然有效

Temporal Statements

- Current 只关心现在的状态，不关心历史变化
 - “What is now?” CURRENT VALIDTIME **SELECT * FROM Policy;**
 - E.g. “*How many products do we currently have in stock?*”
- Sequenced 查询数据的历史变化序列。可以看到每个数据在不同时间点或时间段的状态
 - “What was, and when?”
 - E.g. “*Give the sequence of how many product were in stock.*”
 - Or “*When did the stock level fall below X in the past?*”
- Nonsequenced 查询过去曾经出现过的数据，但不关注顺序
 - “What was at any time?”
 - E.g. “*How many products A did we have at any time in stock?*”

Temporal Support in Databases



THE UNIVERSITY OF
SYDNEY

Temporal Support in Databases

- Limited support for temporal data management in DBMSs
 - Conventional (non-temporal) DBs represent a **static snapshot**
 - Management of temporal aspects is implemented by the application
 - Adds additional complexity to application programs
 - Some time data types and functions available in SQL, e.g., DATE, TIME, DATEADD(), DATEDIFF()
 - SQL:2011 added support for temporal tables
 - Still **very limited query support**

How to represent time in (traditional) databases?

- Basic building blocks: SQL data types DATE, TIMESTAMP
- Modelled as a set of time instants/points with a partial order
- Can add attributes of these time data types to any table, interpretation is application-specific
 - Example 1: Person (id INT, name TEXT, birthday DATE);
 - Example 2: Project (id INT, title TEXT, start DATE);
 - Example 3: TripLog(user INT, car INT, when TIMESTAMP, distance INT);

Date and Time in SQL

SQL Type	Example	Accuracy	Description
DATE	'2012-03-26'	1 day	a date (some systems incl. time)
TIME	16:12:05	ca. 1 ms	a time, often down to nanoseconds
TIMESTAMP	2012-03-26 16:12:05	ca. 1 sec	Time at a certain date: SQL Server: DATETIME
INTERVAL	'5 DAY'	years - ms	a time duration

- Comparisons
 - Normal time-order comparisons with '=', '>', '<', '<=', '>=' , ...

- Constants
 - **CURRENT_DATE** db system's current date 和数据库在的时区的时间相同
 - **CURRENT_TIME** db system's current timestamp

- Example:

```
SELECT *
  FROM Measurements
 WHERE obsdate < CURRENT_DATE;
```

Date and Time in SQL (cont'd)

- Database systems support a variety of date/time related ops
 - Unfortunately, not very standardized – a lot of slight differences
- Main Operations
 - **EXTRACT(component FROM date)**
 - e.g. EXTRACT(year FROM startDate)
 - **DATE string** (Oracle syntax: TO_DATE(*string,template*))
 - e.g. DATE '2012-03-01'
 - Some systems allow templates on how to interpret *string*
 - **+/- INTERVAL:**
 - e.g. '2012-04-01' + INTERVAL '36 HOUR'
- Many more -> check database system's manual

Semi-Temporal Databases

- We want to keep track since when a current fact is true.
 - We add timestamp defining the *period of validity* since a point in time
`Employees (tfn, name, city, ..., since)`
 - Q: Meaning of SINCE attribute? => typically [since, now)
'Ever since day SINCE, supplier s has been under contract'
- Employees has become a '**semi-temporal' table**

tfn	name	city	since
E1	Employee1	London	1996-06-01
E2	Employee2	Sydney	1996-08-01
E3	Employee3	New York	2000-10-01
E4	Employee4	Frankfurt	2005-12-31
E5	Employee5	Singapore	2006-12-31

Adding History

- We want to keep the history of when which employees were assigned to what positions.
 - Need to add data about the *period of validity* to a EmpPosition table
 - As there is no ‘time period’ data type in SQL, we have to simulate it with two DATEs (V_{ts} and V_{te}).
`EmpPosition (tfn, pos, start_date, end_date)`
 - Q: Open or closed time interval? => typically, V_{ts} is included but V_{te} is not [V_{ts}, V_{te})
- EmpPosition has become a **valid-time table**

tfn	pos	start_date	end_date
E1	assistant	1996-06-01	1996-08-01
E1	consultant	1996-08-01	1998-04-01
E1	senior consultant	1998-04-01	1998-10-01
E2	manager	2000-10-01	9999-12-31
E3	sysadmin	2005-12-31	9999-12-31

<- side issue: how to represent ‘till now’?

SQL:2011 added support for Valid and Transaction Time

- SQL:2011 adds period definitions as metadata to tables (and not as a new data type!)
- A period is a conceptual grouping of a physical start time and end time attribute/column
- For instance, a period EPeriod built from existing attributes EStart and EEnd
... **PERIOD FOR** EPeriod (EStart, EEnd) ...
- A period is by default a half-closed interval [EStart, EEnd)
- Distinction between
 - application time (= valid time)
 - system time (= transaction time)
- So far supported only by few commercial DBMS, such as IBM DB2

Example: Temporal Aggregation

- **Input:** external project funding

Name	Dept	Budget	TimePeriod
Tom	CS	100 K	[Feb, Jul]
Tim	CS	180 K	[May, Jul]
Joe	M	150 K	[Apr, Aug]

- **Query:** Amount of external funding per department?
- **Output:** Temporal aggregation (at each point in time)

Dept	Budget	TimePeriod
CS	100 K	[Feb, Apr]
CS	280 K	[May, Jul]
M	150 K	[Apr, Aug]



[Example from: J. Gamper, eBISS 2017]

PostgreSQL

- Key features
 - Supports all standard SQL temporal data types (points and intervals)
in addition:
 - Range types (periods)
 - Indexes on range types
 - Temporal constraints using GiST indexes
 - Predicates and functions for range types
- <http://www.postgresql.org/docs/devel/static/rangetypes.html>

Tables with Range Types in PostgreSQL

- New data type **DATERANGE** to represent (time) intervals

```
CREATE TABLE Emp (
    eName      VARCHAR,
    ePeriod    DATERANGE,
    eDept      VARCHAR
);
INSERT INTO Emp VALUES ('Anton', '[2010,2014]', 'SCS');
```

Indexes on Range Types in PostgreSQL

- Provides a powerful, extensible index: **Generalised Search Tree** (GiST)
 - On periods and combinations of period and other attributes
(eg. GiST is also very useful for spatial data)

```
CREATE INDEX Emp_idx ON Emp USING GIST (EPeriod)
```

```
CREATE INDEX Emp_idx ON Emp USING GIST (EName, EPeriod)
```

- Improves range indexing for many predicates
 - e.g., EQUAL, OVERLAP, LESS THAN, CONTAINS, . . .

Query Examples on PostgreSQL Range Types

- SQL with predicates, e.g., OVERLAPS (&&), BEFORE, AFTER

```
SELECT *
  FROM Emp JOIN Dept ON EPeriod && DPeriod
```

```
SELECT *
  FROM Dept
 WHERE upper_inf(DPeriod) = TRUE
```

- Some additional functions on ranges: UNION (+), INTERSECTION (*), DIFFERENCE (-)

```
SELECT '[2010, 2013)' + '[2012, 2015)'
```

Temporal Support in Current DBMS

age > 18

DBS	Period	Keys	Update	Predicates	Queries
SQL:2011	✓	✓	✓	✓	✗
IBM DB2	✓	✓	✓	✓	✗
MySQL	✗	✗	✗	✗	✗
MS SQL Server	✗	✗	✗	✗	✗
PostgreSQL	✓	~	✗	✓	✗
Oracle	✓	✓	✓	✓	✗
Teradata	✓	✓	✓	✓	~

[J. Gamper, eBISS 2017]

Summary

- Temporal data is ubiquitous – but only limited support in database systems
- Challenges:
 - Limited temporal data types in SQL
 - Limited support for temporal statements / queries
 - Time Periods not part of standard SQL, hence such queries difficult to represent and encode
 - Timeseries support
 - Transaction vs. Valid Time



Representing Time Series Data



THE UNIVERSITY OF
SYDNEY

Timeseries Data

- Series of observations (data points) made sequentially over time at discrete time points, usually at equal intervals
- Examples
 - weather observations
 - trip data
 - road usage
 - sleep patterns
 - cytometry / panel data



Working with Time Series Data

- Data ingress (file import, database access, web scraping, ...)
- Storage
 - Standard file formats, databases
- Analysis
 - In database vs. applications (eg. Python programs)
- Visualisation, reporting

Timeseries Data – Storage Approaches

- Point-based representation
 - Multiple rows with atomic data types
- Sequence-based representation
 - Single row with array of time point data
 - Requires array datatype
 - UNNEST array operation, array_agg with filter
- Dedicated time-series database
 - Eg. TimescaleDB, InfluxDB, or also Scidb

Point-Based Representation in Flat Table

- Each tuple is timestamped with a time point/instant
- Most basic and simple data model
- Timestamps are atomic values and can be easily compared with $=, !=, <, >, \geq, \leq$
- Syntactically different relations store different information
- *If a fact is valid at several time points, multiple tuples are needed*
- Good for time domain analysis
- Allows for indexing

```
CREATE TABLE Observations (
    ts           TIMESTAMP,
    station      INT,
    temperature  FLOAT,
    wind         CHAR(2)
) ;
```

ts	station	temperature	wind

Working with Point-Based Representations

- Individual *INSERT*, *UPDATE* statements (typically append-only)
- OLAP-style queries for filtering, aggregation, grouping

```
SELECT station, AVG(temperature)  
      FROM Observations  
     WHERE ts BETWEEN start_ts AND end_ts;
```

- Careful with granularity and timezone of data
- Joins to dimension tables
- Indexing – typically B-Tree based

```
CREATE INDEX TimeIdx ON Observations(ts);
```

Sequence-Based Representation in Nested-Table

- Single row with array of time point data
- Requires array datatype
 - => outside 1NF idea of relational model
- How to represent discrete time points now?
 - Part of the array? So tuple (ts, value)?
 - Or implicit by position?
 - Then need to make sure that different arrays represent same data and time points
- Special functions needed to
 - Create arrays from set of data array_agg()
 - Pivot data back into sets UNNEST

```
CREATE TABLE Observations2 (
    station      INT,
    temperature  FLOAT[],
    wind         CHAR(2)[],
    obs_start    TIMESTAMP,
    obs_end      TIMESTAMP
) ;
```

station	temp	wind	start, end
	[...]	[...]	,

Cf. <https://www.postgresql.org/docs/9.2/functions-array.html>

Working with Array-Data

- INSERTs with array data

```
INSERT INTO Observations2  
VALUES(4711, '{79,82,90,69,75,80,81}', '2018-01-01', ...);
```

- UPDATE of array content

```
UPDATE Obs2 SET temperature[4] = 22.5 WHERE station=..
```

- Indexing => GIN (generalised inverted index)

```
CREATE INDEX Timelidx2 ON Observations2 USING GIN("ts")
```

- Querying

- Turning arrays to rows: Unnesting

数组 → 行 (Unnest)

```
SELECT UNNEST(temperature) FROM Observations2 WHERE ...
```

行 → 数组 (array_agg)

- Turn rows into array data: array_agg()